deeplearning.ai

Introduction to ML strategy

Why ML Strategy?

Motivating example






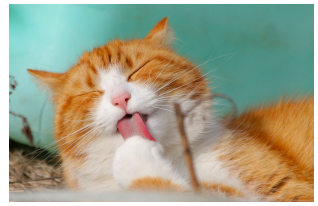

90%

Ideas:

- Collect more data
- Collect more diverse training set
- Train algorithm longer with gradient descent
- Try Adam instead of gradient descent
- Try bigger network
- Try smaller network
- Try dropout
- Add L_2 regularization
- Network architecture
  - Activation functions
  - # hidden units
  - …

Andrew Ng

deeplearning.ai

Introduction to ML strategy

Orthogonalization

TV tuning example

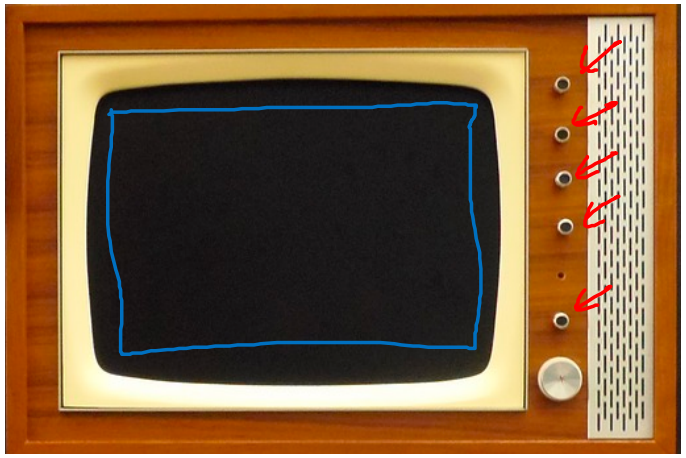

Orthogonalization

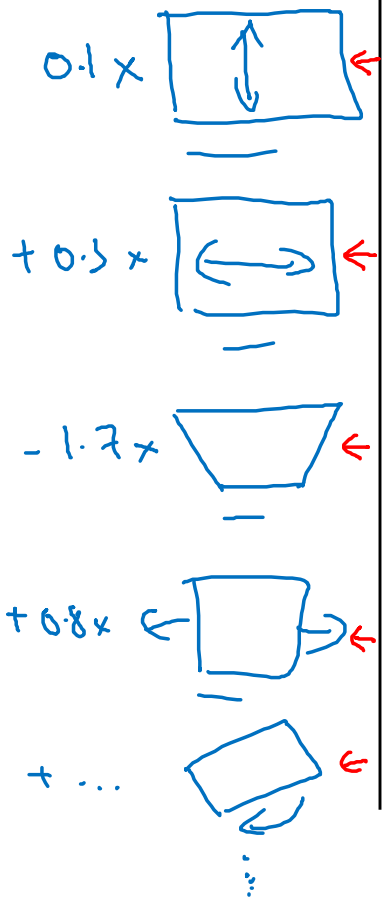

Car

→ Steering

→ { Accelerator, Braking }

→ 0.3 x angle − 0.8 speed

→ 2 x angle + 0.9 speed.

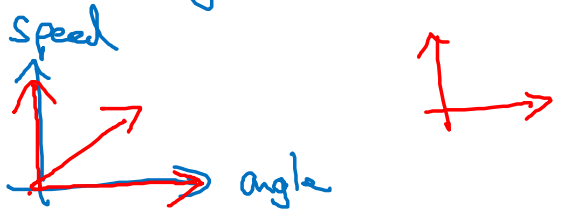

Andrew Ng

Chain of assumptions in ML

→ Fit training set well on cost function

(≈ human-level performance)

- bigger network
- Adam
- ...

early stopping

↓

→ Fit dev set well on cost function

- Regularization
- Bigger training set

↓

→ Fit test set well on cost function

Bigger dev set

↓

→ Performs well in real world

(Happy cat pic app users.)

Change dev set or cost function

Andrew Ng

deeplearning.ai

Setting up
your goal

Single number evaluation metric

Using a single number evaluation metric

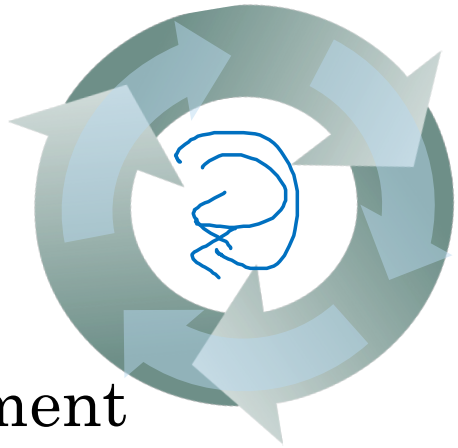

Of examples recognized as cat, what % actually are cats?

What % of actual cats are correctly recognized

Classifier	Precision	Recall
A	95%	90%
B	98%	85%

F_1 score = "Average" of P and R.

$\left(\frac{2}{\frac{1}{P}+\frac{1}{R}}\right.$. "Harmonic mean")

Dev set + Single number evaluation metric

real

Speed up iterating

Another example

Algorithm	US	China	India	Other
A	3%	7%	5%	9%
B	5%	6%	5%	10%
C	2%	3%	4%	5%
D	5%	8%	7%	2%
E	4%	5%	2%	4%
F	7%	11%	8%	12%

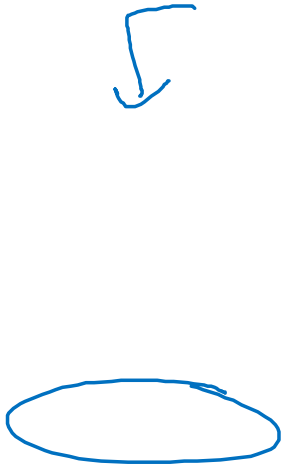

deeplearning.ai

Setting up
your goal

Satisficing and optimizing metrics

Another cat classification example

optimizing ↓ ↓ satisficing

Classifier	Accuracy	Running time
A	90%	80ms
B	92%	95ms
C	95%	1,500ms

Cost = accuracy − 0.5 × running Time

Maximize accuracy

Subject to running Time ≤ 100 ms.

N metrics: 1 optimizing

N−1 satisficing

Wakewords / Trigger words

Alexa, OK Google,

Hey Siri, nihaobaidu

你好百度

accuracy.

#false positive

Maximize accuracy.

s.t. ≤ 1 false positive

every 24 hours.

Andrew Ng

deeplearning.ai

Setting up your goal

Train/dev/test distributions

Cat classification dev/test sets

development set, hold out cross validation set

Regions:

- US
- UK
- Other Europe
- South America
- India
- China
- Other Asia
- Australia

Dev

Test

Randomly shuffle into dev/test

dev set + metric

Idea

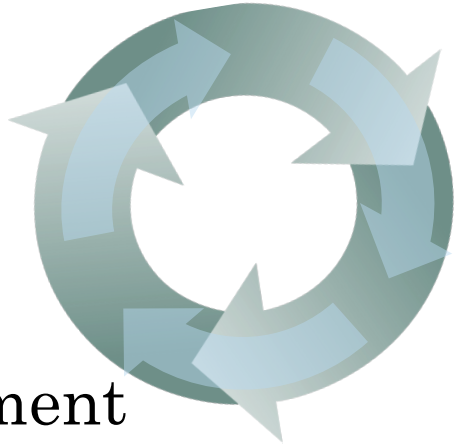

Experiment

Code

Andrew Ng

True story (details changed)

Optimizing on dev set on loan approvals for medium income zip codes

$x \rightarrow y$ (repay loan?)

Tested on low income zip codes

~3 month

Guideline

Same distribution

Choose a dev set and test set to reflect data you expect to get in the future and consider important to do well on.

training

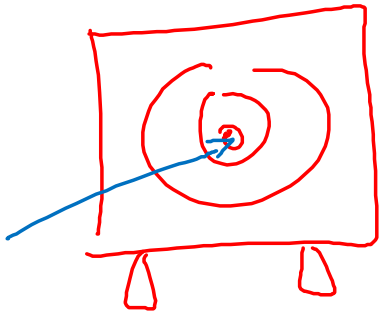

dev

metric

test

deeplearning.ai

Setting up
your goal

Size of dev and test sets

Old way of splitting data

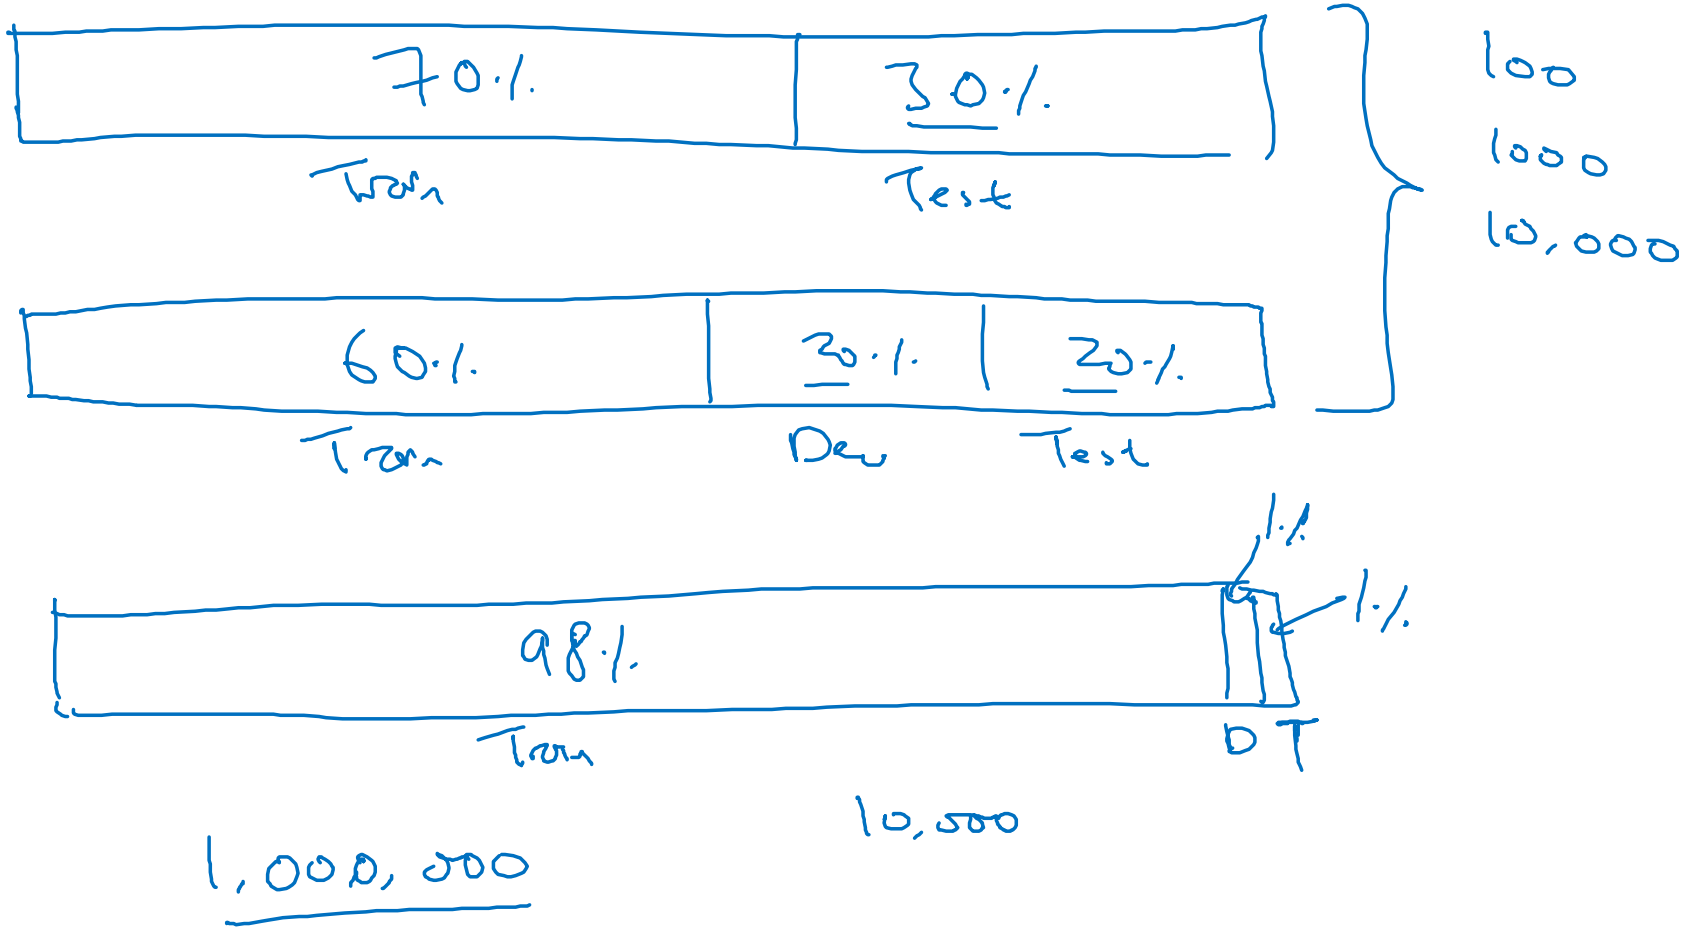

Andrew Ng

Size of dev set

A B

Set your dev set to be big enough to detect differences in algorithm/models you're trying out.

A B

97% → 97.1%

0.1%

100: small

↳ 1%

1,000

10,000

100,000

0.01%

0.001%

Online advertising

Size of test set

- Set your test set to be big enough to give high confidence in the overall performance of your system.

10,000

100,000

Train + Dev.

No test set.

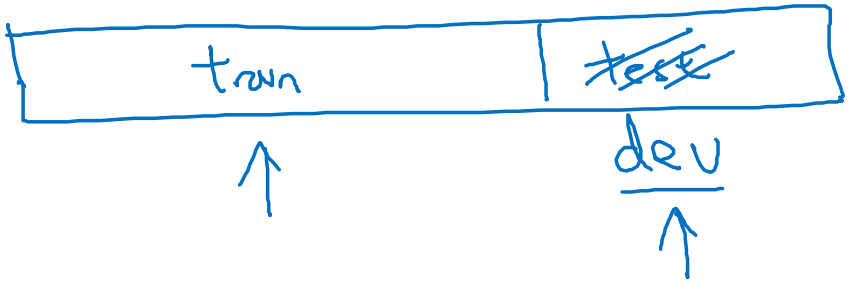

deeplearning.ai

Setting up your goal

When to change dev/test sets and metrics

Cat dataset examples

Metric + Dev : Prefer A
You/users : Prefer B.

→ Metric: classification error

Algorithm A: 3% error → pornographic

✓ Algorithm B: 5% error

$$\text{Error}: \frac{1}{\sum w^{(i)}} \frac{1}{m_{dev}} \sum_{i=1}^{m_{dev}} w^{(i)} \mathbb{1}\{y_{pred}^{(i)} \neq y^{(i)}\}$$

$y_{pred}^{(i)}$: predicted value (0/1)

$$\rightarrow w^{(i)} = \begin{cases} 1 & \text{if } x^{(i)} \text{ is non-porn} \\ 10 & \text{if } x^{(i)} \text{ is porn} \end{cases}$$

Orthogonalization for cat pictures: anti-porn

1. So far we've only discussed how to define a metric to evaluate classifiers. ← Place target

2. Worry separately about how to do well on this metric. ← Aim/shoot at target

$$J = \frac{1}{\sum w^{(i)}} \sum_{i=1}^{m} w^{(i)} \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Another example

Algorithm A: 3% error

Algorithm B: 5% error

Dev/test

User images

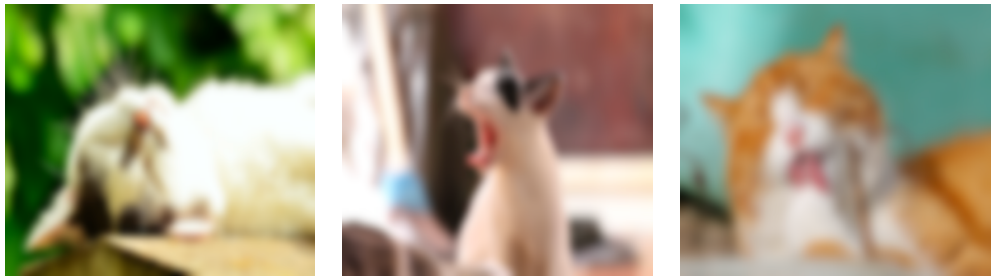

If doing well on your metric + dev/test set does not correspond to doing well on your application, change your metric and/or dev/test set.

deeplearning.ai

Comparing to human-level performance

Why human-level performance?

Comparing to human-level performance

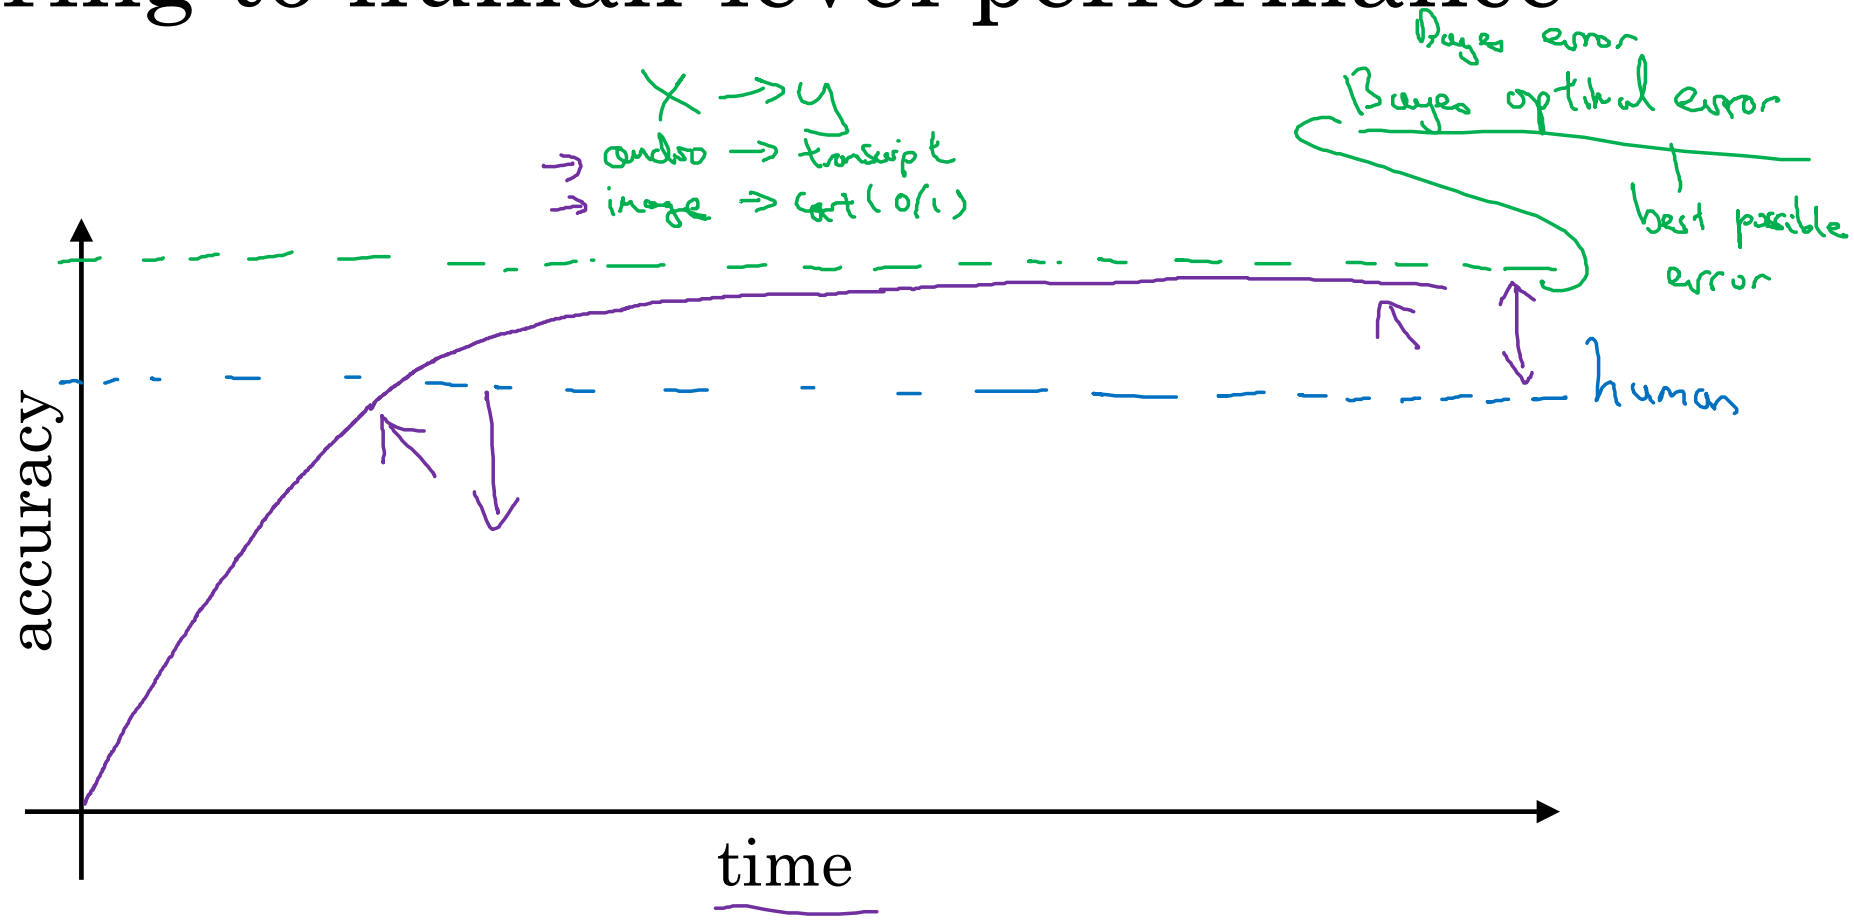

Why compare to human-level performance

Humans are quite good at a lot of tasks. So long as ML is worse than humans, you can:

- Get labeled data from humans. (x, y)
- Gain insight from manual error analysis: Why did a person get this right?
- Better analysis of bias/variance.

deeplearning.ai

Comparing to human-level performance

Avoidable bias

Bias and Variance

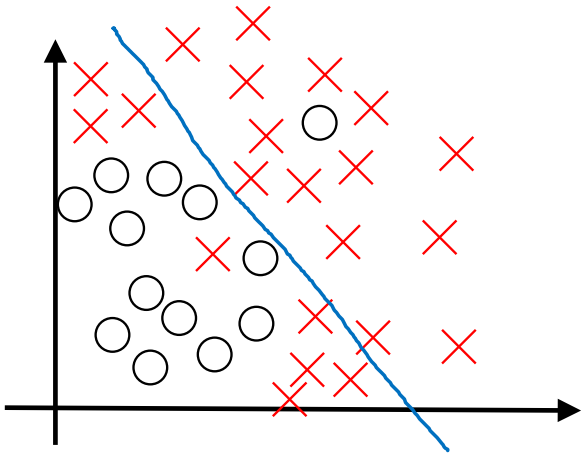

high bias

Underfitting

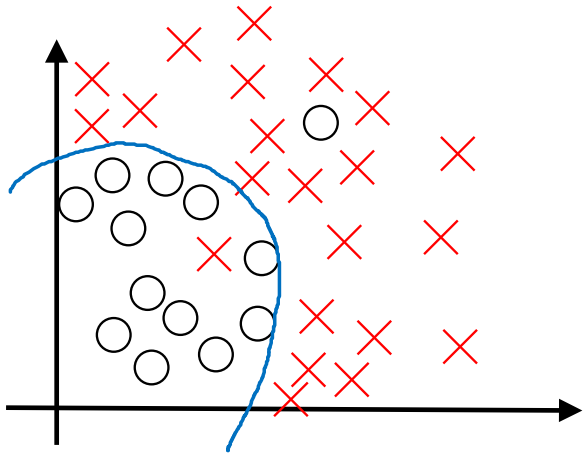

"just right"

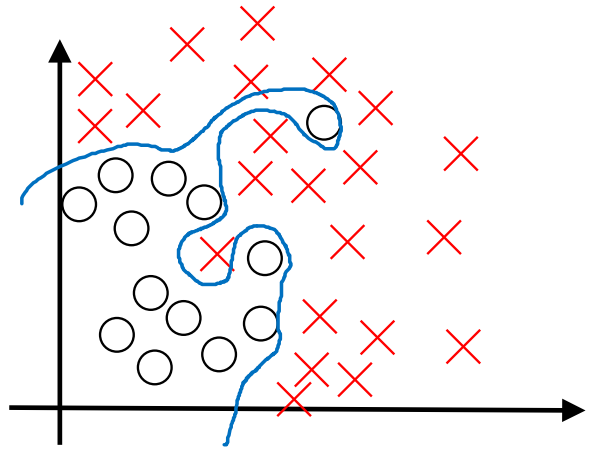

high variance

Overfitting.

Bias and Variance

Cat classification

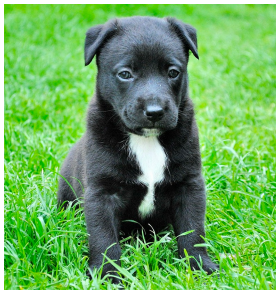

Human-level ≈ 0% ....

Training set error:

Dev set error:

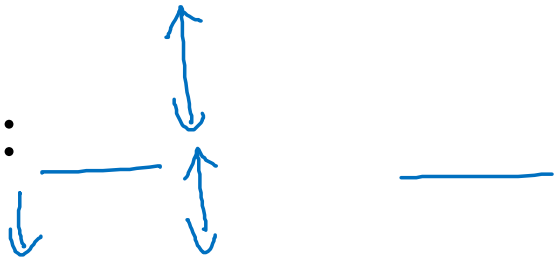

high variance | high bias | high bias, high variance | low bias, low variance

Cat classification example

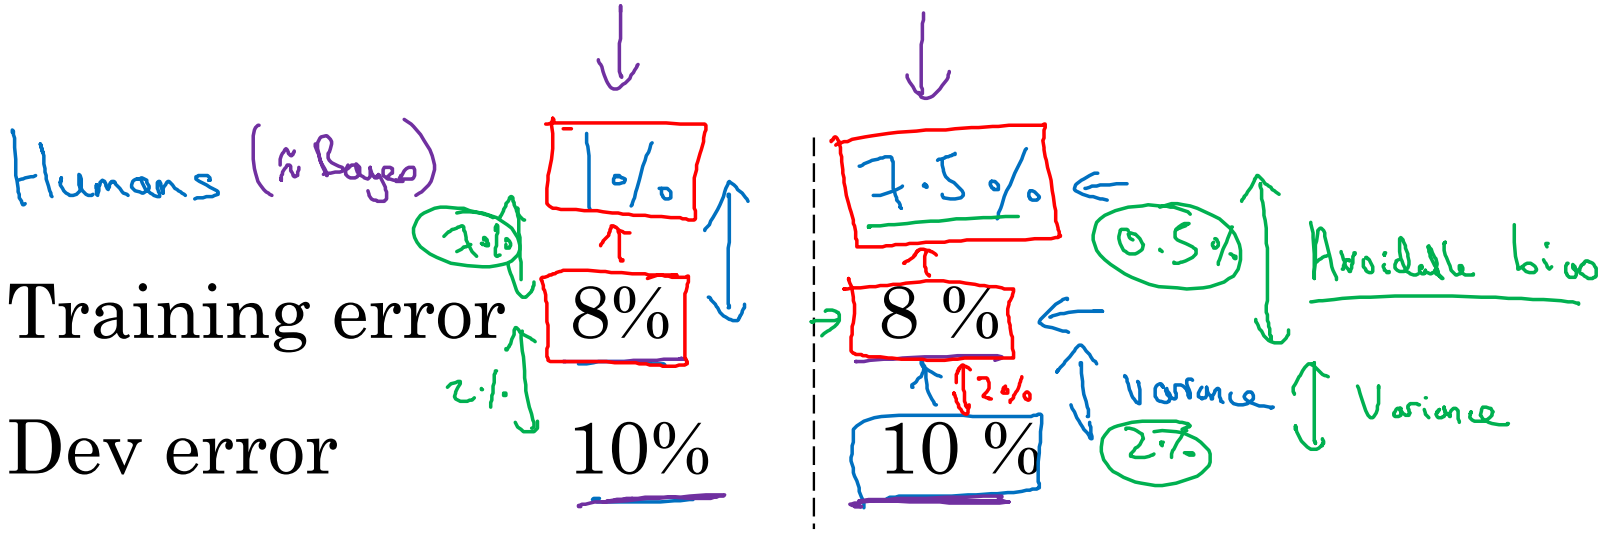

Human-level error as a proxy for Bayes error.

deeplearning.ai

Comparing to human-level performance

Understanding human-level performance

Human-level error as a proxy for Bayes error

Medical image classification example:

Suppose:

(a) Typical human ..... 3 % error

(b) Typical doctor ..... 1 % error

(c) Experienced doctor ..... 0.7 % error

(d) Team of experienced doctors .. 0.5 % error

Bayes error $\leq$ 0.5%

What is "human-level" error?

Error analysis example

Human (proxy for Bayes error)

Avoidable bias

Training error

Variance

Dev error

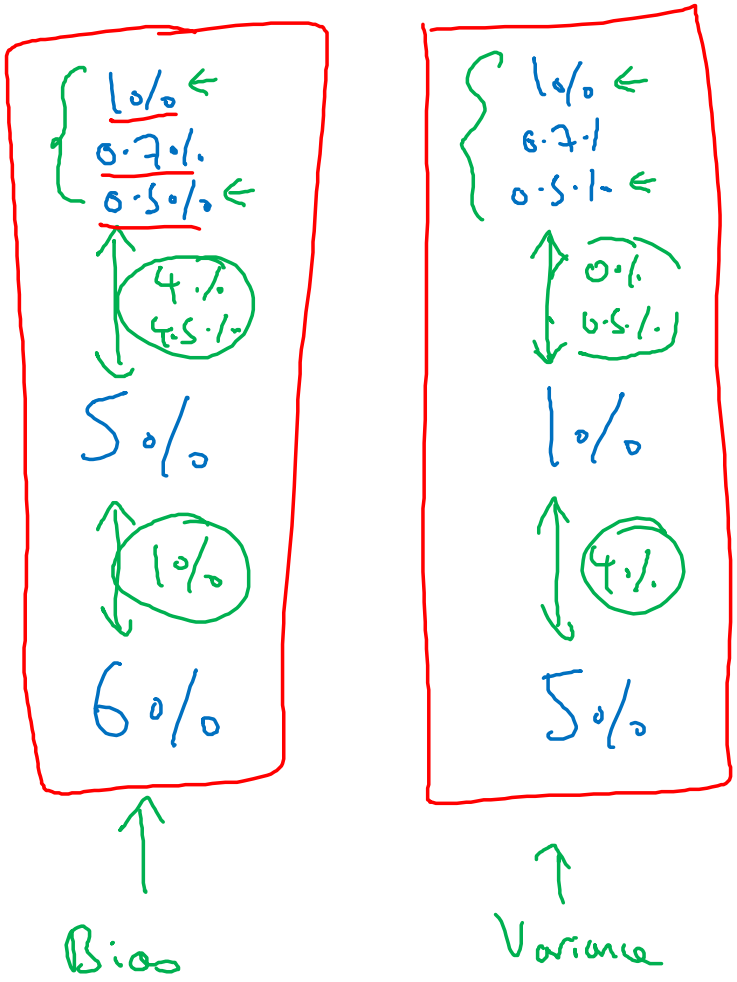

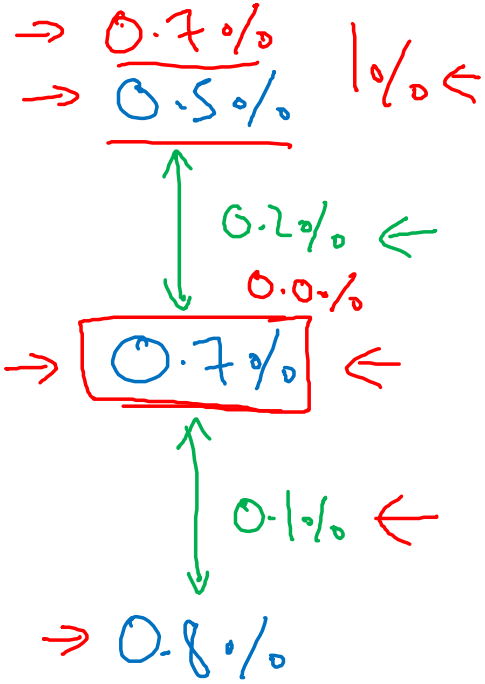

Summary of bias/variance with human-level performance

Human-level error

(proxy for Bayes error)

Training error

Dev error

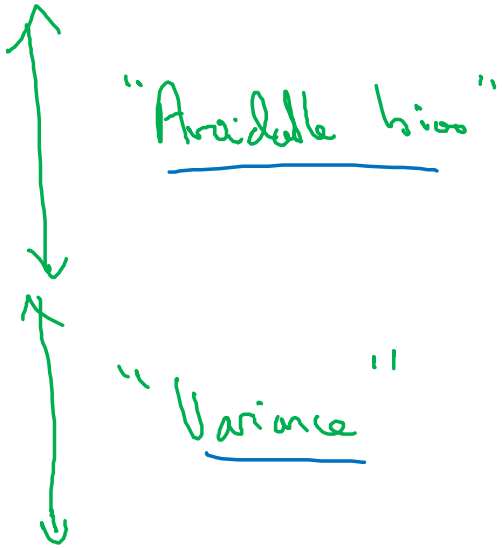

deeplearning.ai

Comparing to human-level performance

Surpassing human-level performance

Surpassing human-level performance

Team of humans

One human

Training error

Dev error

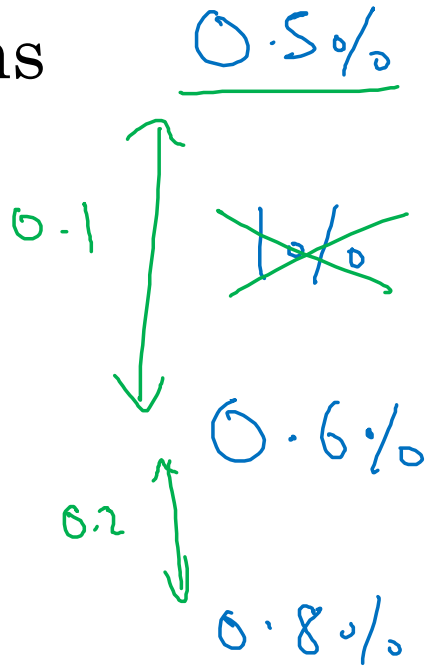

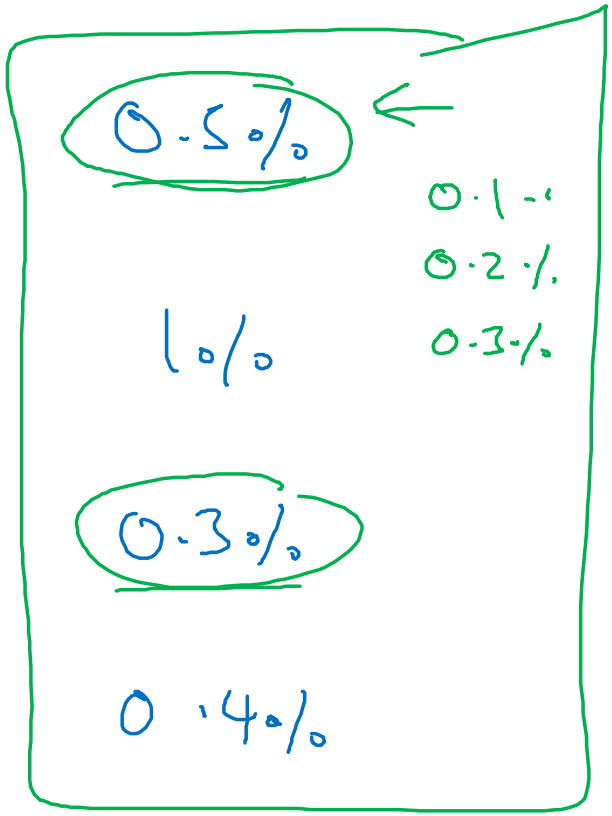

What is avoidable bias?

Problems where ML significantly surpasses human-level performance

- Online advertising
- Product recommendations
- Logistics (predicting transit time)
- Loan approvals

Structured data

Not natural perception

Lots of data

- Speech recognition
- Some image recognition
- Medical
  - ECG, Skin cancer, ...

Andrew Ng

deeplearning.ai

Comparing to human-level performance

Improving your model performance

The two fundamental assumptions of supervised learning

1. You can fit the training set pretty well.

   ~ Avoidable bias

2. The training set performance generalizes pretty well to the dev/test set.

   ~ Variance

Reducing (avoidable) bias and variance

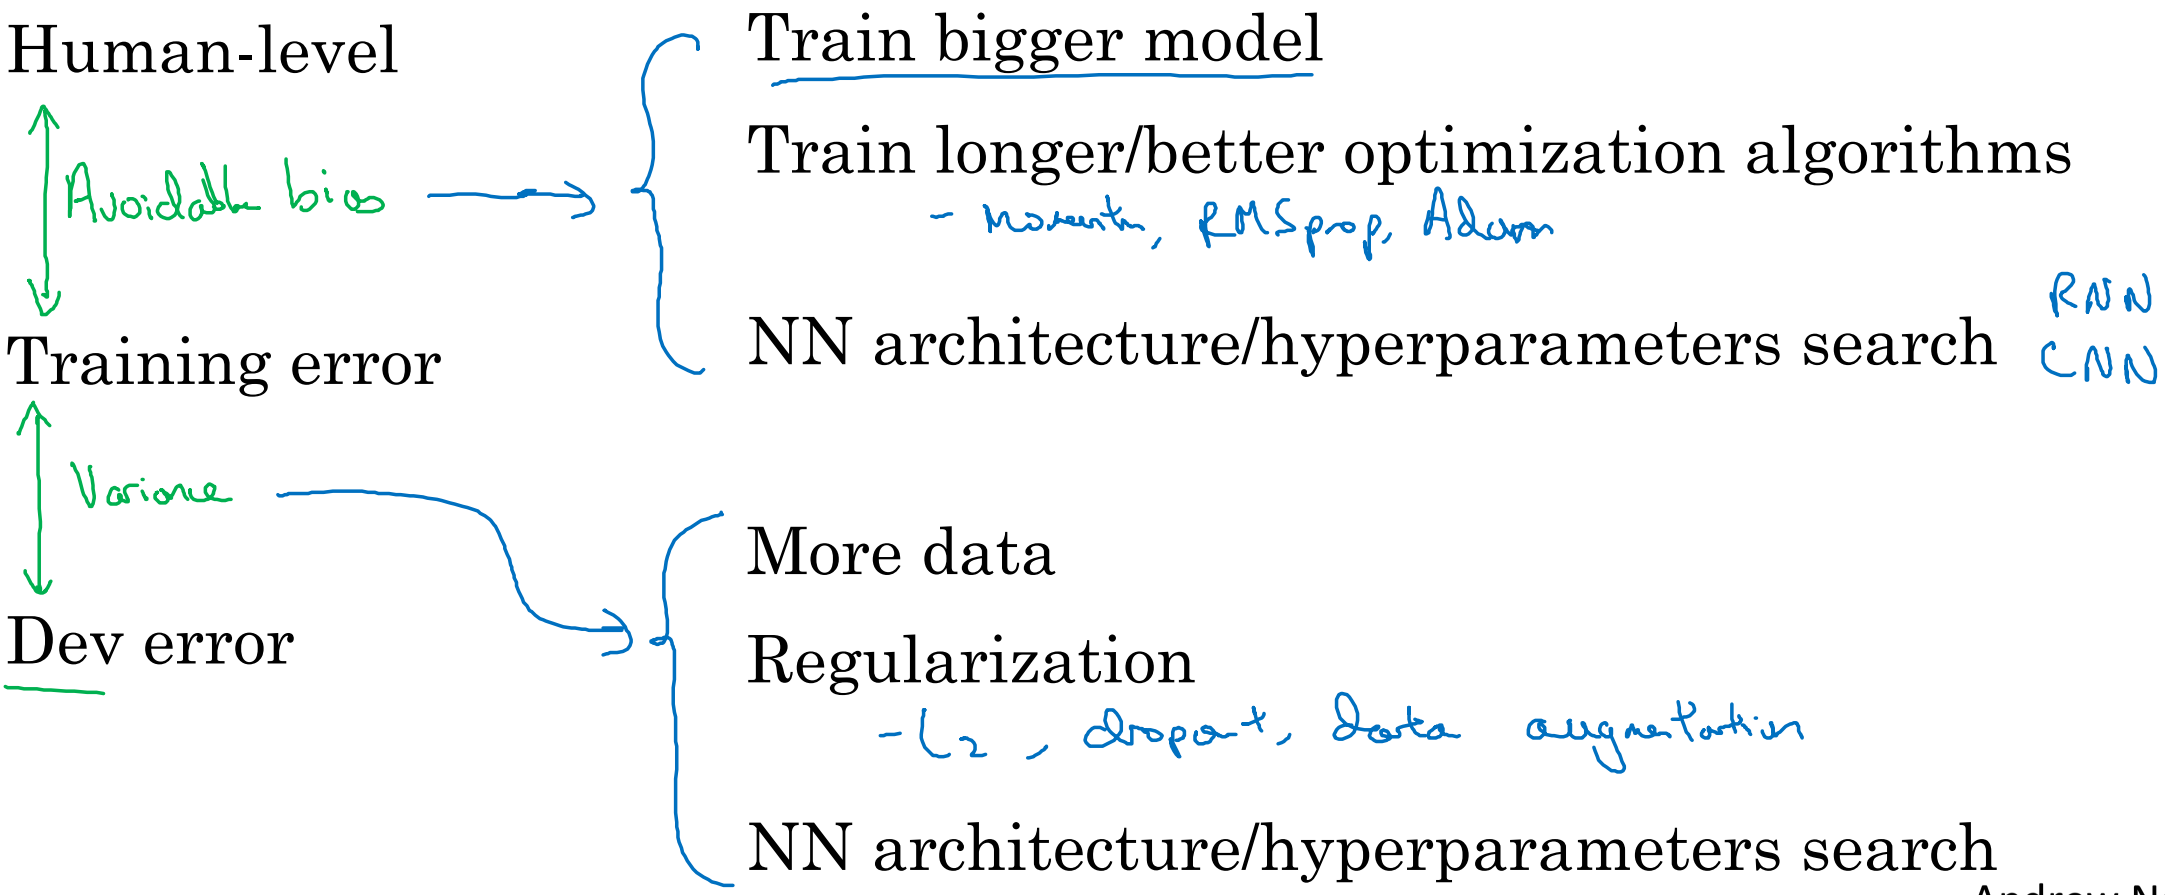

Andrew Ng