



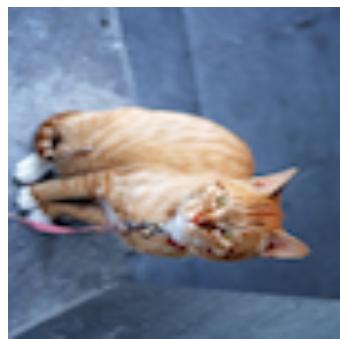
deeplearning.ai

Computer vision

Convolutional
Neural Networks

Computer Vision Problems

Image Classification



→ Cat? (0/1)

64x64

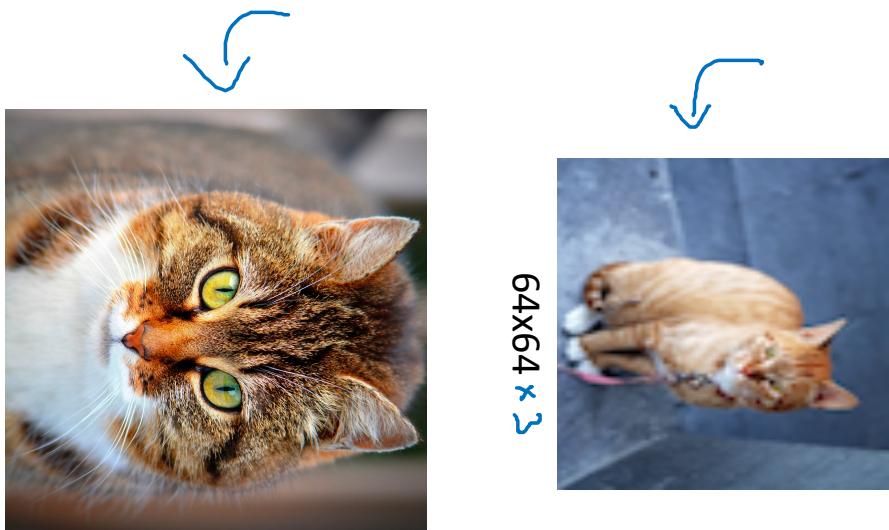
Object detection



Neural Style Transfer ↴



Deep Learning on large images

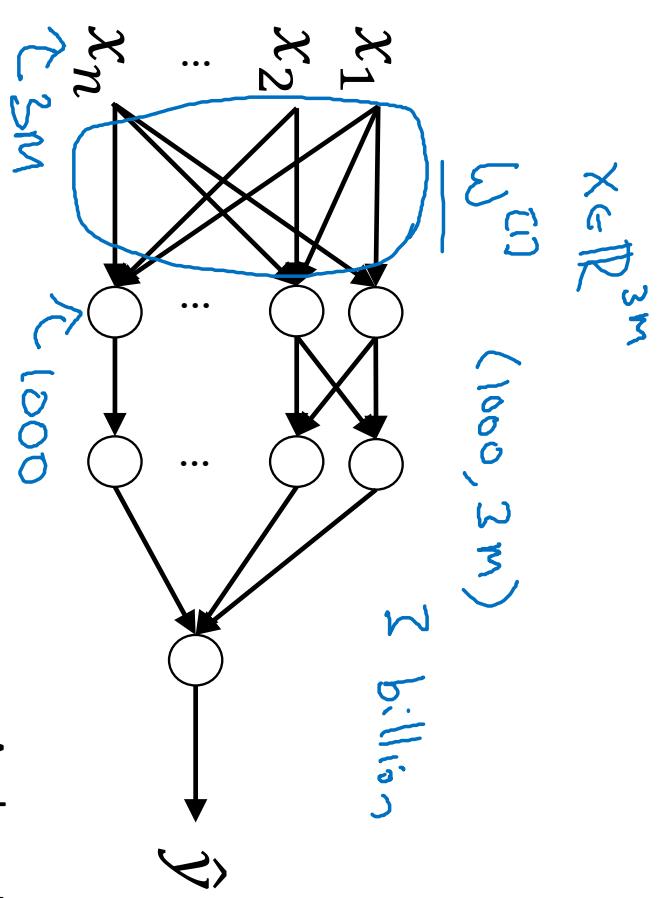


→ Cat? (0/1)

$64 \times 64 \times 3$

12288

$1000 \times 1000 \times 3$
= 3 million



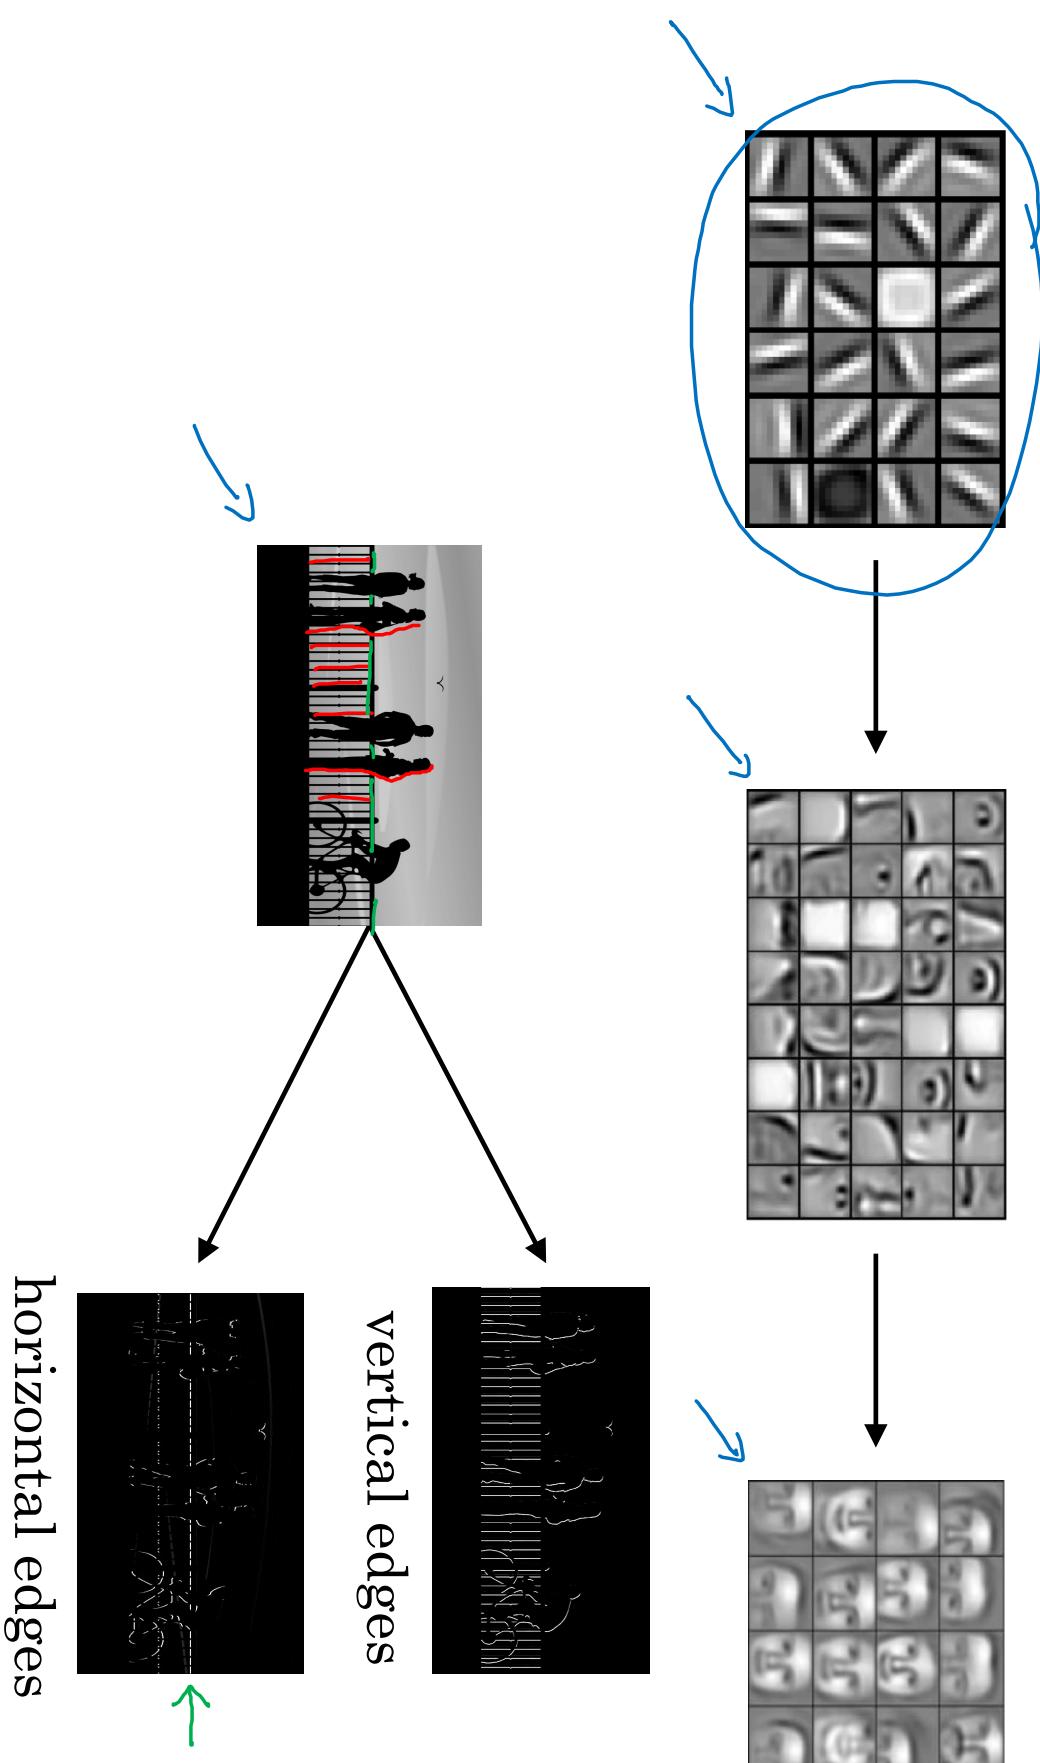


Convolutional Neural Networks

Edge detection
example

deeplearning.ai

Computer Vision Problem



horizontal edges

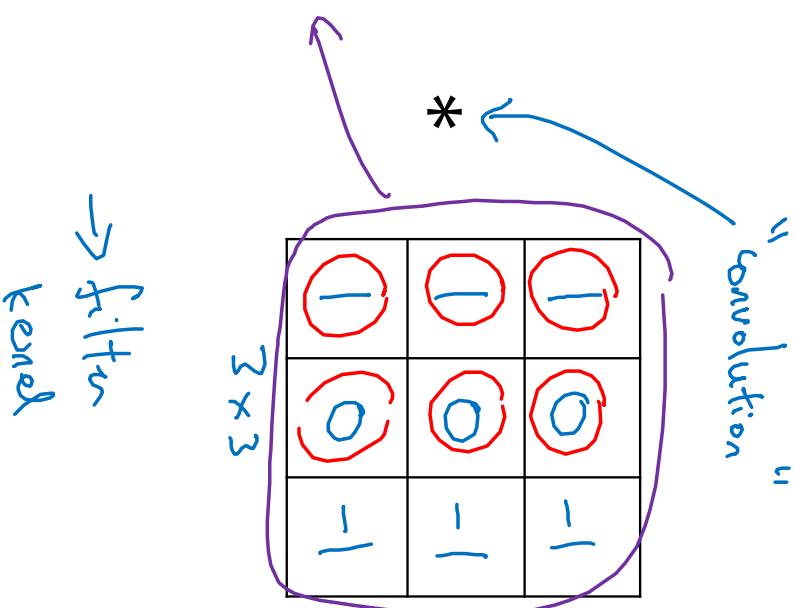
vertical edges

Vertical edge detection

$$\rightarrow 3 \times 1 + 1 \times 1 + 2 \times 1 + 0 \times 0 + 5 \times 0 + 7 \times 0 + 1 \times -1 + 8 \times -1 + 2 \times -1 = -5$$

6×6

3	1	0	1	2	7	4
1	5	0	1	0	-1	-1
2	7	0	1	0	1	1
0	1	3	1	1	3	-1
2	3	7	1	7	8	-1
4	2	1	6	2	8	9

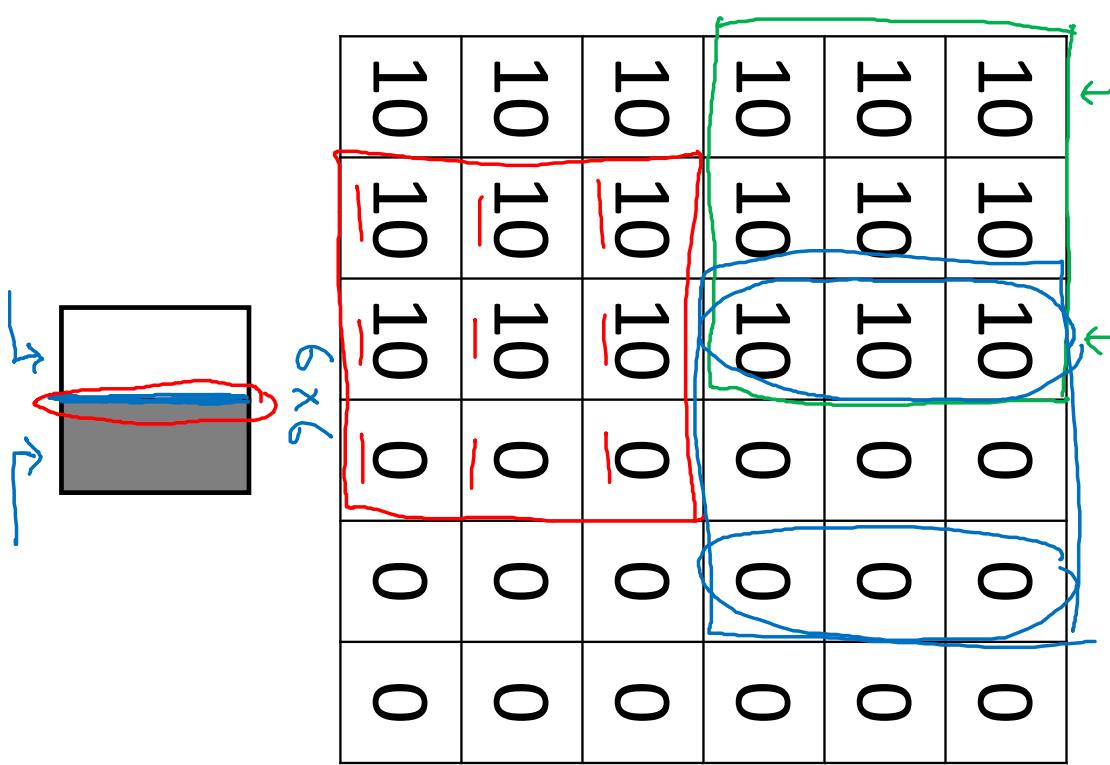


=

4×4

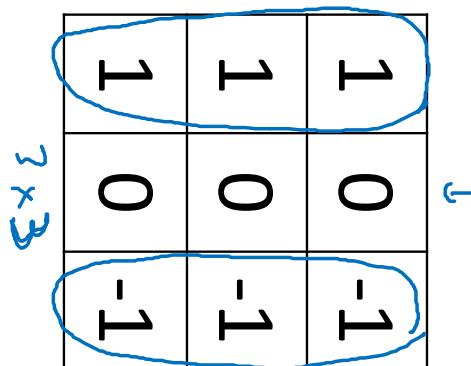
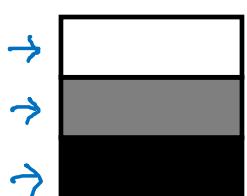
-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

Vertical edge detection

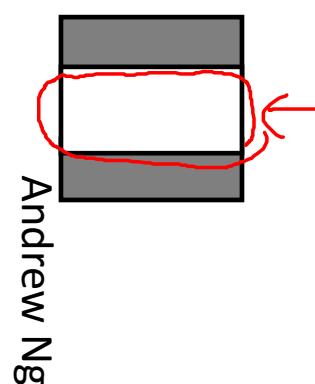


*

*



=



The output image has a maximum value of 30, indicated by a red box. The input image is also shown with a red box around the vertical edge.

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



deeplearning.ai

Convolutional Neural Networks

More edge detection

Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0



0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

*

*

1	0	-1
1	0	-1
1	0	-1

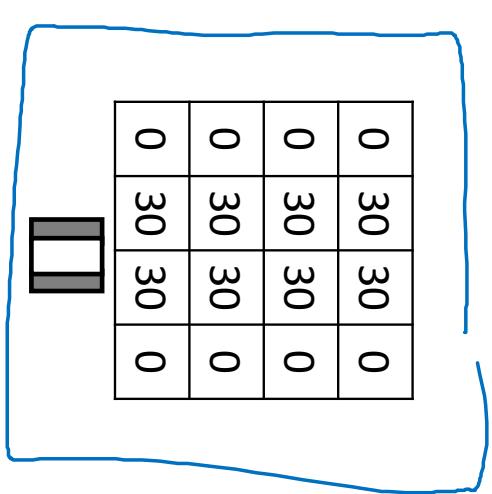
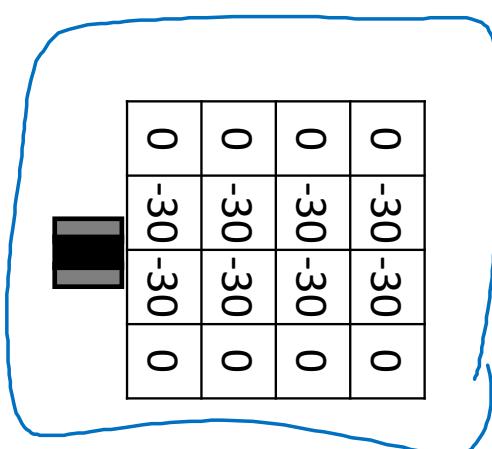
0	30	30
0	30	30
0	30	30

=

=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0



Vertical and Horizontal Edge Detection

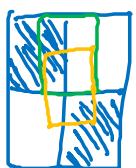


1	0	-1
1	0	-1
1	0	-1

Vertical

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

6x6



*

1	1	1
0	0	0
-1	-1	-1



1	1	1
0	0	0
-1	-1	-1

Horizontal

0	0	0	0	0	0
30	10	-10	-30		
30	10	-10	-30		
0	0	0	0		
0	0	0	0		

=

Learning to detect edges

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

1	0	-1
1	0	-1
1	0	-1



Sobel filter

convolution



1	0	-1
0	0	-1
1	0	-1



Scharr filter

$$= \begin{bmatrix} 45^\circ & 0^\circ & 45^\circ \\ 0^\circ & 0 & 0 \\ -45^\circ & 0^\circ & -45^\circ \end{bmatrix}$$



3	0	-3
0	0	-10
3	0	-2



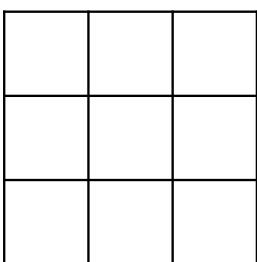
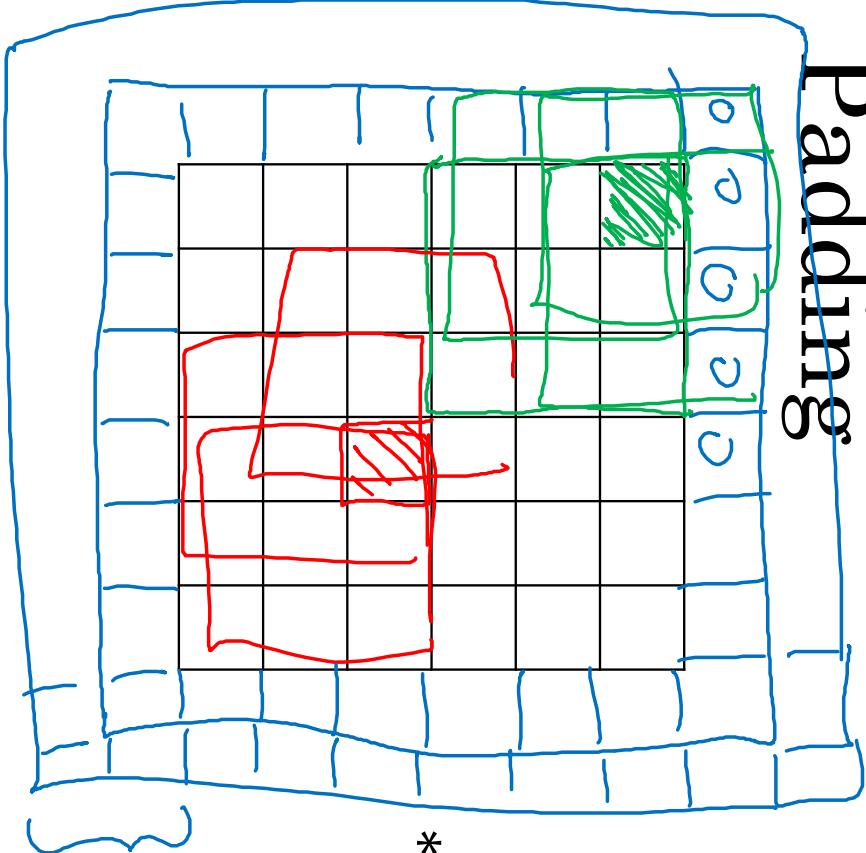
deeplearning.ai

Convolutional Neural Networks

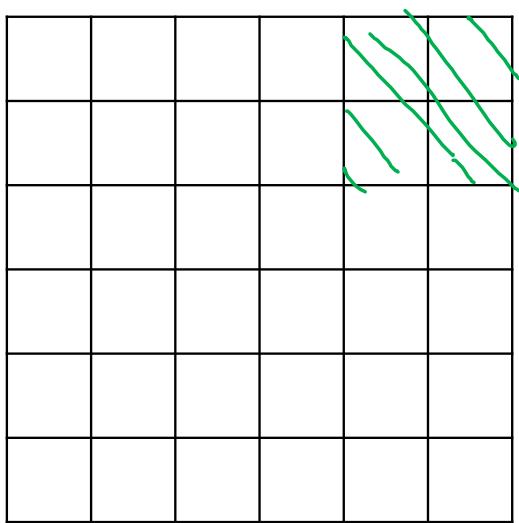
Padding

Padding

- Shrink output
- Throw away info from edge



=



6x6

4x4

$$6 \times 6 \rightarrow 8 \times 8$$

n x n

$$n - f + 1 \times n - f + 1$$

$$6 - 3 + 1 = 4$$

$$\begin{aligned} & n + 2p - f + 1 \times n + 2p - f + 1 \\ & 6 + 2 - 3 + 1 \times \underline{\quad} = 6 \times 6 \end{aligned}$$

$$p = \text{padding} = 1$$

Andrew Ng

Valid and Same convolutions

→ no padding

“Valid”: $n \times n$ $*$ $f \times f$ \rightarrow $\frac{n-f+1}{f} \times n-f+1$
 6×6 $*$ 3×3 \rightarrow 2×2

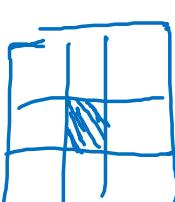
“Same”: Pad so that output size is the same as the input size.

$$n + 2p - f + 1 = n \Rightarrow p = \frac{f-1}{2}$$

$n + 2p - f + 1 \times n + 2p - f + 1$

$3 \times 3 \quad p = \frac{3-1}{2} = 1 \quad | \quad 5 \times 5 \quad p=2$

f is usually odd
3x1
5x3
5x5
7x7





deeplearning.ai

Convolutional Neural Networks

Strided convolutions

Strided convolution

2	3	3	4	7	3	4	4	6	3	2	4	9	4
6	1	6	0	9	1	8	0	7	1	4	0	3	2
3	3	4	4	8	3	3	4	8	3	9	4	7	4
7	1	8	0	3	1	6	0	6	1	3	0	4	2
4	-3	2	4	1	-3	8	4	3	3	4	4	6	4
3	1	2	0	4	1	1	0	9	1	8	0	3	2
0	-1	1	0	3	-1	9	0	2	-1	1	0	4	3

*

3	4	4
1	0	2
-1	0	3

*

Stride = 2

3x3

$$\lfloor \frac{2}{2} \rfloor = \text{floor}(\frac{2}{2})$$

91	100	83
69	41	127
44	72	74



7x7

$n \times n$ * $f \times f$

paddig p stride s

$$S=2$$

$$\frac{n+2p-f}{s} + 1 = \frac{4+0-3}{2} + 1 = \frac{4}{2} + 1 = 3$$

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

Summary of convolutions

$n \times n$ image $f \times f$ filter

padding p stride s

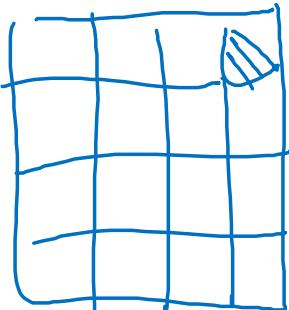
Output size:

$$\left\lceil \frac{n+2p-f}{s} + 1 \right\rceil \quad \times \quad \left\lceil \frac{n+2p-f}{s} + 1 \right\rceil$$

Technical note on cross-correlation vs. convolution

Convolution in math textbook:

$$\begin{array}{c} \begin{array}{|c|c|c|c|c|c|} \hline 2 & 3 & 7 & 4 & 6 & 2 \\ \hline 6 & 6 & 9 & 8 & 7 & 4 \\ \hline -1 & 4 & 8 & 3 & 8 & 9 \\ \hline 7 & 8 & 3 & 6 & 6 & 3 \\ \hline 4 & 2 & 1 & 8 & 3 & 4 \\ \hline 3 & 2 & 4 & 1 & 9 & 8 \\ \hline \end{array} \quad * \quad \begin{array}{|c|c|c|} \hline 3 & 4 & 5 \\ \hline 1 & 0 & 2 \\ \hline -1 & 9 & 7 \\ \hline \end{array} \\ = \end{array}$$

$$(A * B) * C = A * (B * C)$$


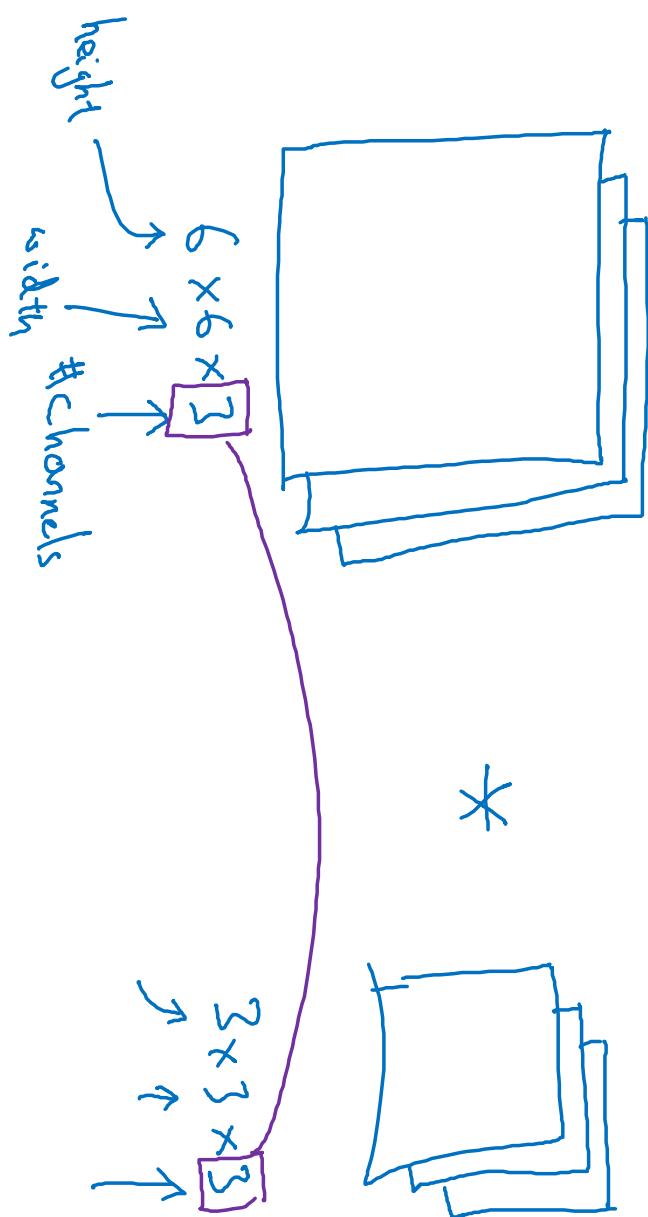


Convolutional
Neural Networks

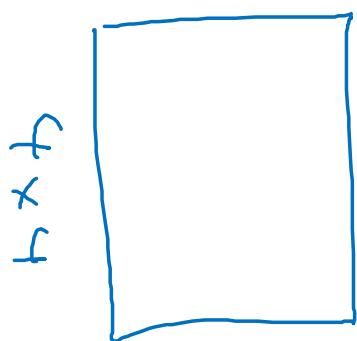
Convolutions over
volumes

deeplearning.ai

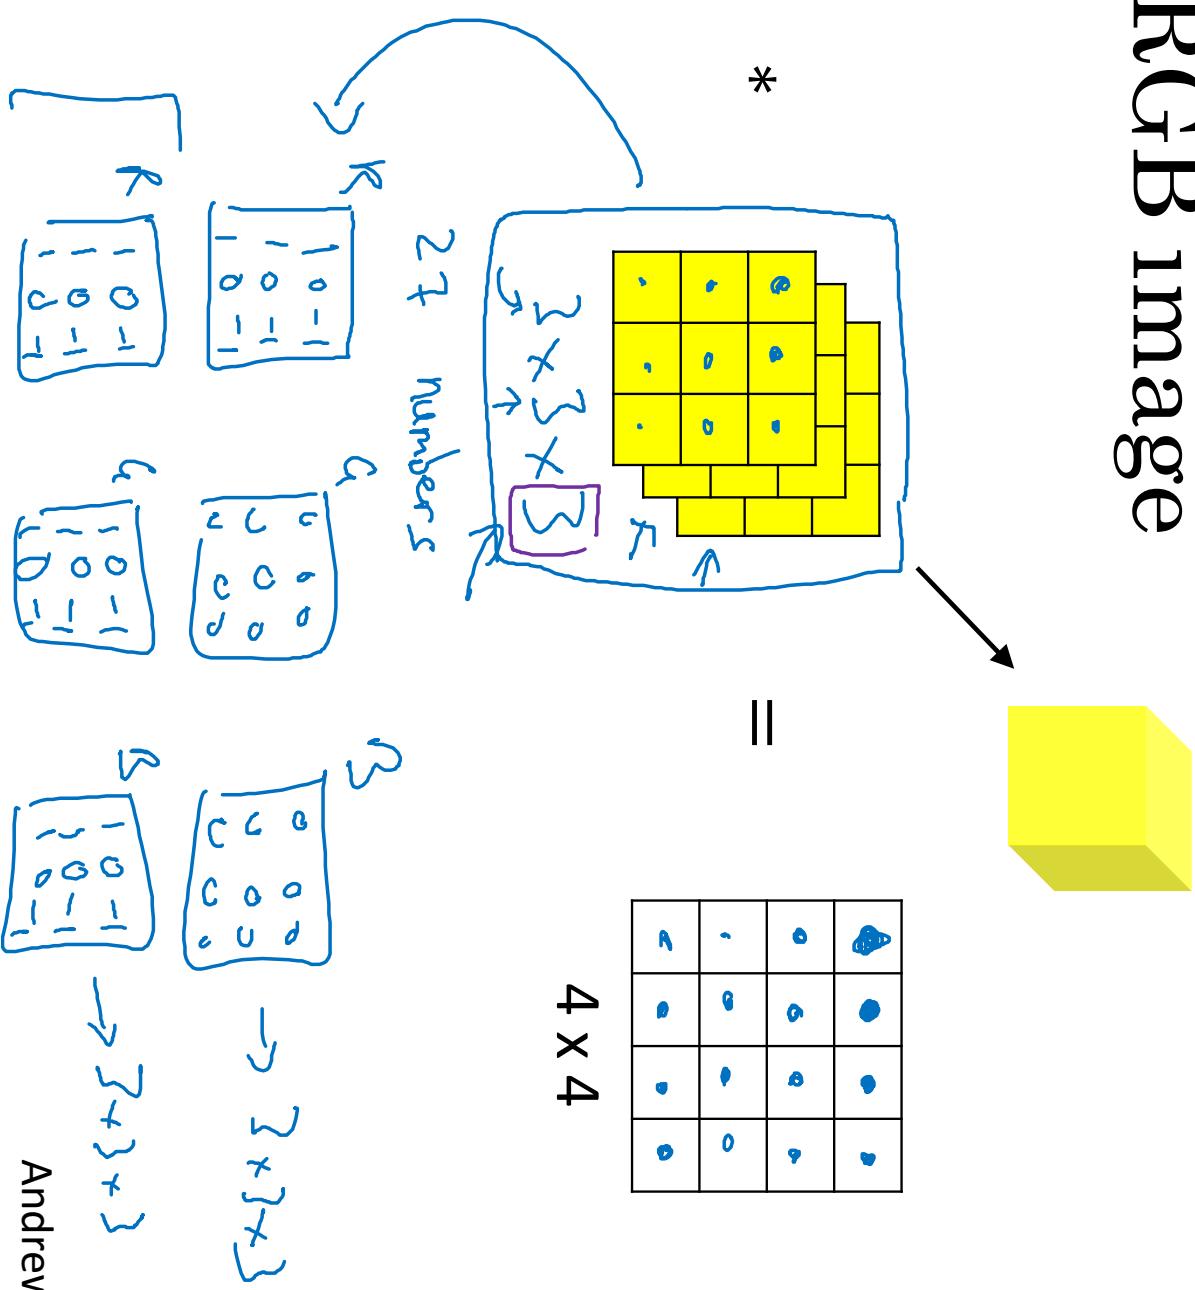
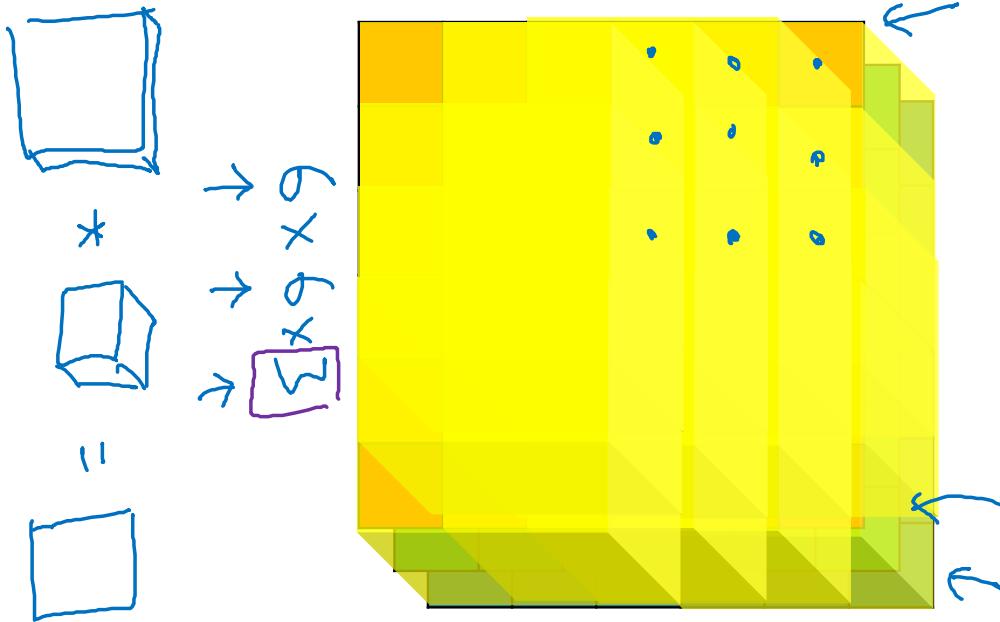
Convolutions on RGB images



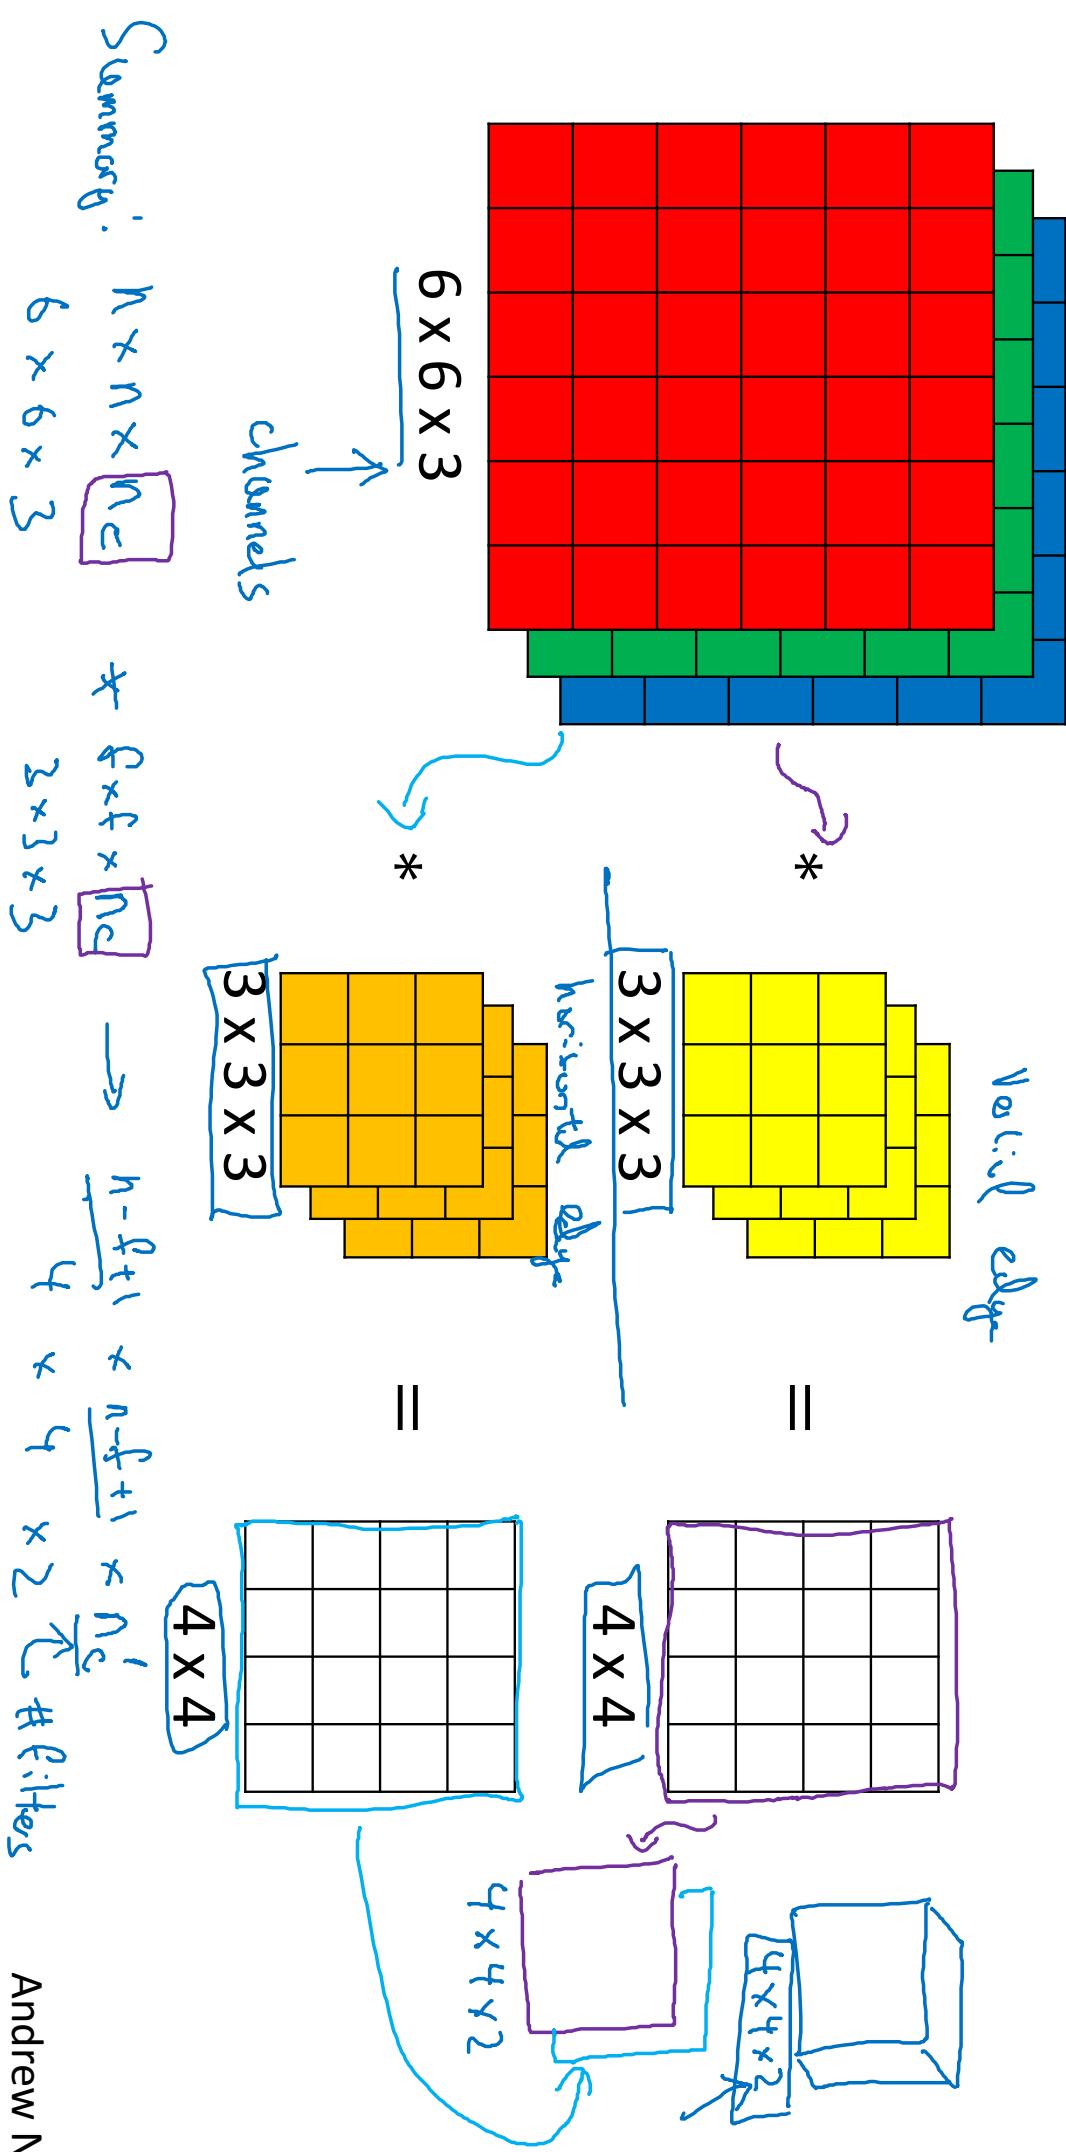
=



Convolutions on RGB image



Multiple filters

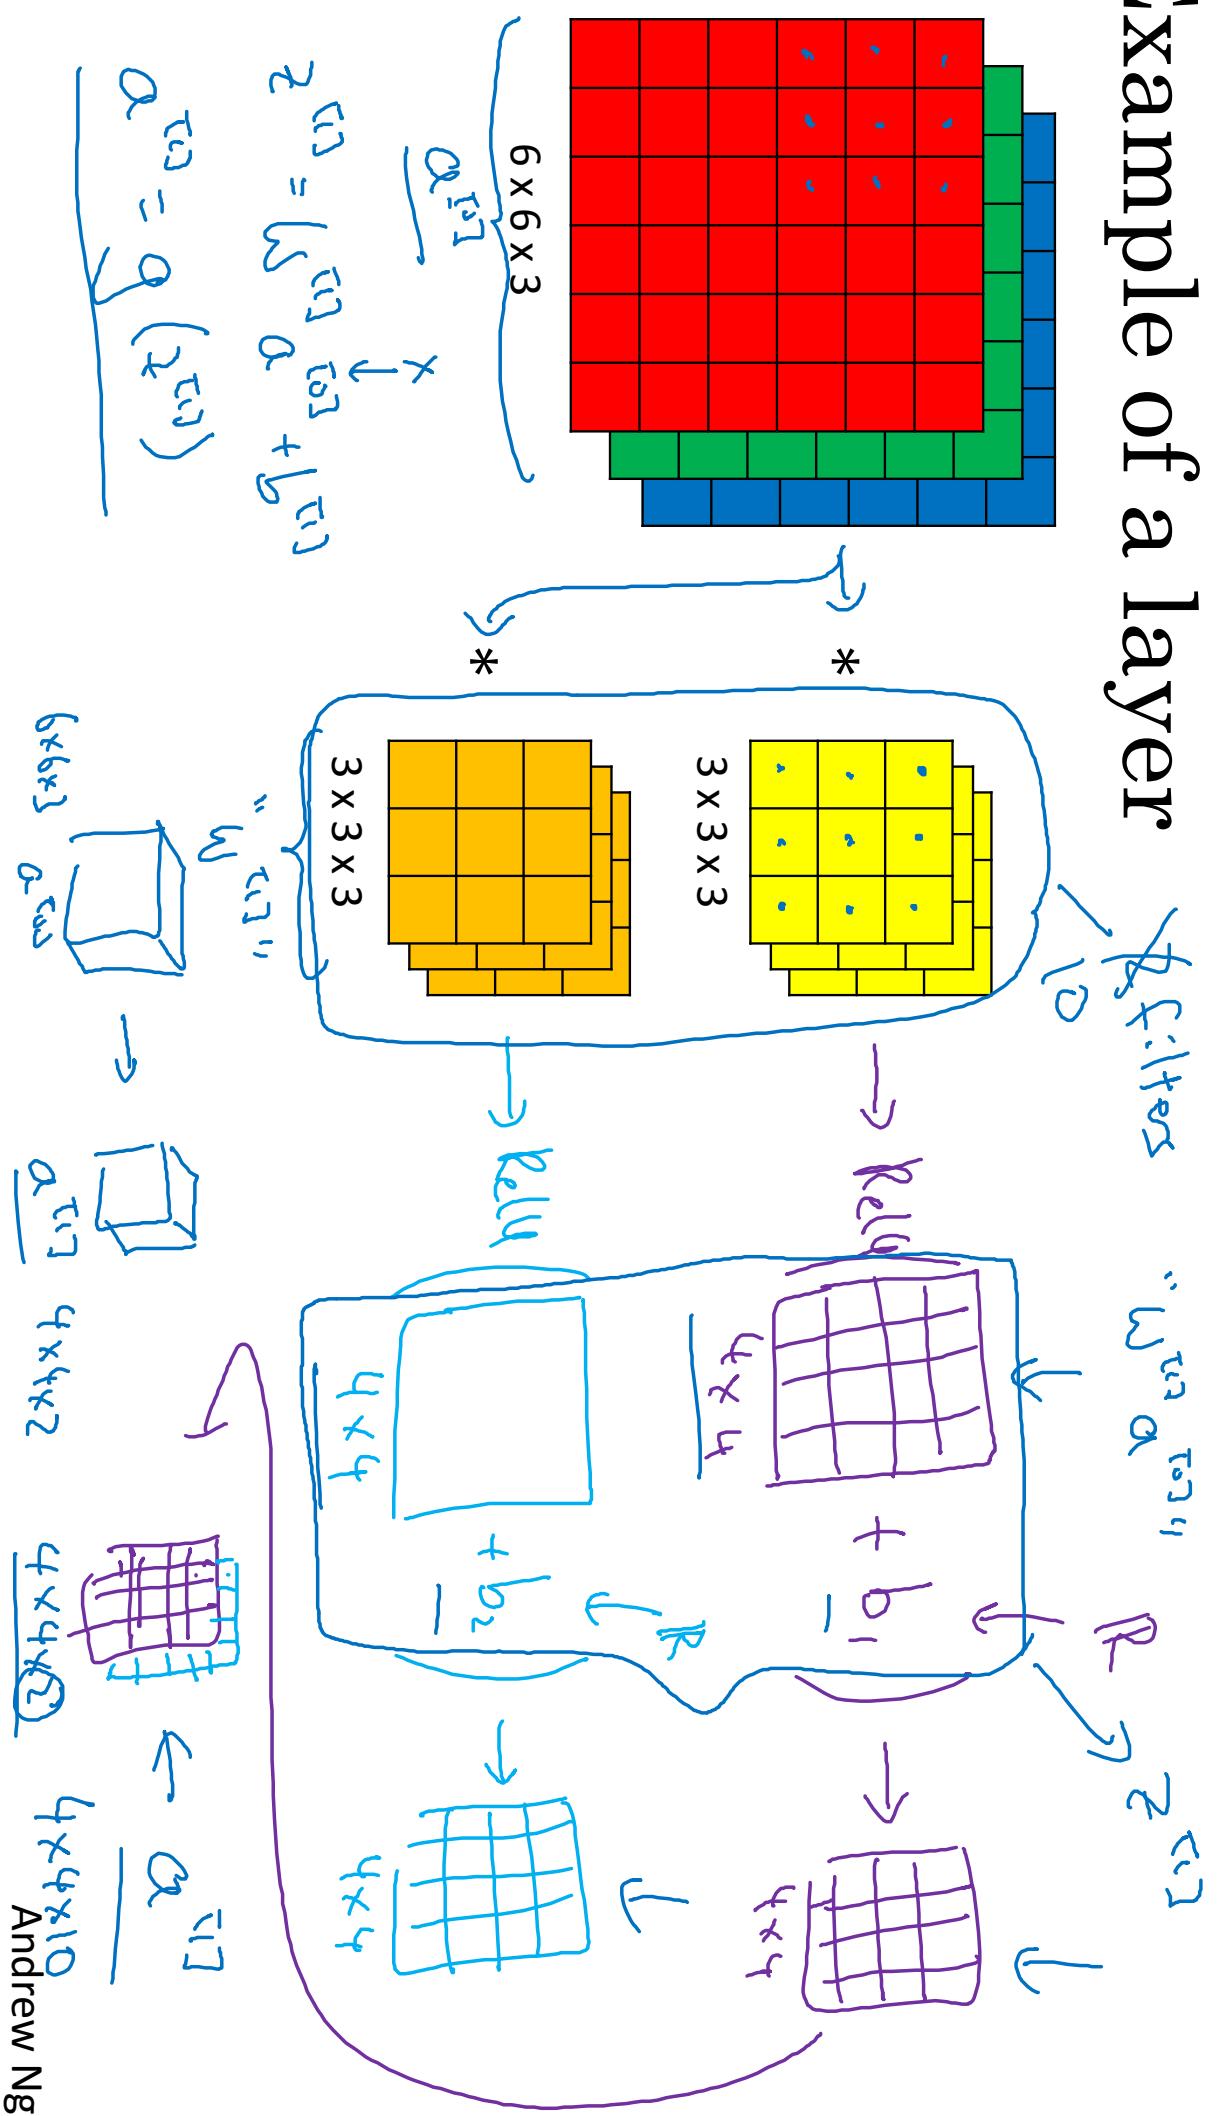




Convolutional Neural Networks

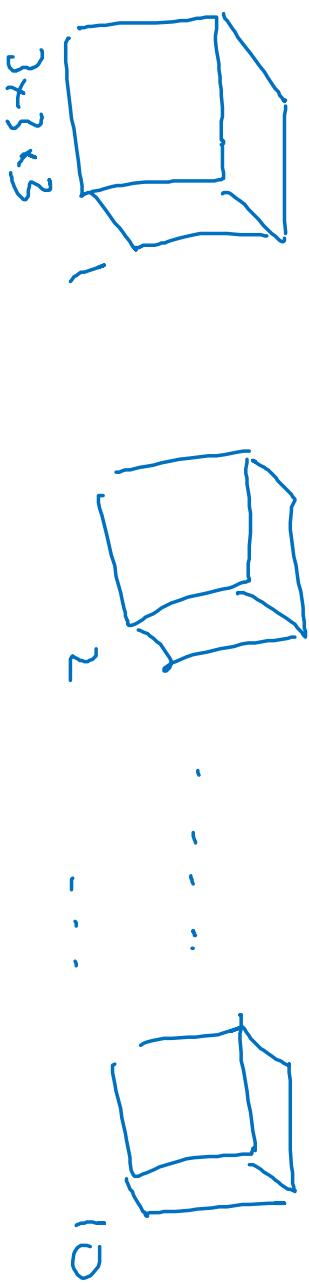
One layer of a
convolutional
network

Example of a layer



Number of parameters in one layer

If you have 10 filters that are $3 \times 3 \times 3$ in one layer of a neural network, how many parameters does that layer have?



~ 7 parameters.
+ bias
 $\rightarrow \sim 28$ parameters.
 ~ 280 parameters.

Summary of notation

If layer l is a convolution layer:

$f^{[l]}$ = filter size

$p^{[l]}$ = padding

$s^{[l]}$ = stride

$n_c^{[l]}$ = number of filters

Each filter is: $\underbrace{f^{[l] \times f^{[l]} \times n_c^{[l]}}}_{n_{kw}}$

Activations: $a^{[l]} \rightarrow n_H \times n_W \times n_c^{[l]}$.

Weights: $f^{[l+1] \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]}}$

bias: $n_c^{[l]} - (1, 1, 1, n_c^{[l]})$ #filter in layer l . $\underbrace{n_c^{[l]} \times n_H \times n_W}$

$$\text{Input: } \underbrace{n_H \times n_W \times n_c^{[l-1]}}_{\text{Input}} \leftarrow$$

$$\text{Output: } \underbrace{n_H \frac{n_H}{s^{[l]}} \times n_W \frac{n_W}{s^{[l]}} \times n_c^{[l]}}_{\text{Output}} \leftarrow$$

$$n_H \frac{n_H}{s^{[l]}} = \left\lfloor \frac{n_H + 2p^{[l]} - f^{[l]}}{s^{[l]}} \right\rfloor + 1$$

$$n_W \frac{n_W}{s^{[l]}} = \left\lfloor \frac{n_W + 2p^{[l]} - f^{[l]}}{s^{[l]}} \right\rfloor + 1$$

$$n_c^{[l]} = \left\lfloor \frac{n_c^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} \right\rfloor + 1$$

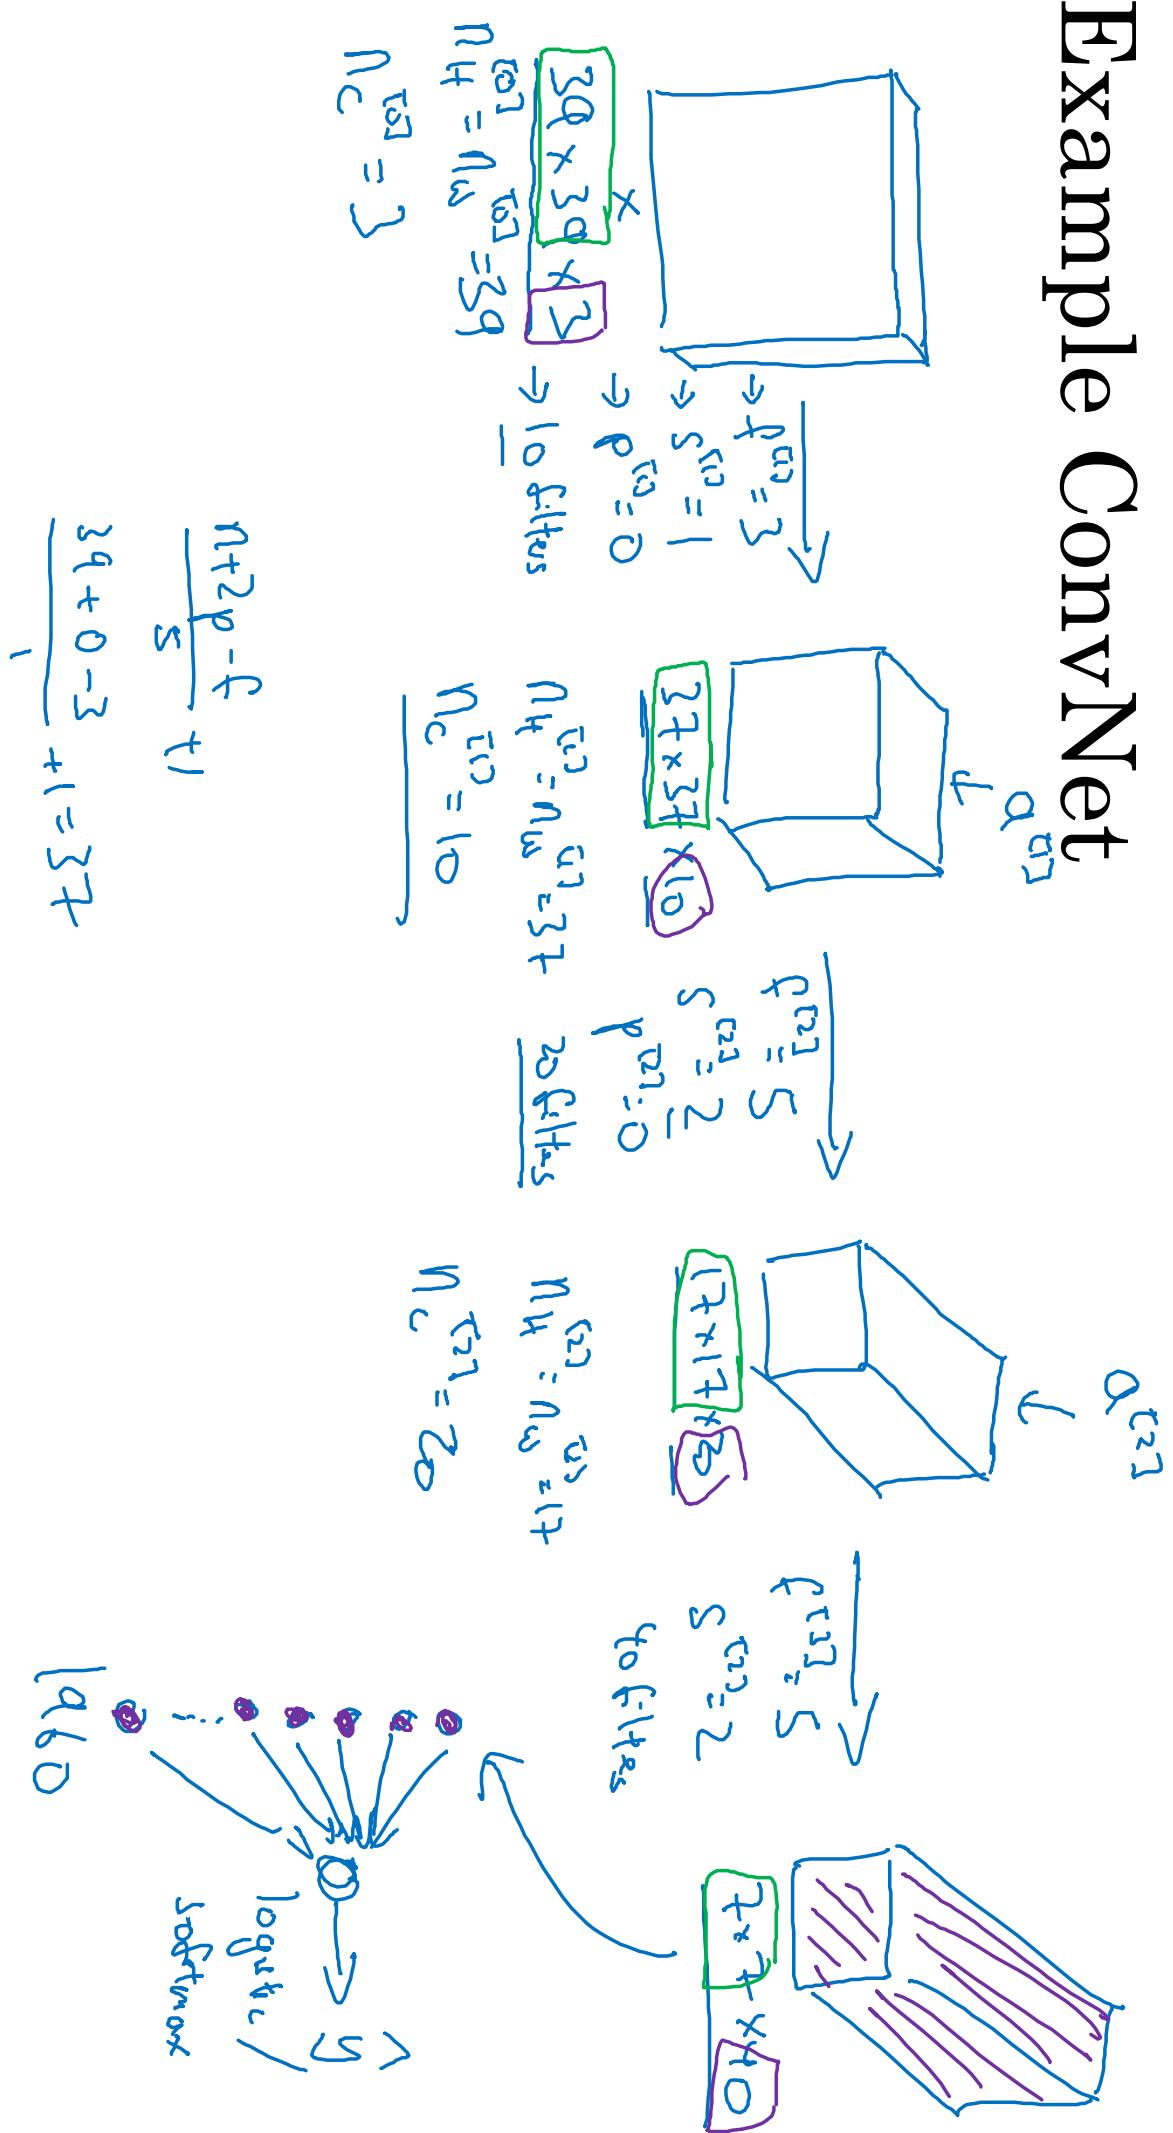


Convolutional Neural Networks

A simple convolution
network example

deeplearning.ai

Example ConvNet



Types of layer in a convolutional network:

- Convolution (conv) ←
 - Pooling (pool) ←
 - Fully connected (Fc) ←
- 



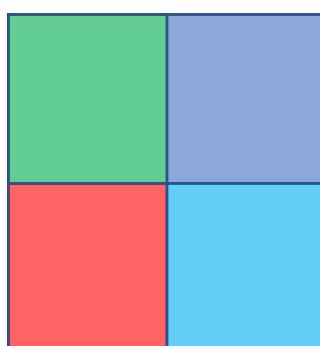
deeplearning.ai

Convolutional Neural Networks

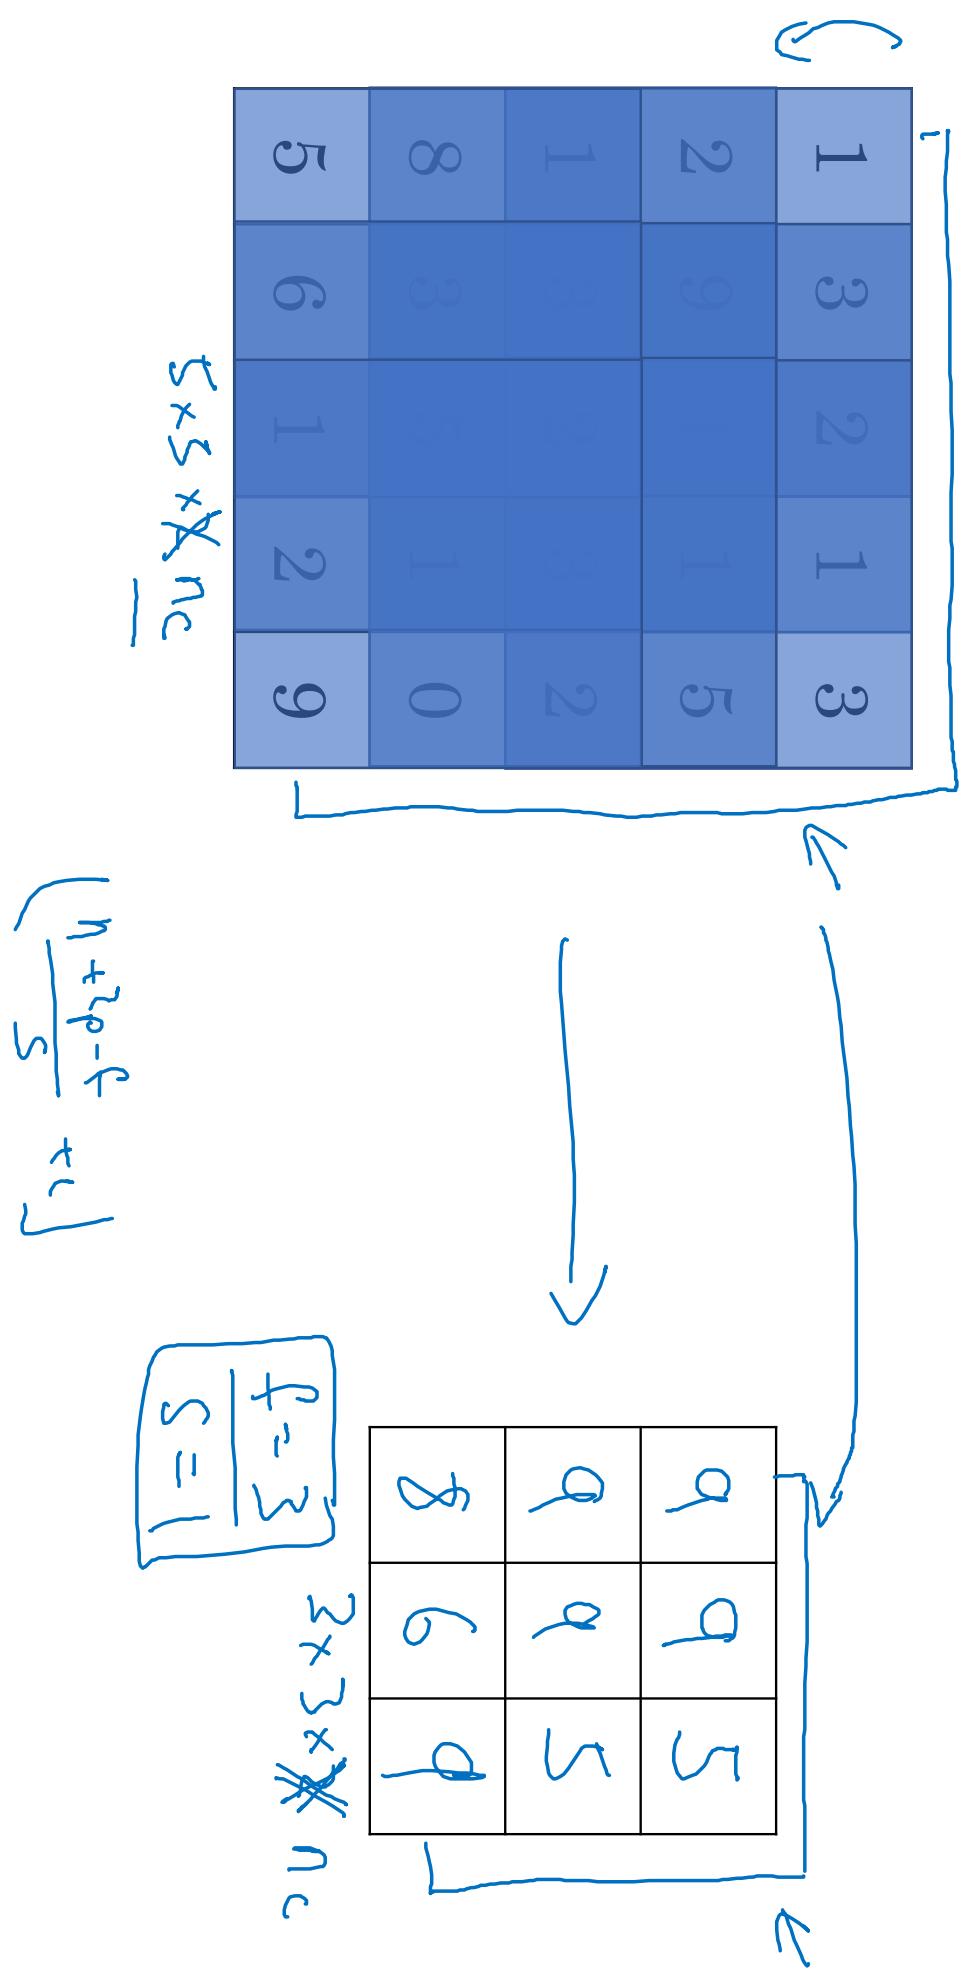
Pooling layers

Pooling layer: Max pooling

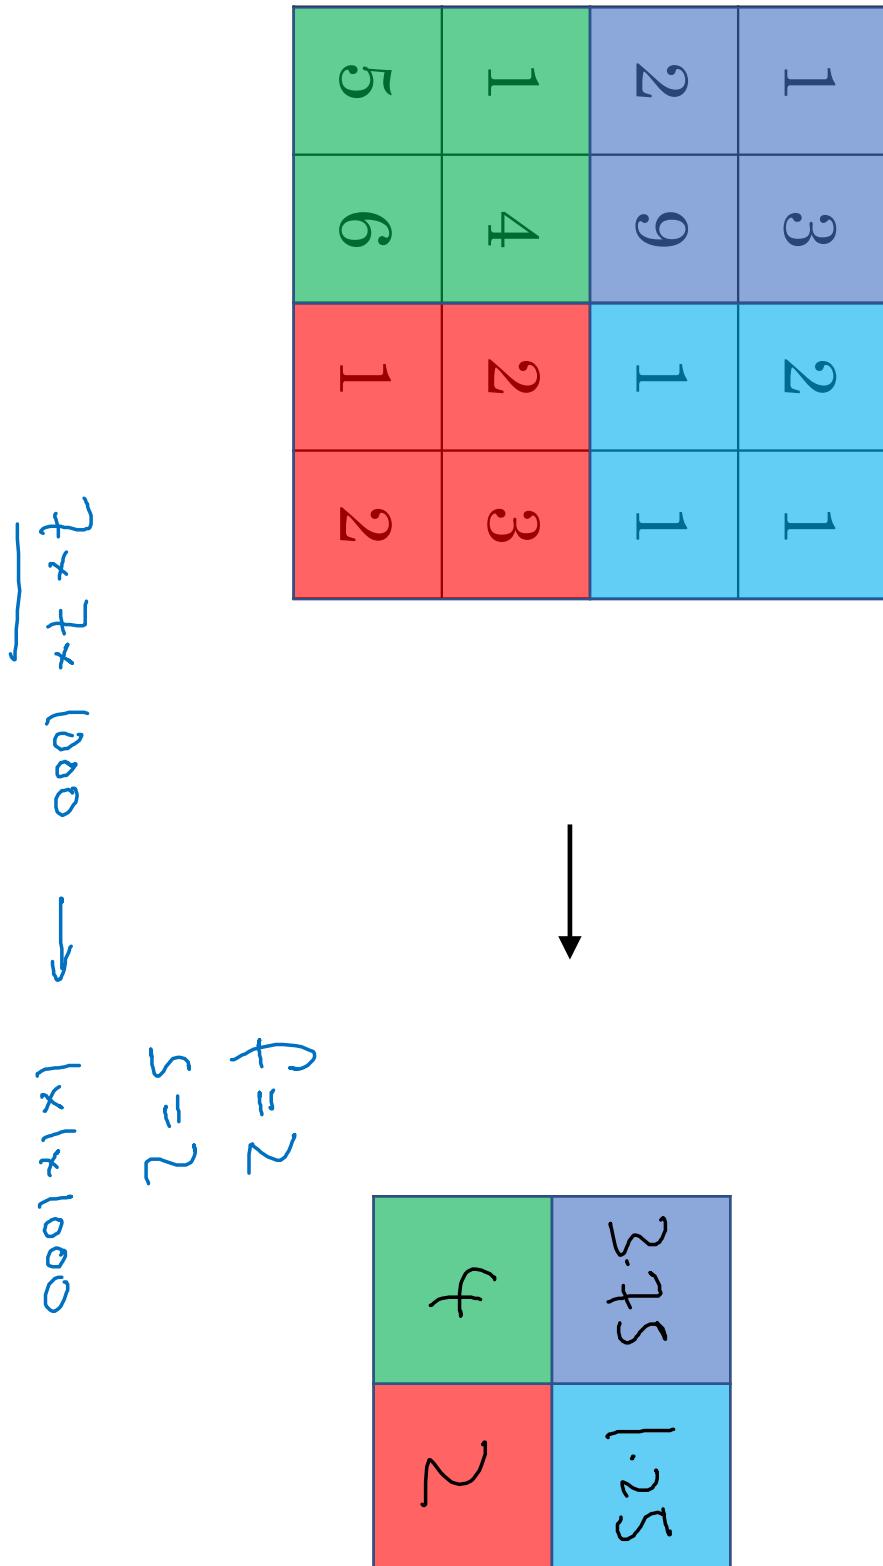
1	3	2	1
2	9	1	1
1	3	2	1
5	6	1	2



Pooling layer: Max pooling



Pooling layer: Average pooling



Summary of pooling

Hyperparameters:

f : filter size

$$f=2, s=2$$

s : stride

$$f=3, s=2$$

Max or average pooling

$$\frac{n_H-f+1}{s} \times \frac{n_W-f+1}{s} \times n_C$$

~~p : padding~~

No parameters to learn!

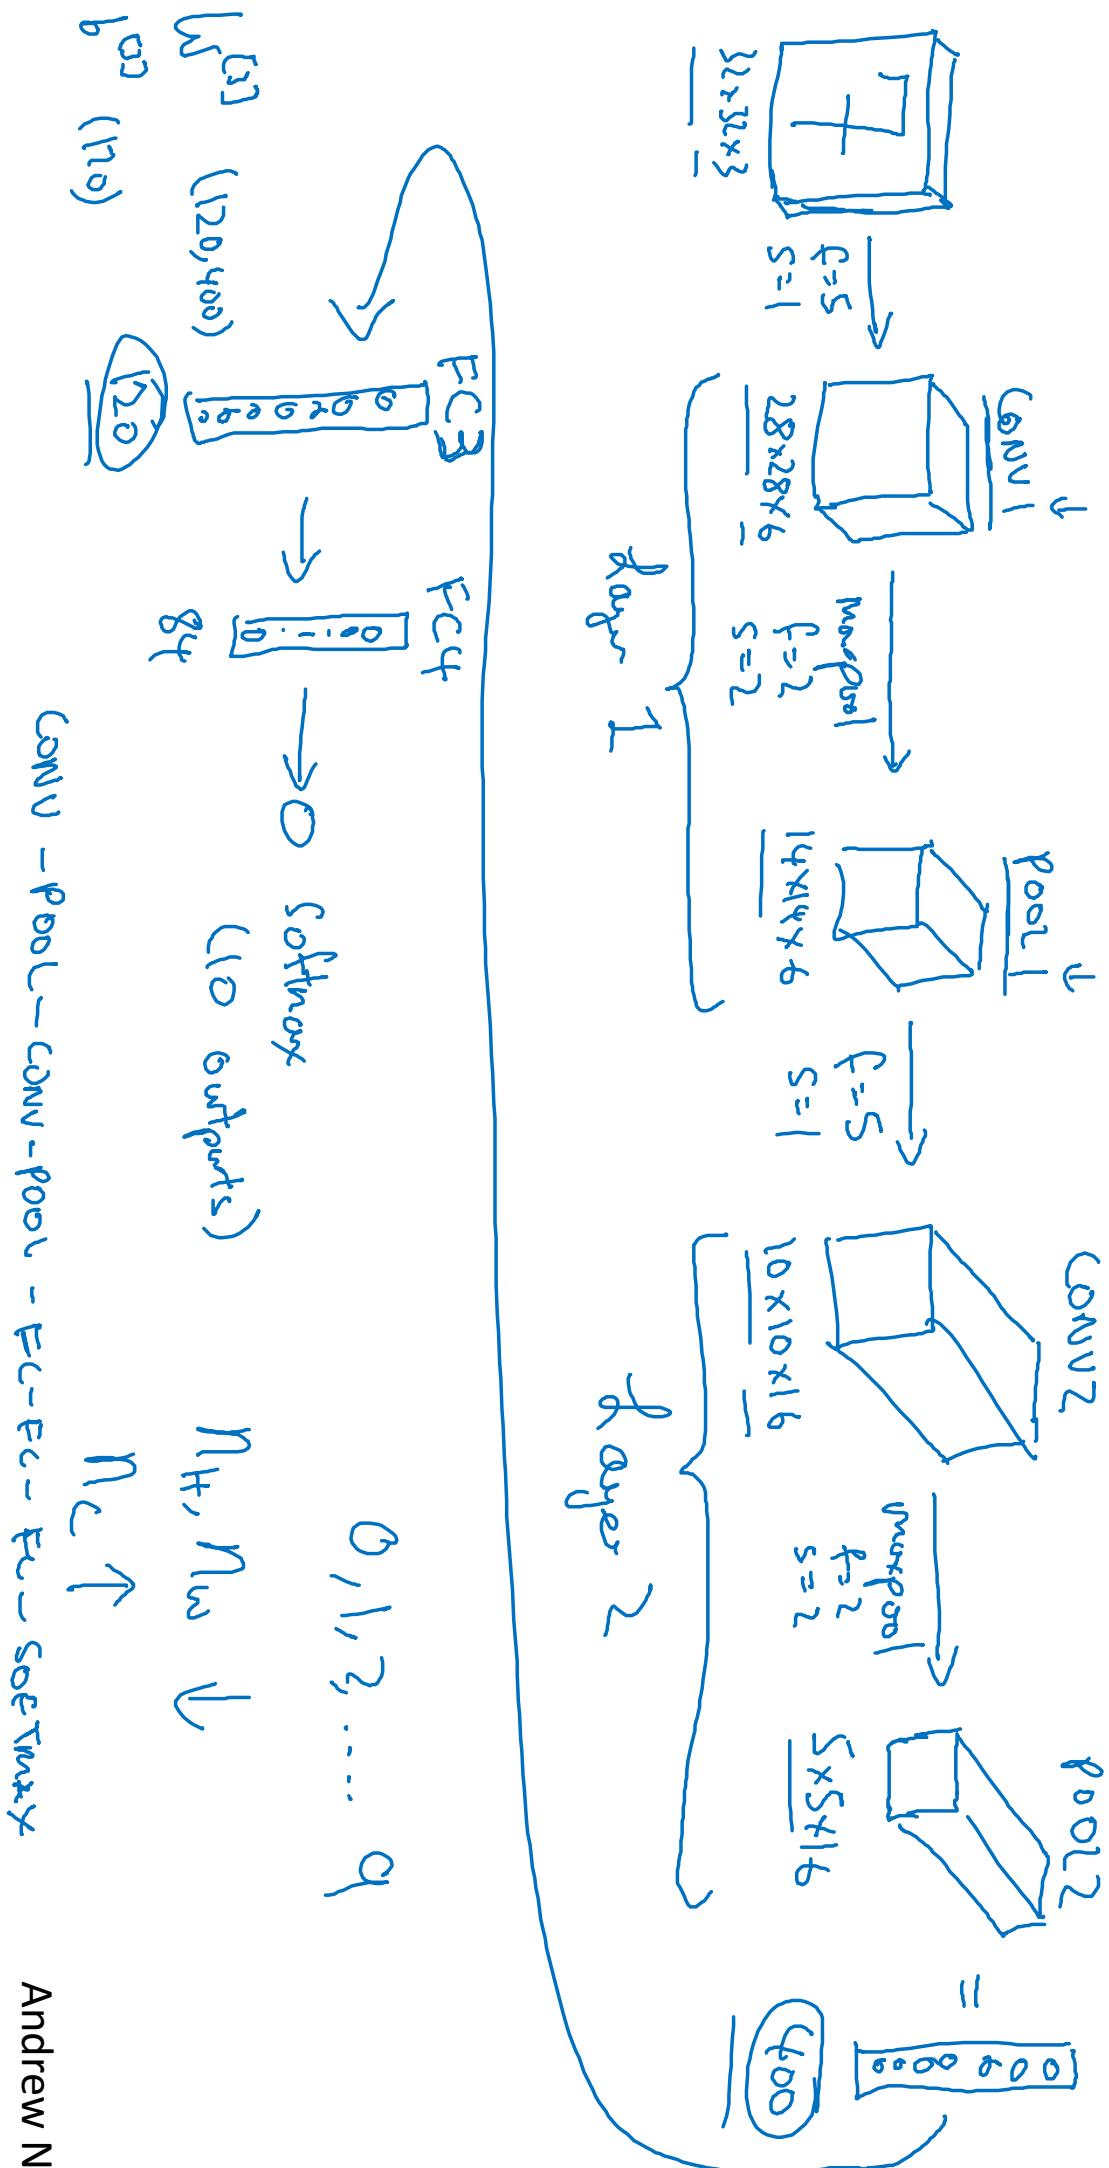


Convolutional Neural Networks

Convolutional neural
network example

deeplearning.ai

Neural network example (LeNet-5)



Neural network example

	Activation shape	Activation Size	# parameters
Input:	(32,32,3)	3,072	α_{w1} 0
CONV1 (f=5, s=1)	(28,28,8)	6,272	608 \leftarrow
POOL1	(14,14,8)	1,568	0 \leftarrow
CONV2 (f=5, s=1)	(10,10,16)	1,600	3216 \leftarrow
POOL2	(5,5,16)	400	0 \leftarrow
FC3	(120,1)	120	48120 {
FC4	(84,1)	84	10164 }
Softmax	(10,1)	10	850

Andrew Ng

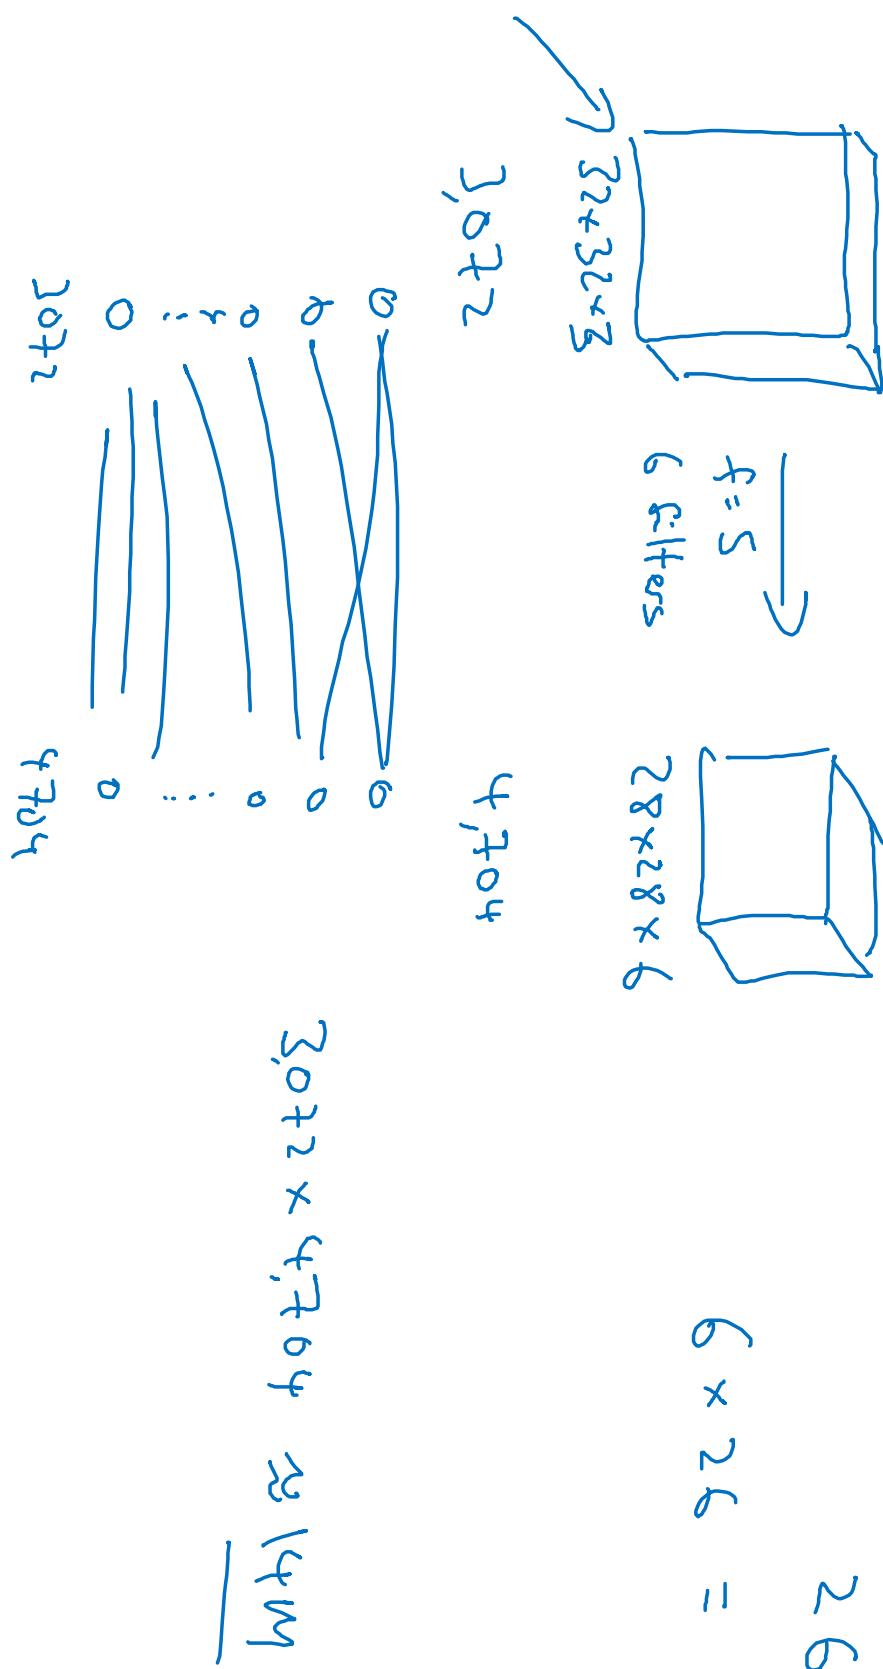


deeplearning.ai

Convolutional Neural Networks

Why convolutions?

Why convolutions



Why convolutions

Translation invariance

The diagram illustrates a convolution operation. It shows an input image of size 6×6 with values 10 or 0, and a kernel of size 3×3 with values 1, 0, -1. The operation is labeled $*$ (convolution) and $=$ (output). Handwritten annotations include green arrows pointing to the first row of the input and the kernel, and blue arrows pointing to the output values 30 and 0. A red circle highlights the value 30 in the output, which is circled in green.

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

1	0	-1
1	0	-1
1	0	-1

=

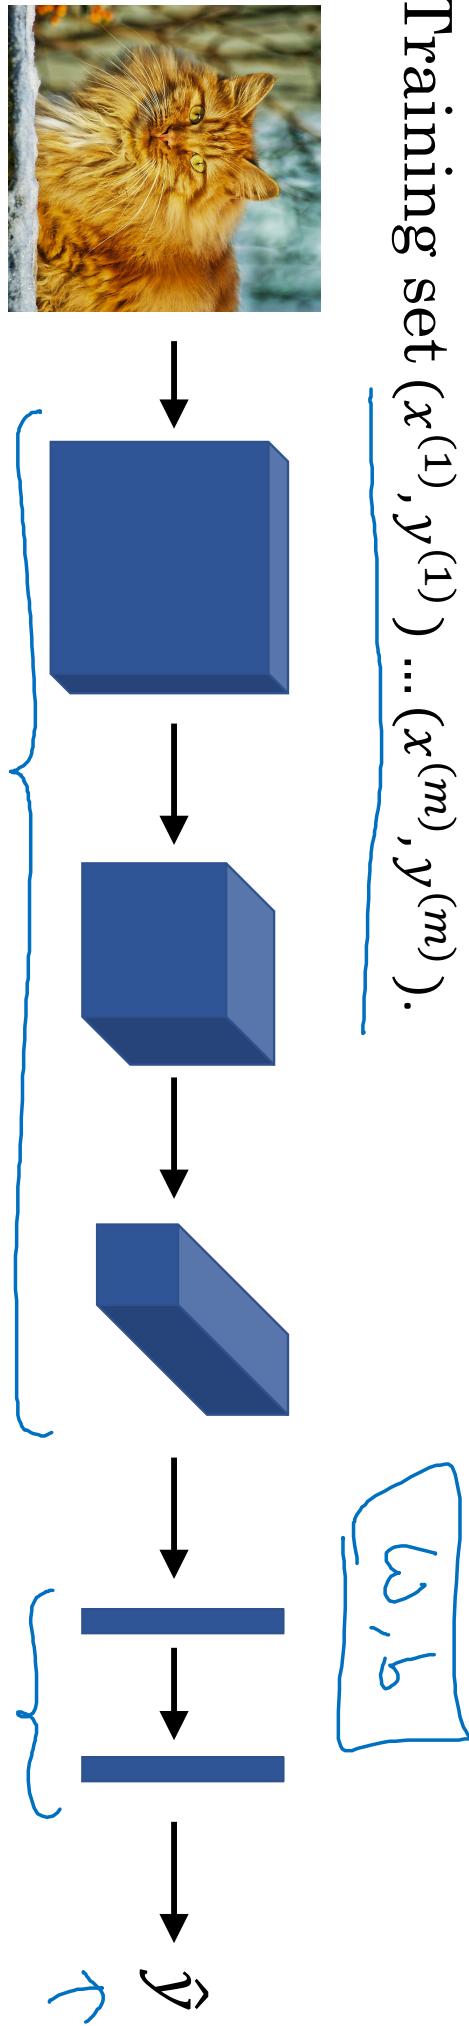
0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

Parameter sharing: A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

→ **Sparsity of connections:** In each layer, each output value depends only on a small number of inputs.

Putting it together

Training set $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$.



$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Use gradient descent to optimize parameters to reduce J