



deeplearning.ai

## Case Studies

---

Why look at  
case studies?

# Outline

Classic networks:

- LeNet-5 ←
- AlexNet ←
- VGG ←

ResNet (152)

Inception



deeplearning.ai

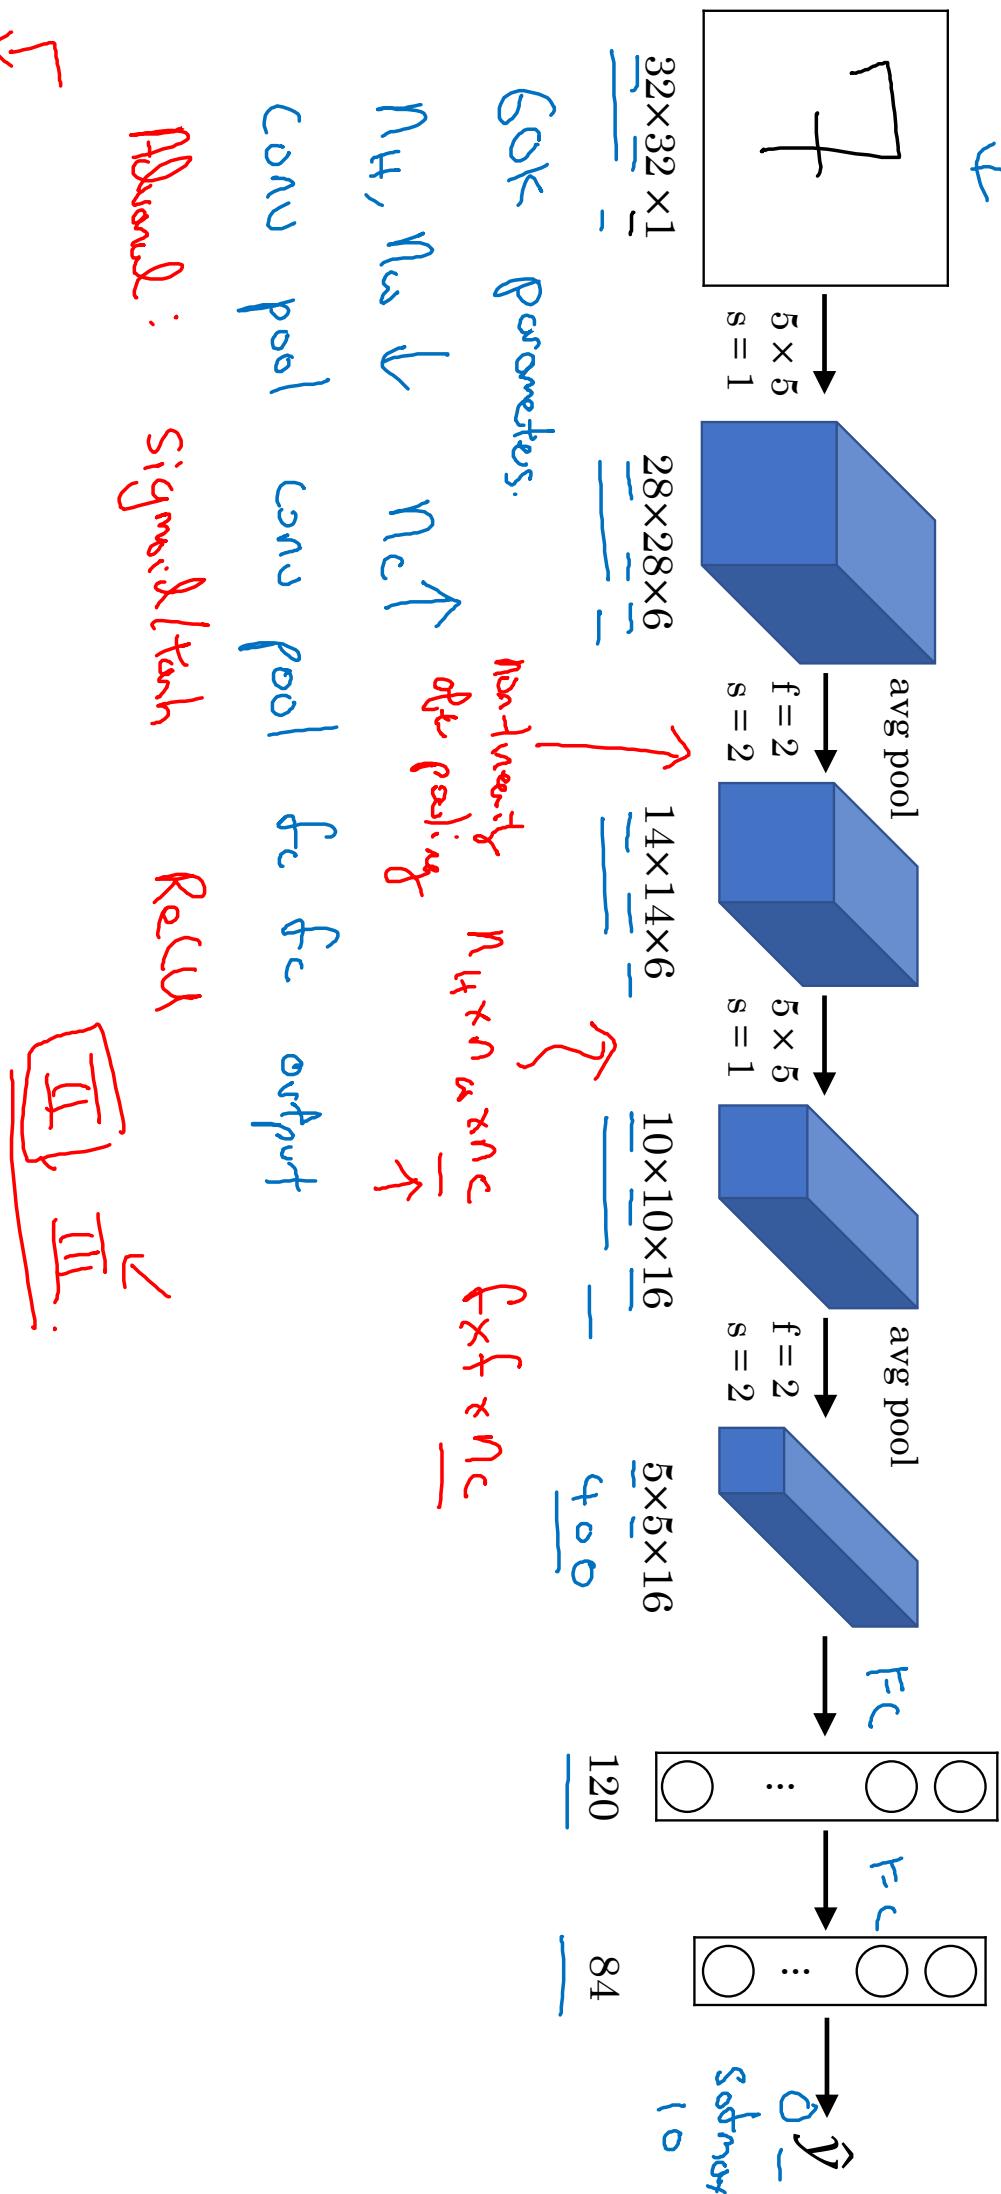
---

# Case Studies

---

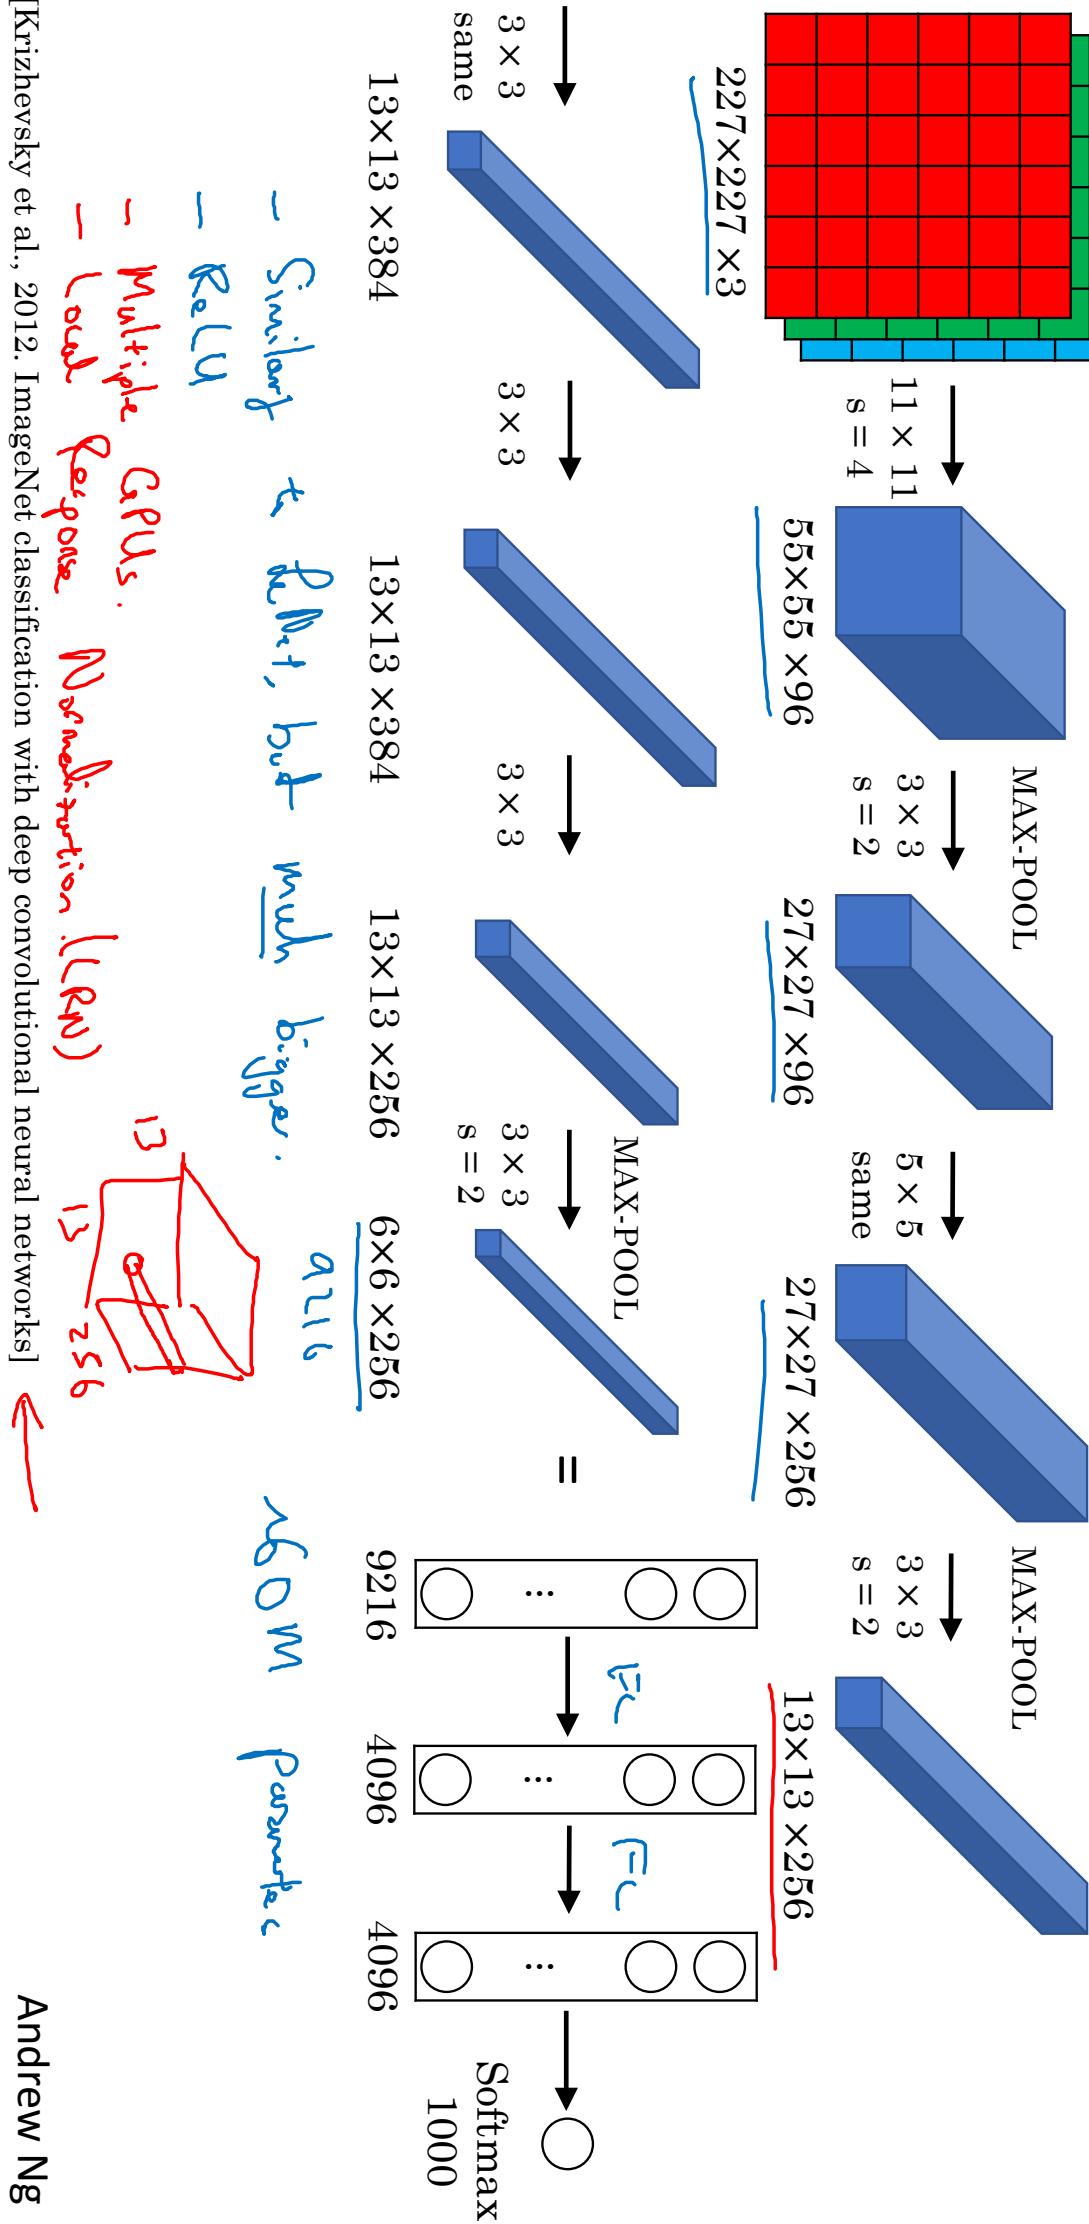
# Classic networks

# LeNet - 5



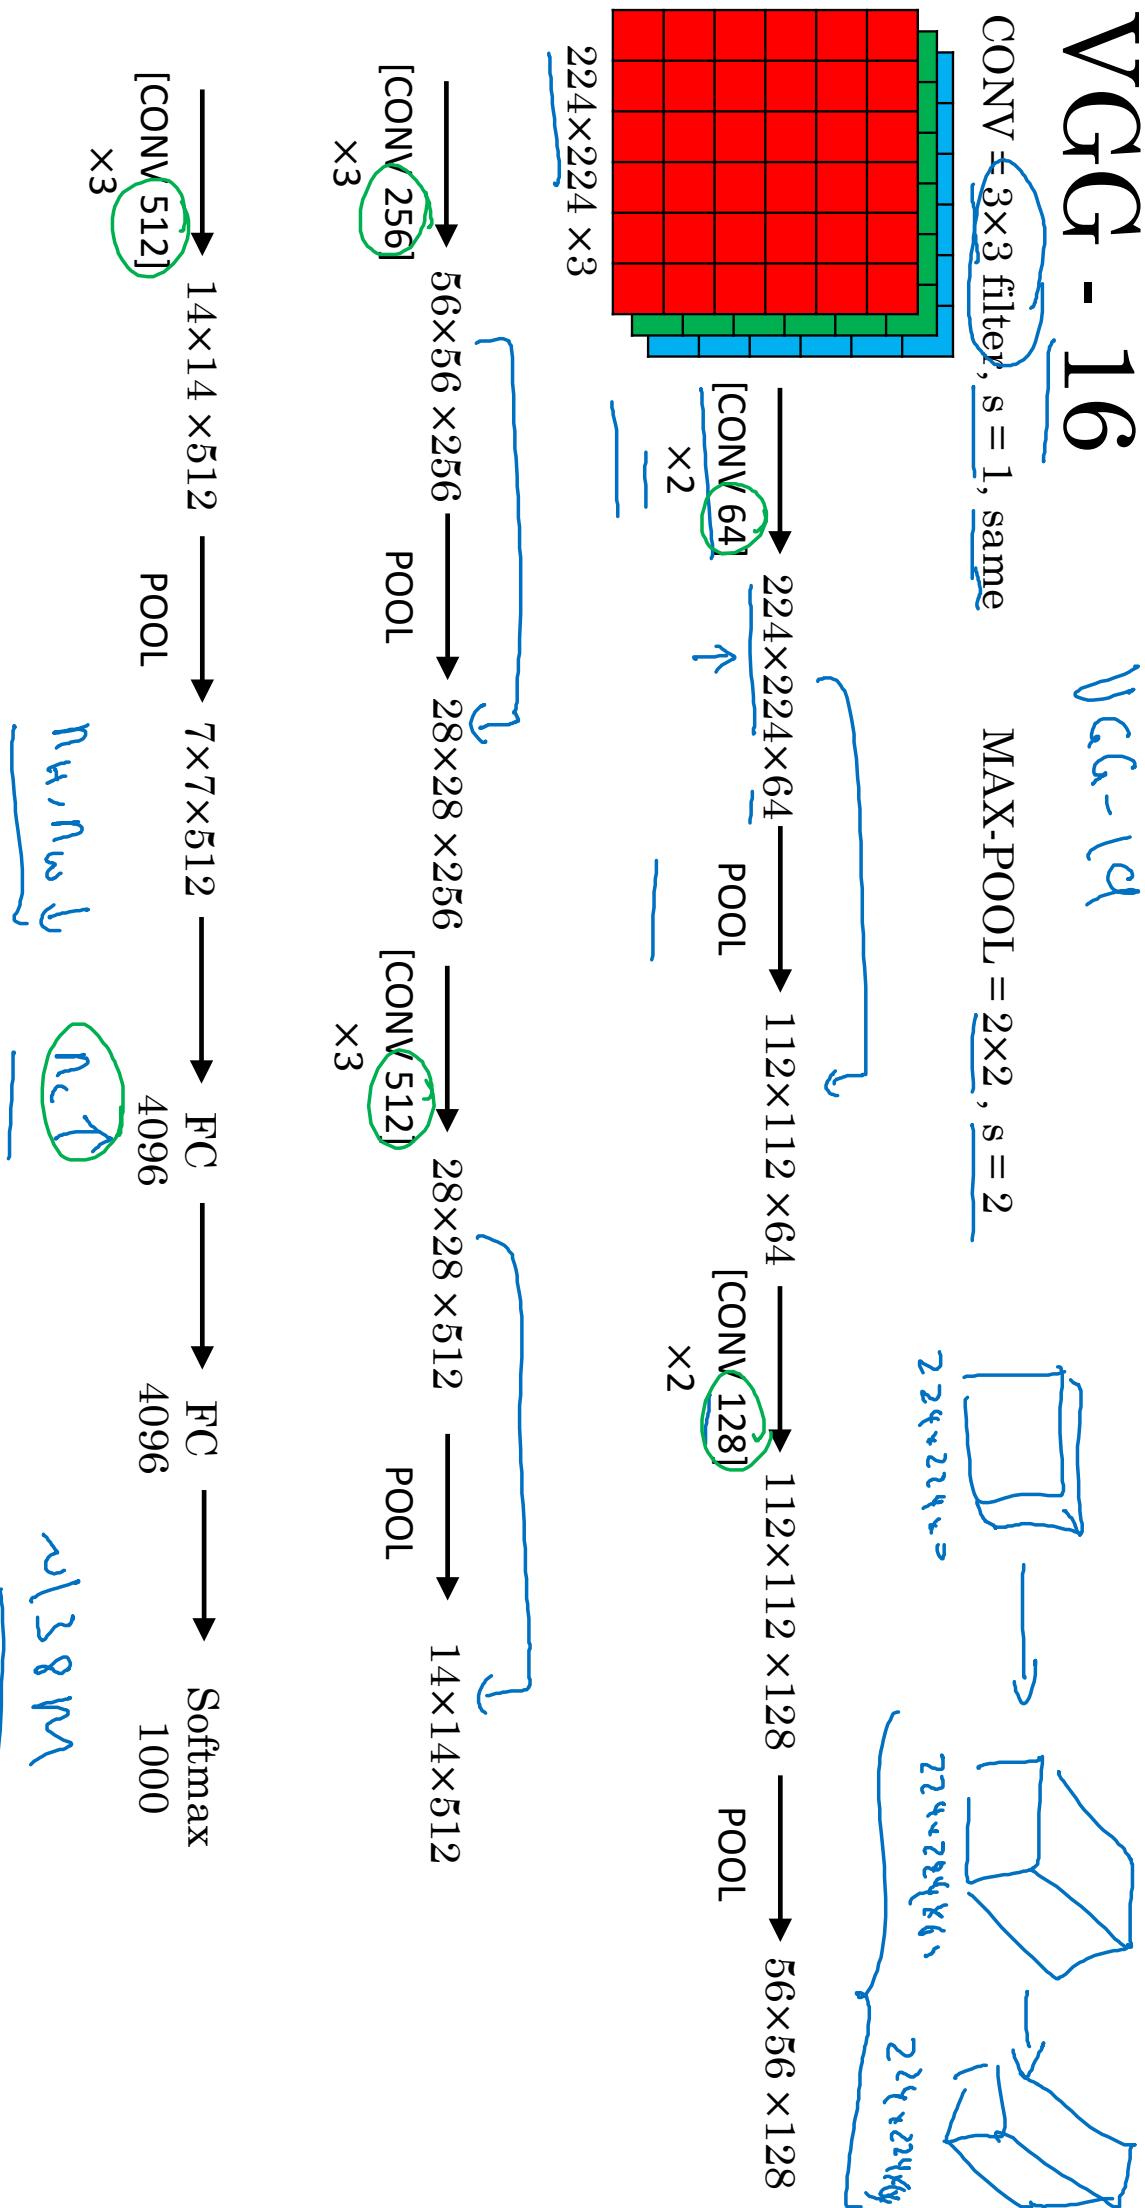
[LeCun et al., 1998. Gradient-based learning applied to document recognition]

# AlexNet



# VGG - 16

VGG-16



[Simonyan & Zisserman 2015. Very deep convolutional networks for large-scale image recognition]

Andrew Ng



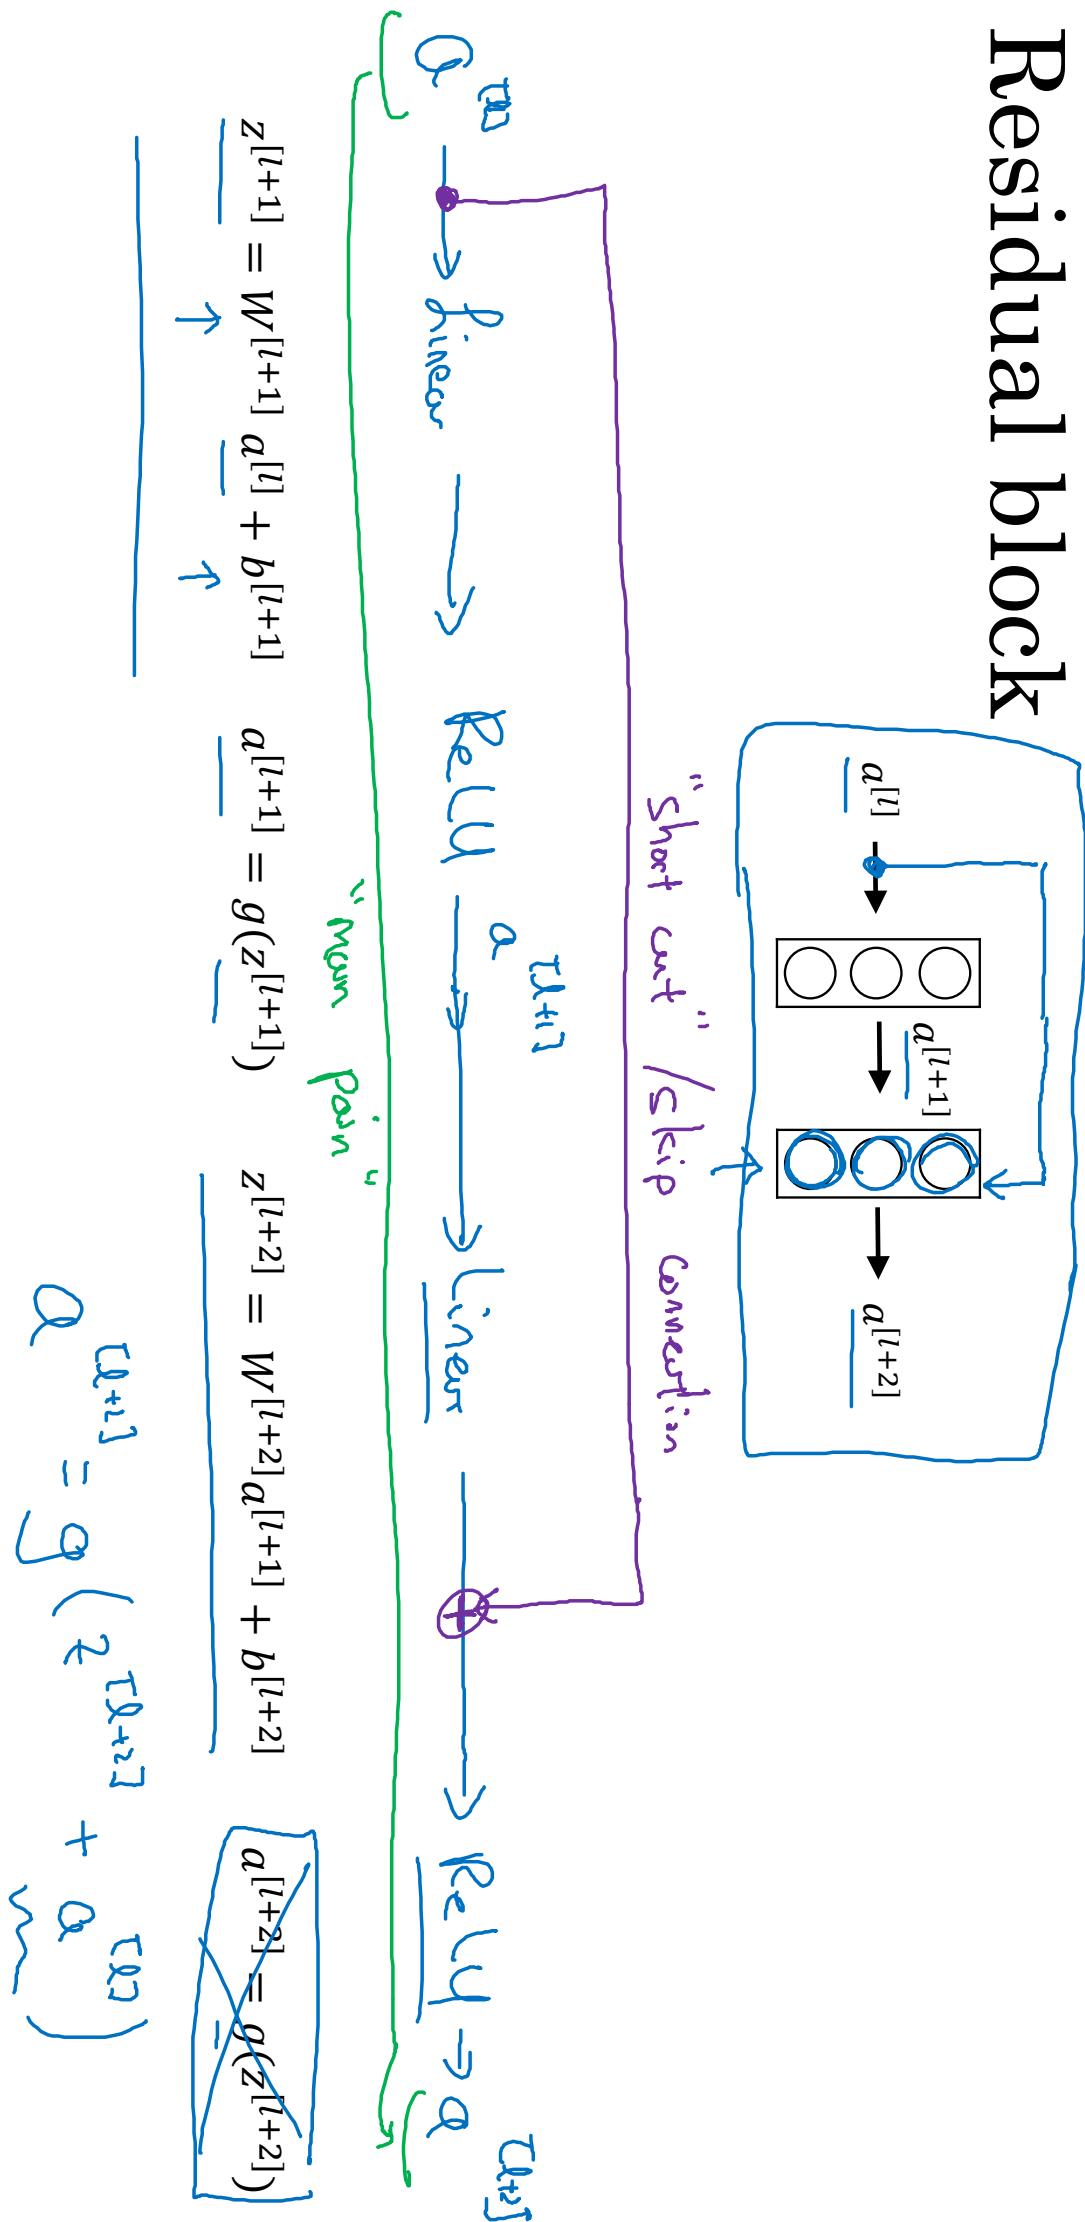
deeplearning.ai

## Case Studies

---

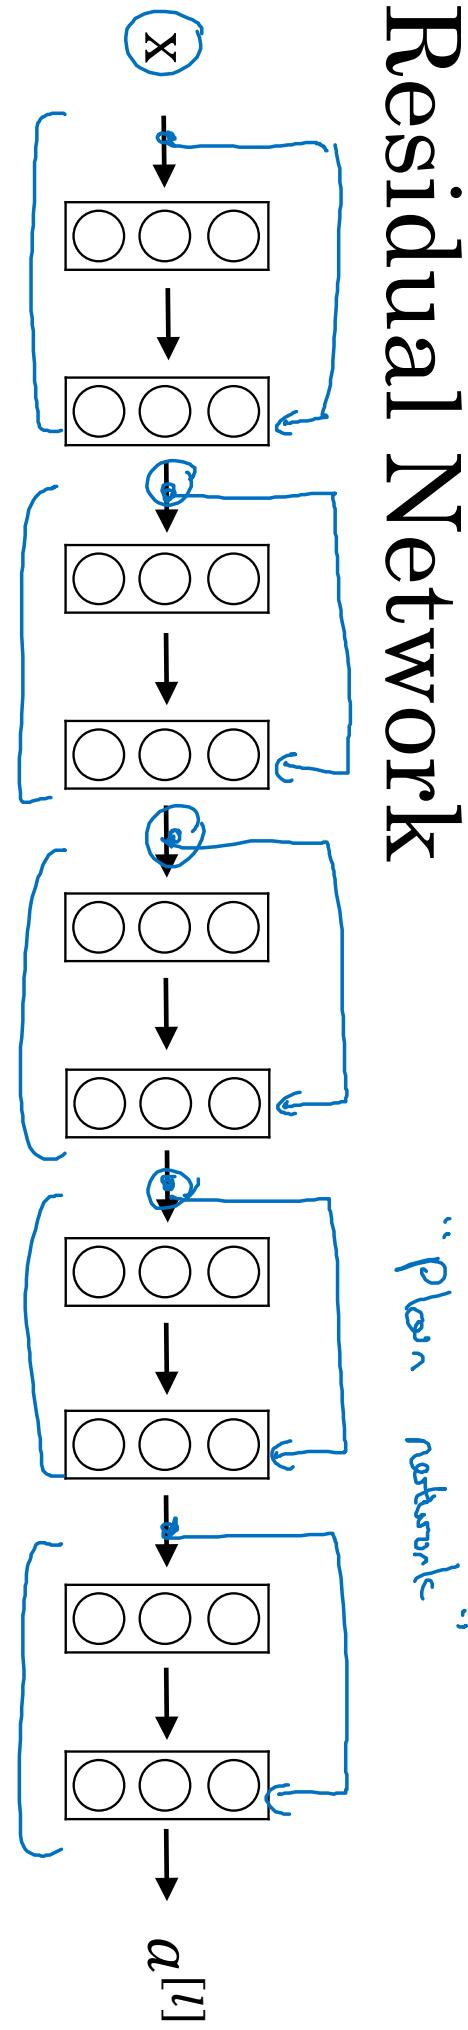
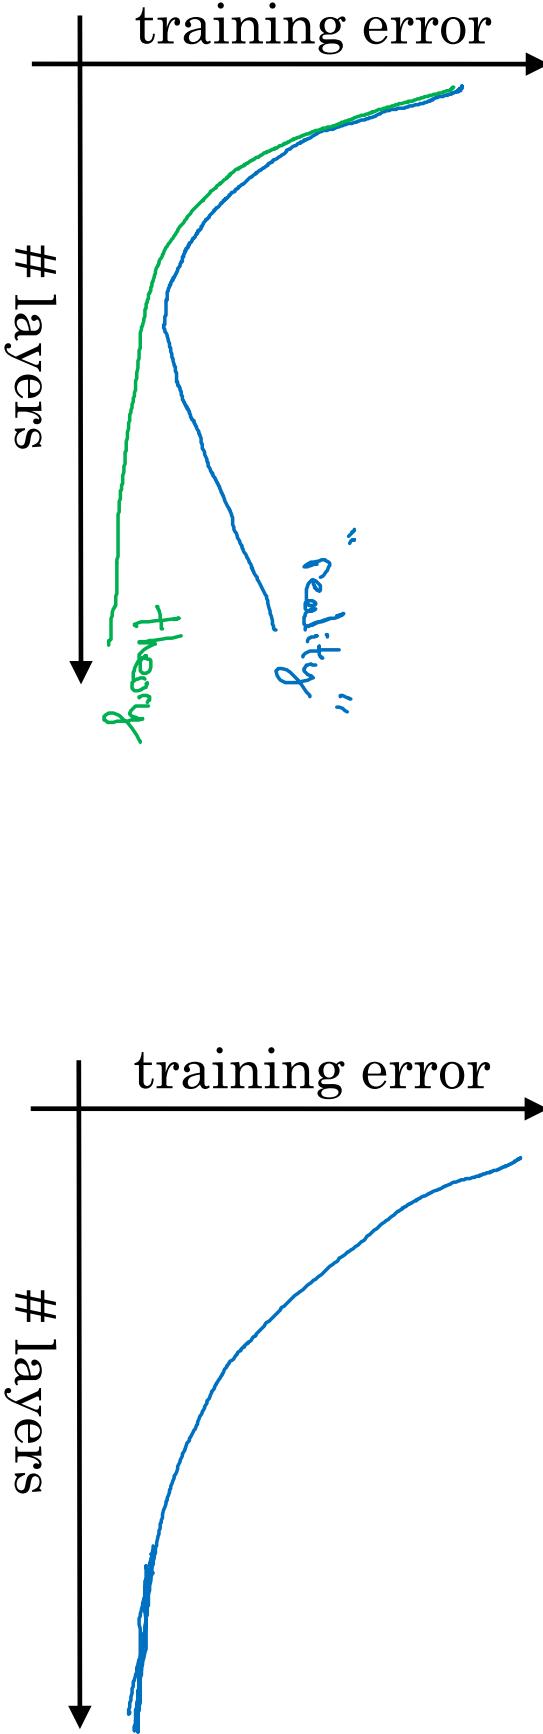
# Residual Networks (ResNets)

# Residual block



[He et al., 2015. Deep residual networks for image recognition]

# Residual Network



[He et al., 2015. Deep residual networks for image recognition]



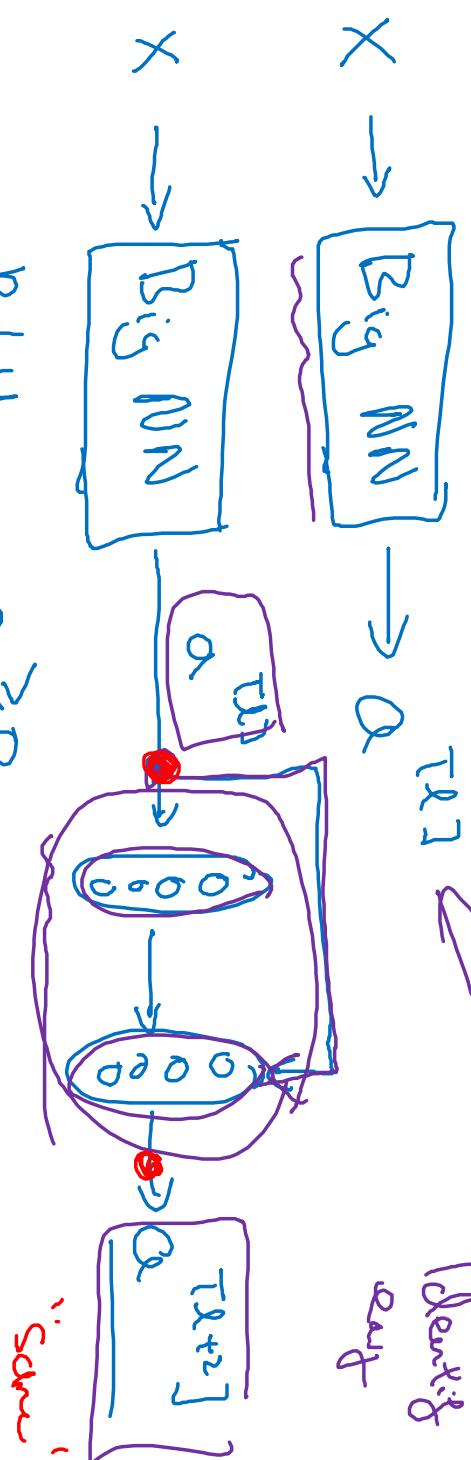
deeplearning.ai

# Why ResNets work

---

## Case Studies

# Why do residual networks work?

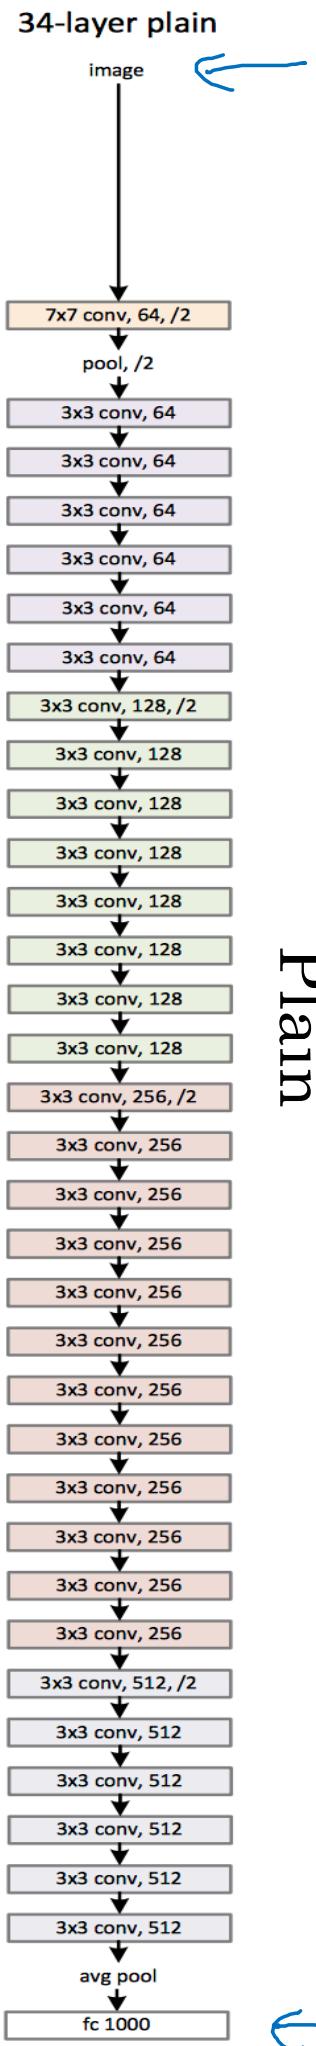


Identifying fraction & ratio  
for Residual blocks  
to learn!

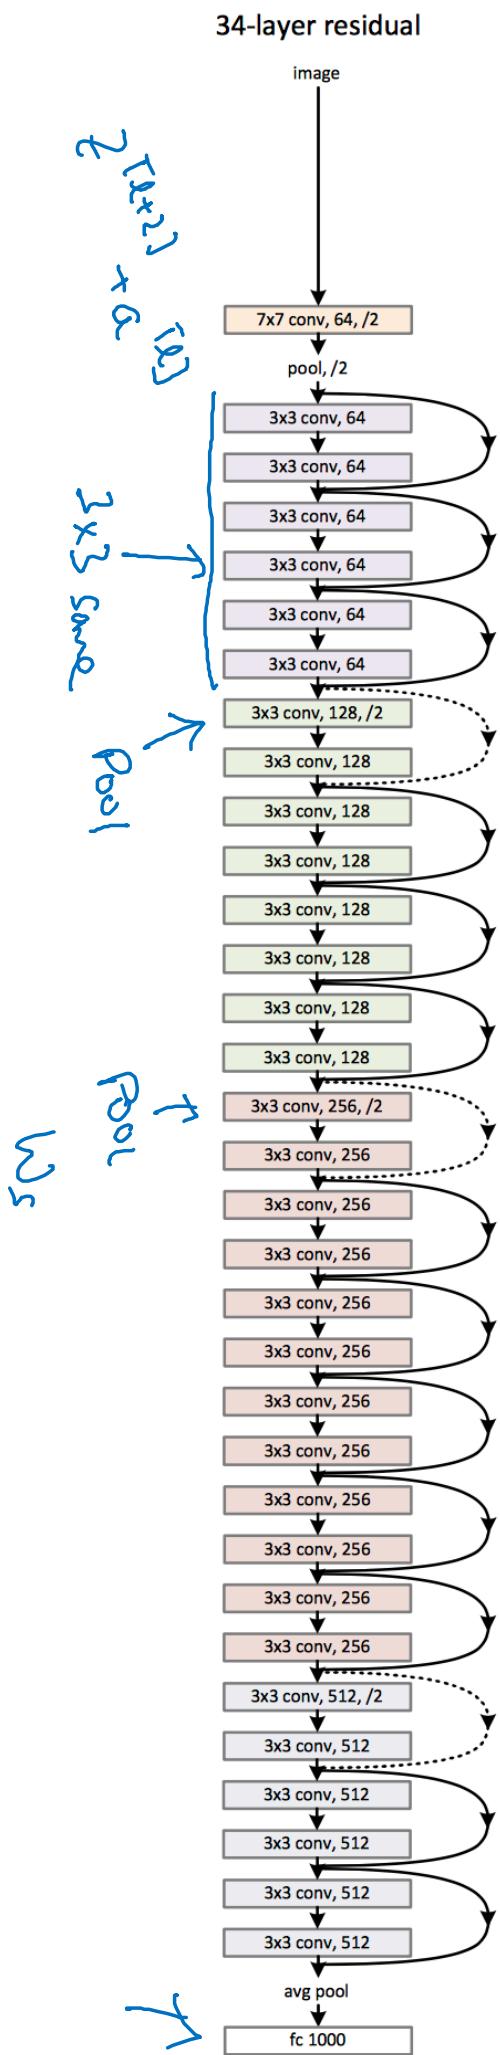
$$\begin{aligned}
 Q^{[L+2]} &= g\left(\frac{\omega^{[L+2]} + a}{2} + b\right) \\
 &= g\left(\frac{\omega^{[L+2]} + \bar{\omega}^{[L+2]}}{2} + \frac{\omega_s a^{[L+2]}}{2}\right) = g(a^{[L+2]}) \\
 \text{If } \omega^{[L+2]} = 0, b^{[L+2]} = 0 \quad \text{128} &\qquad = a^{[L+2]}
 \end{aligned}$$

# ResNet

## Plain



## ResNet



[He et al., 2015. Deep residual networks for image recognition]



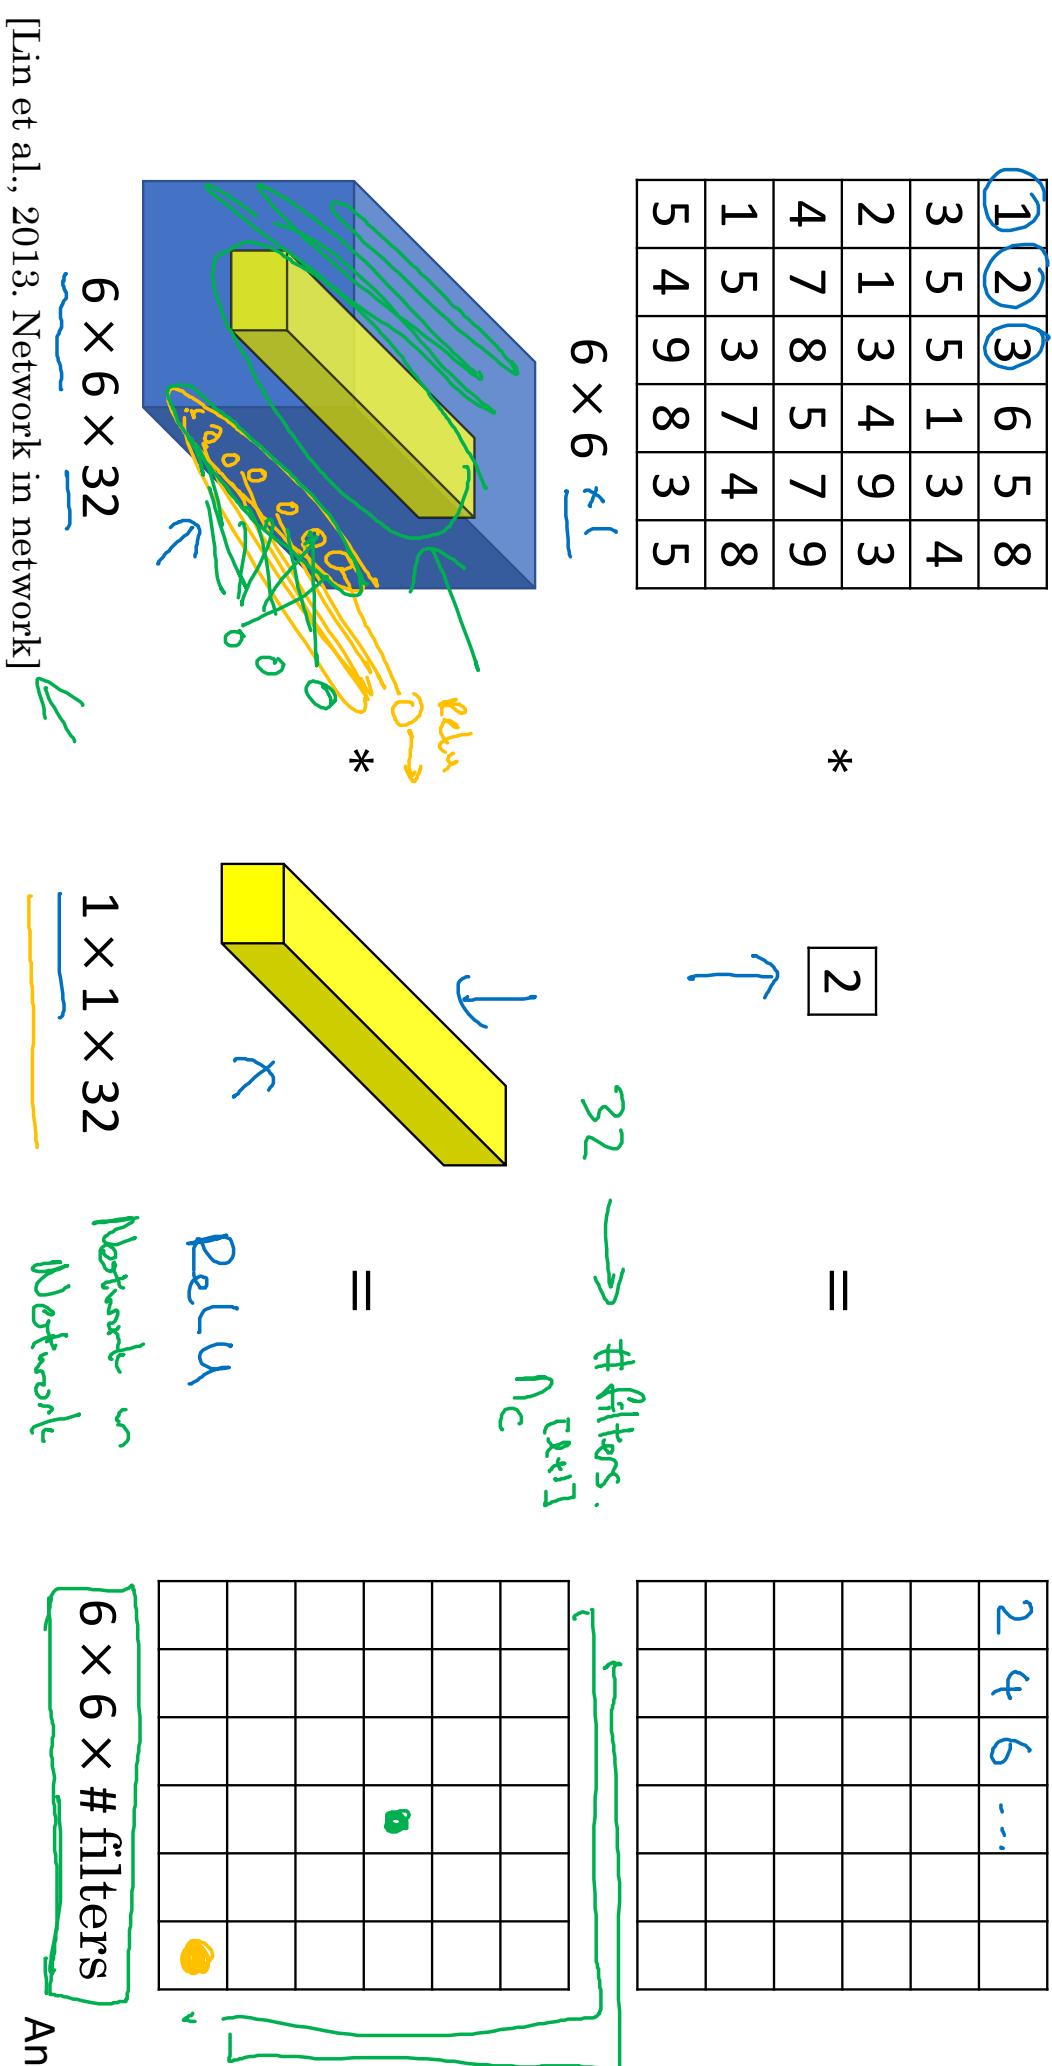
## Case Studies

---

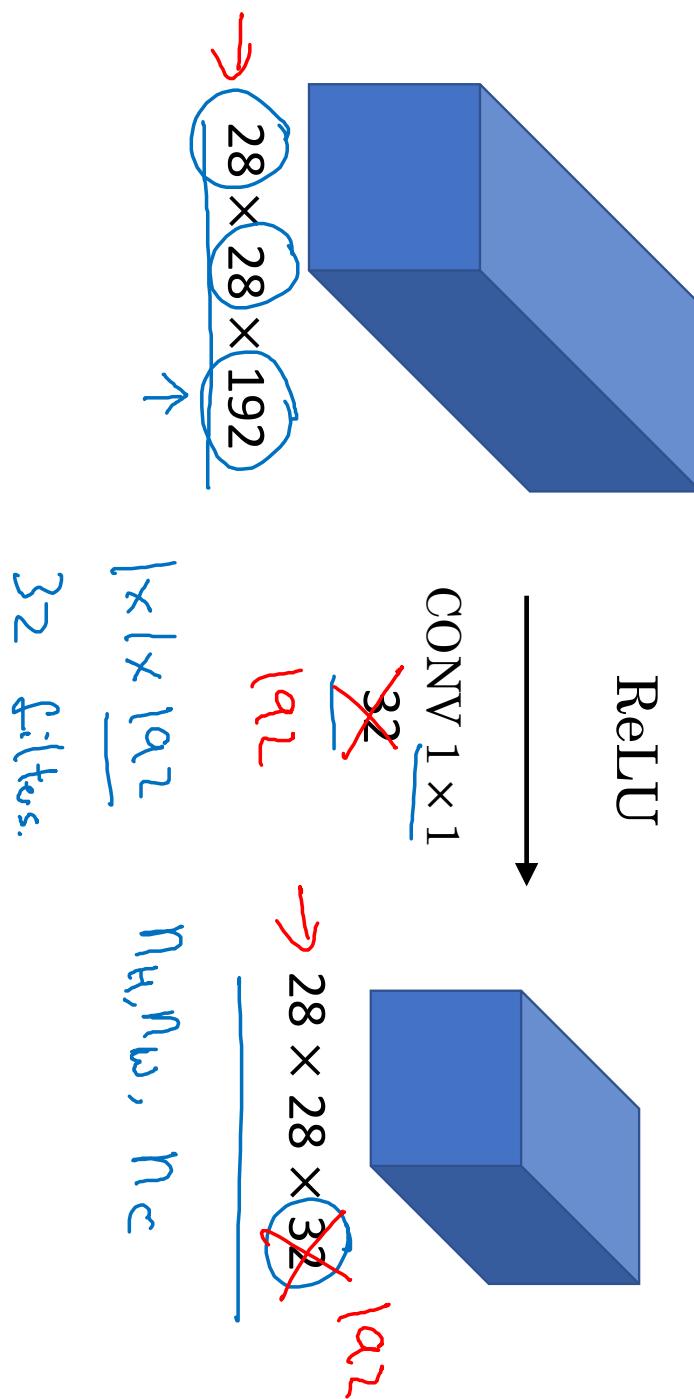
Network in Network  
and  $1 \times 1$  convolutions

deeplearning.ai

# Why does a $1 \times 1$ convolution do?



# Using $1 \times 1$ convolutions



[Lin et al., 2013. Network in network]



deeplearning.ai

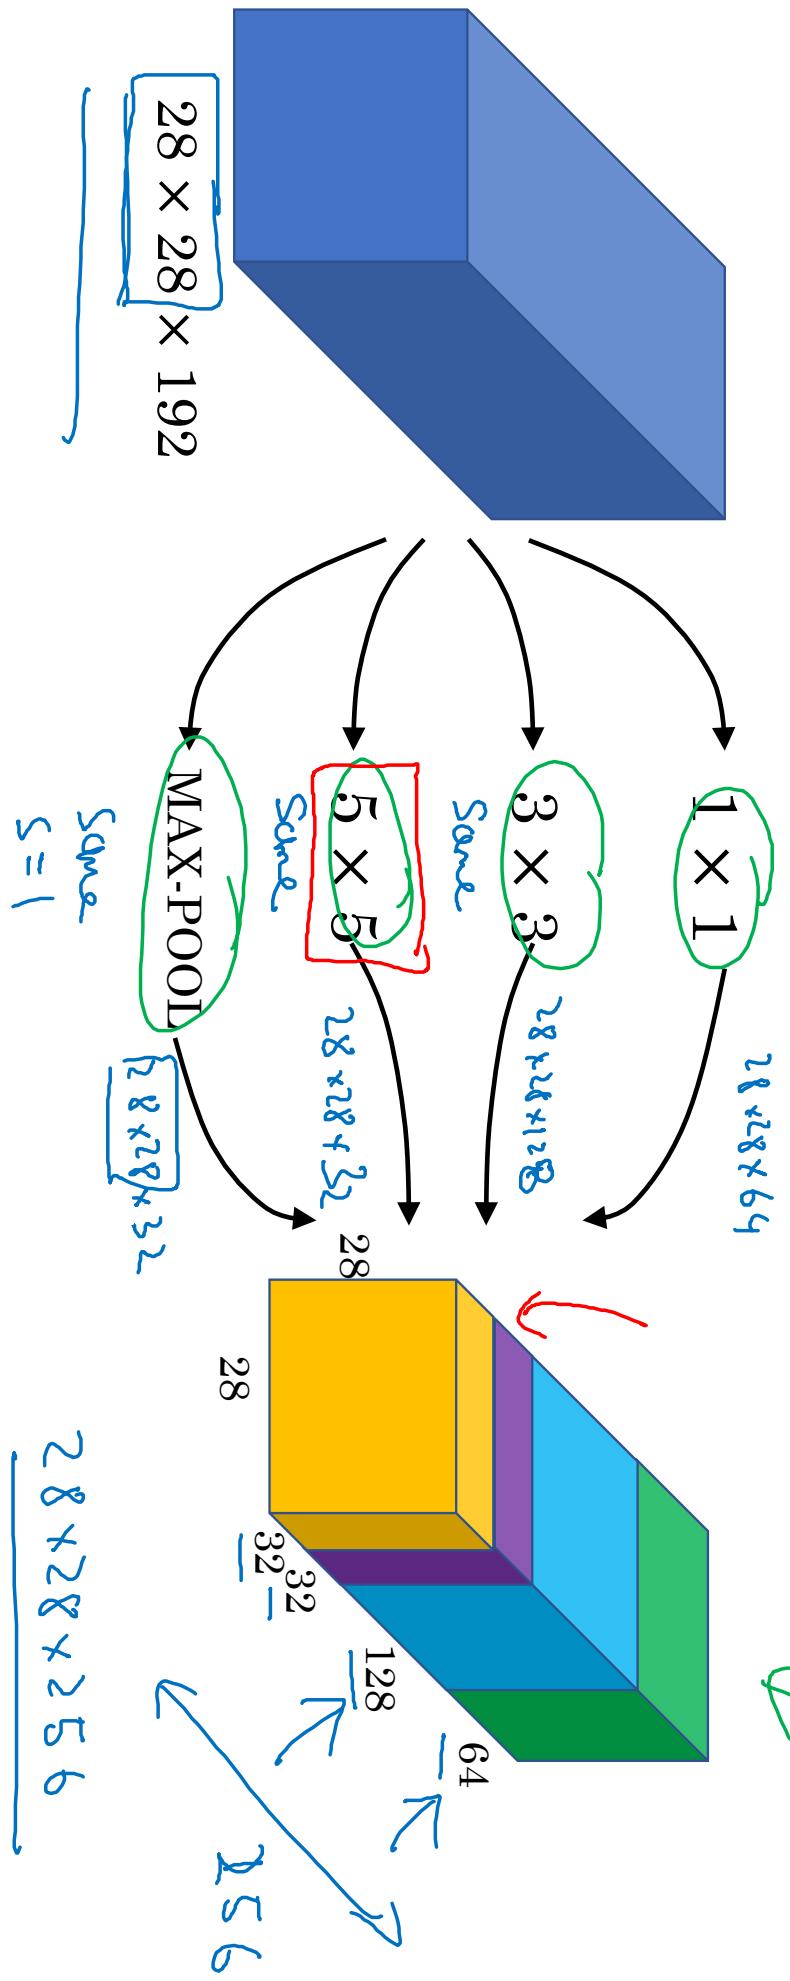
## Case Studies

---

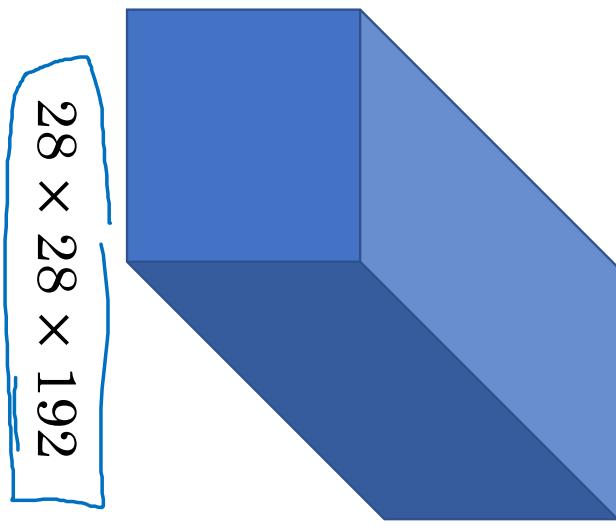
# Inception network motivation

# Motivation for inception network

[Szegedy et al. 2014. Going deeper with convolutions]



# The problem of computational cost

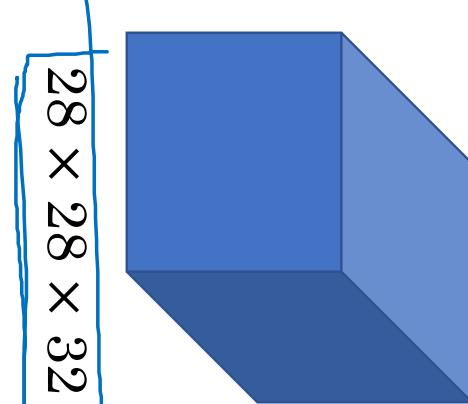


CONV

$5 \times 5$ ,

same,

32



32 filters.

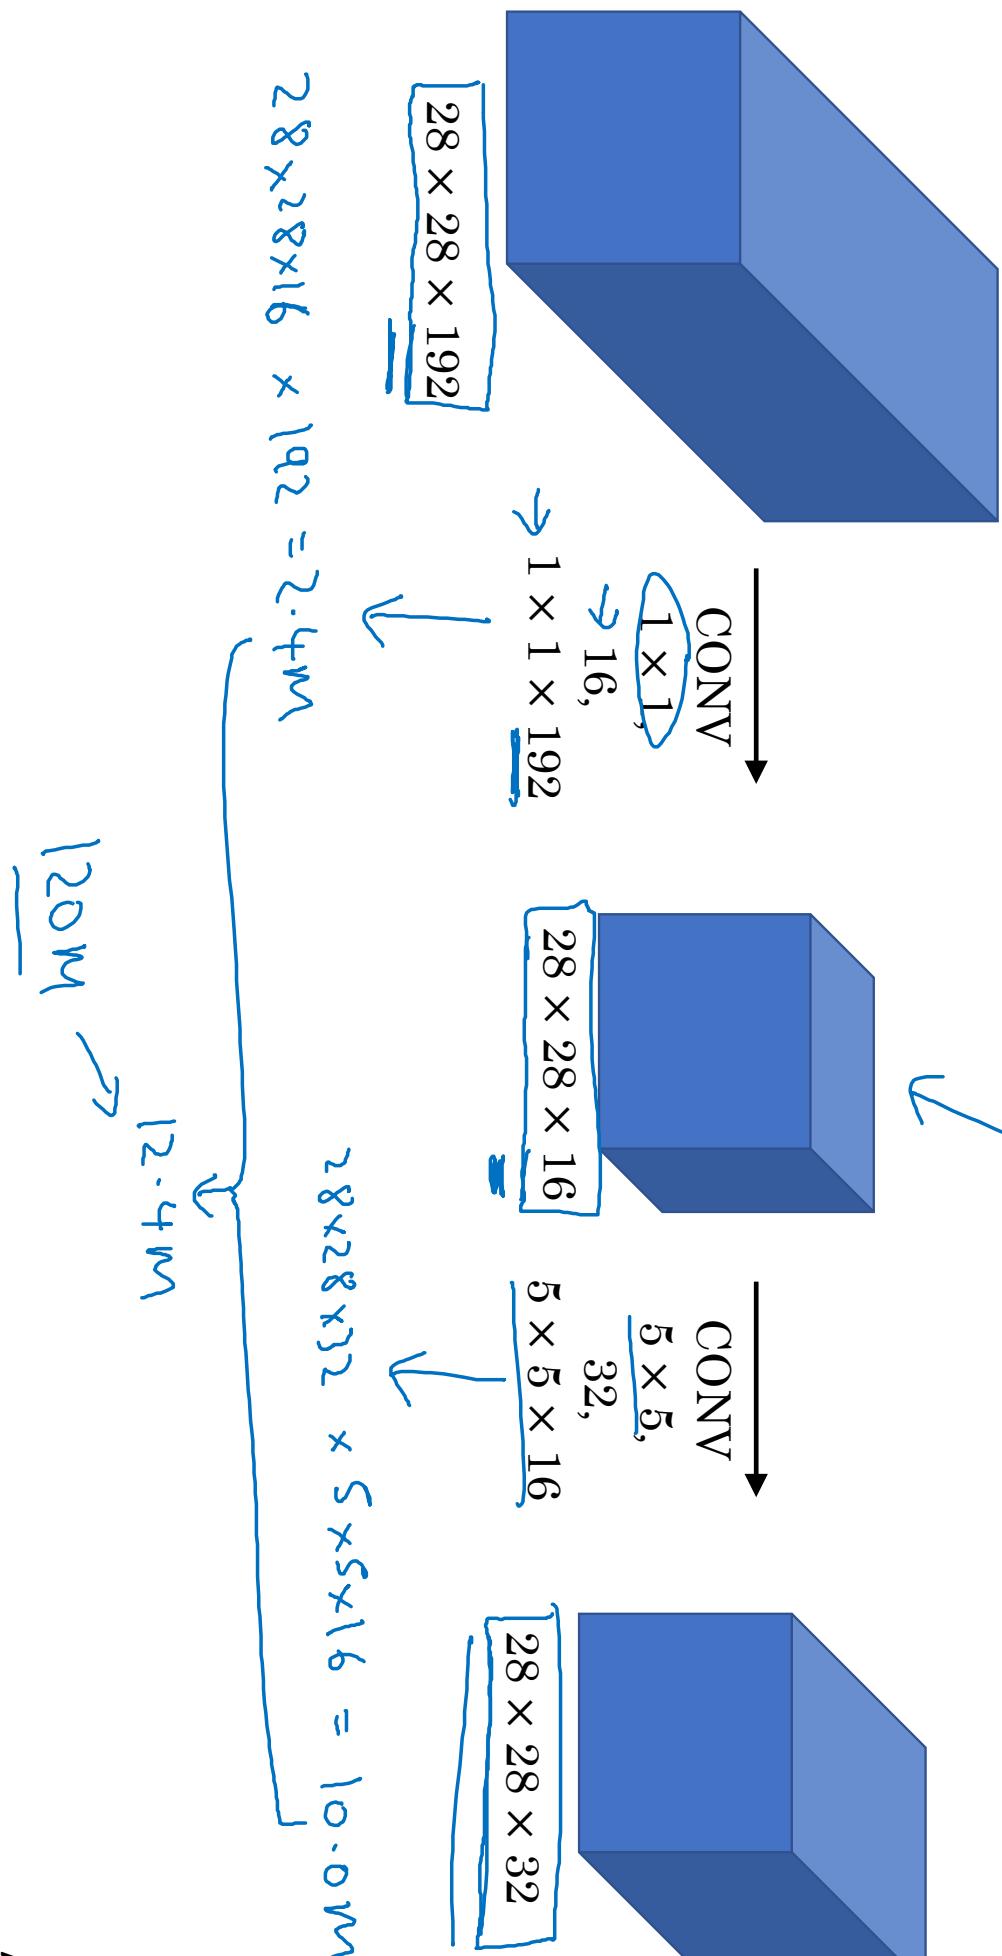
filters are

$5 \times 5 \times 192$ .

$$\frac{28 \times 28 \times 32 \times 5 \times 5 \times 192}{32} = 120M.$$

# Using $1 \times 1$ convolution

"bottleneck layer"





deeplearning.ai

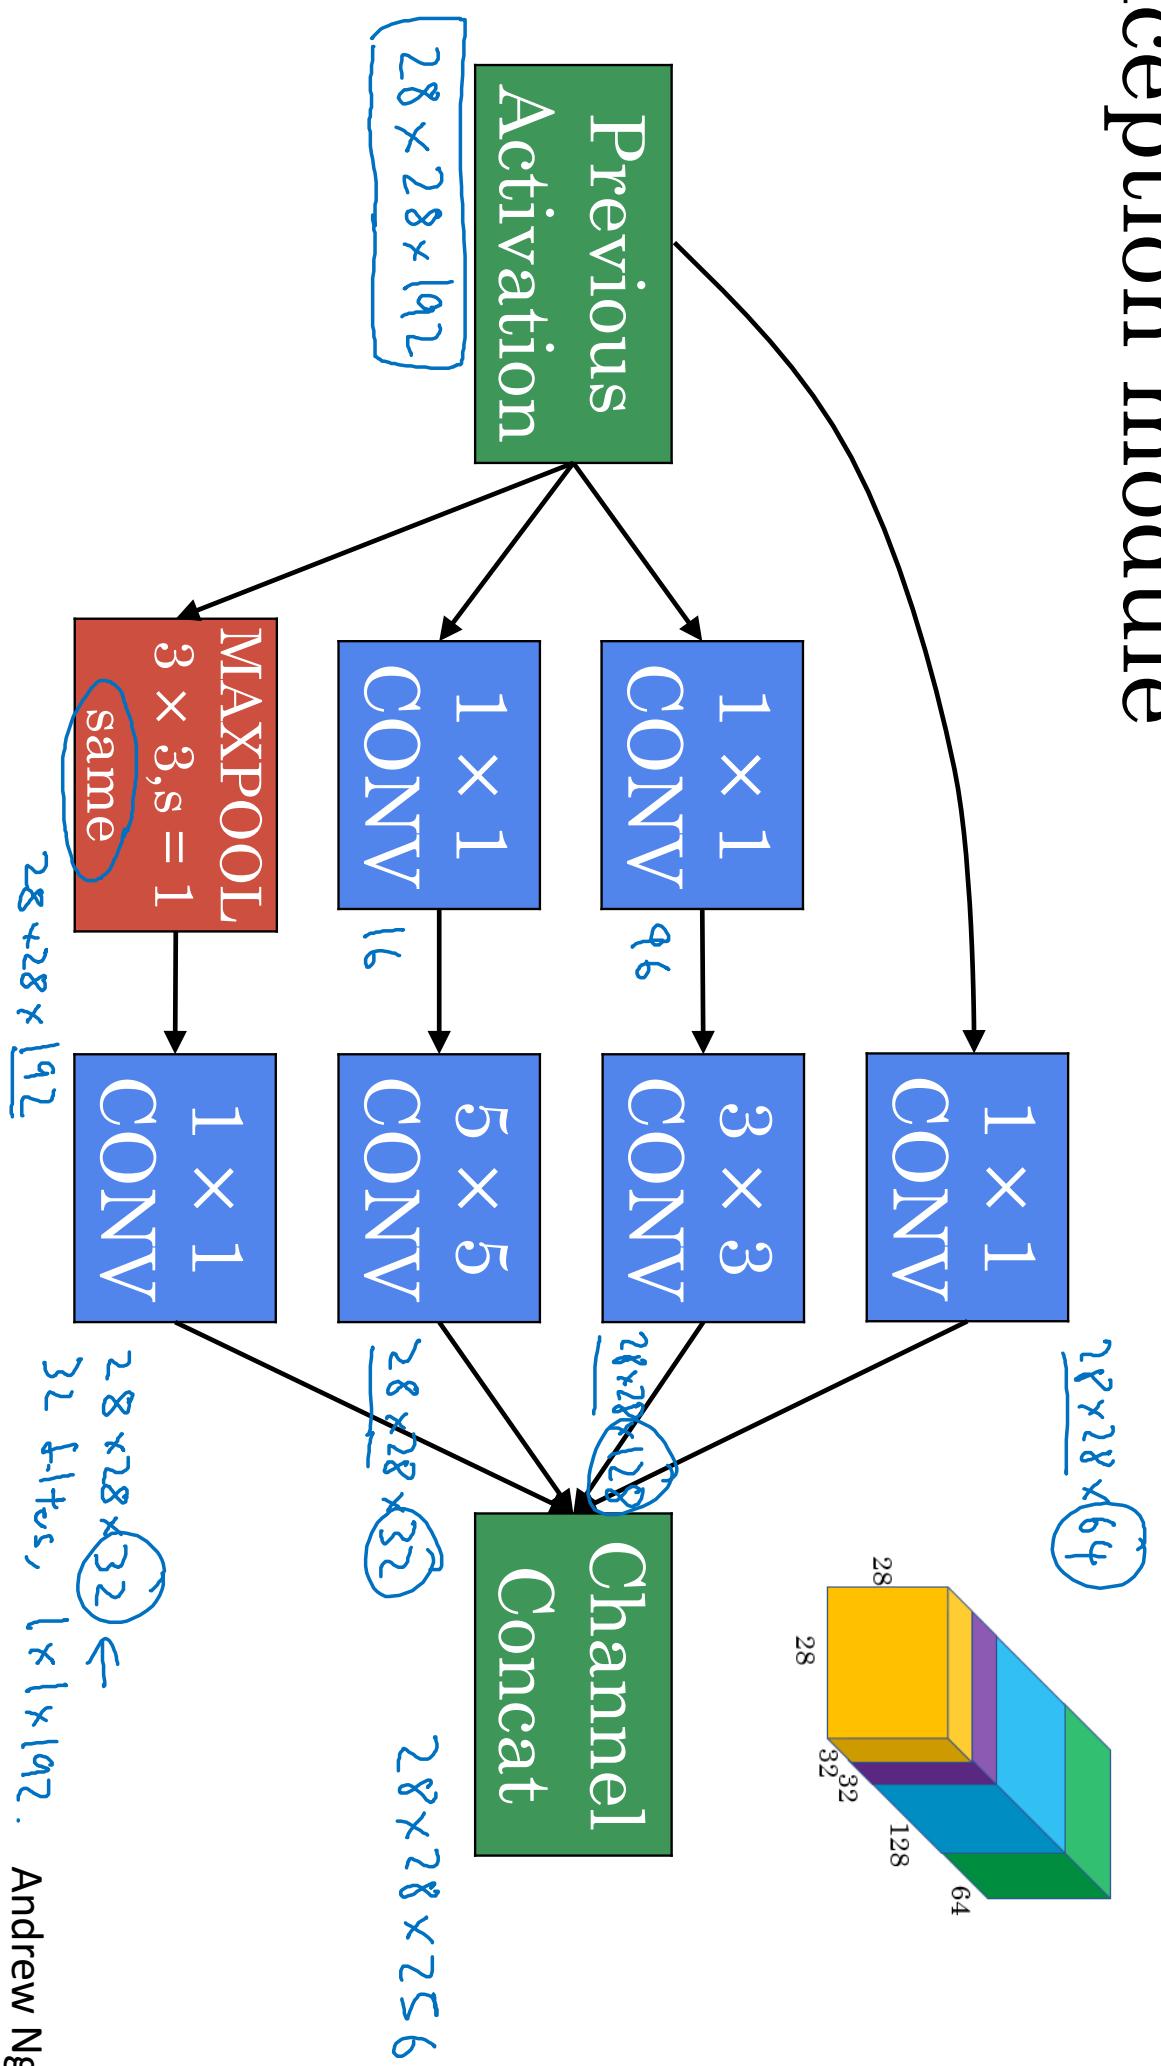
## Case Studies

---

## Inception network

---

# Inception module

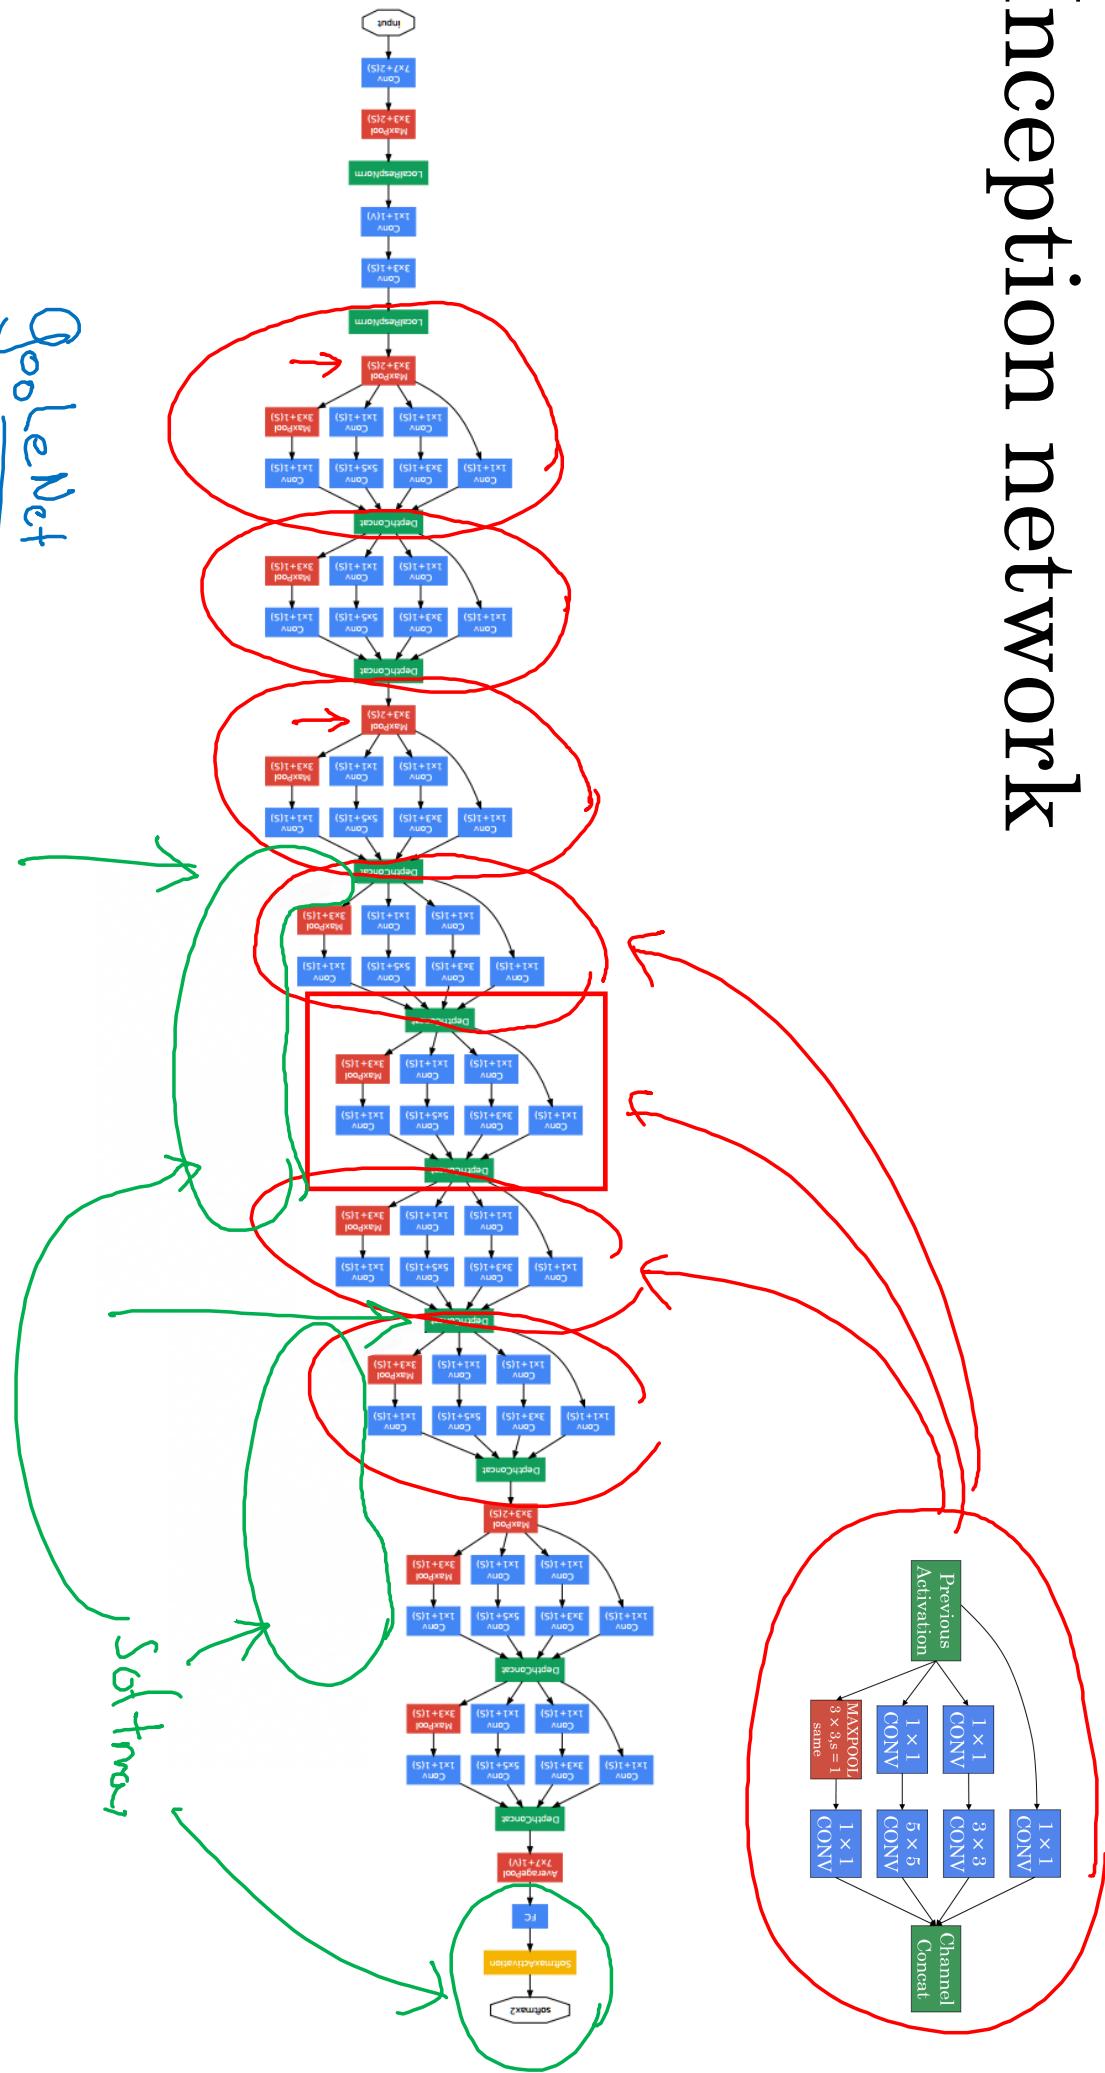


# Inception network

GooleNet

Sofmax

[Szegedy et al., 2014, Going Deeper with Convolutions]



<http://knowyourmeme.com/memes/we-need-to-go-deeper>



Andrew Ng



---

# Convolutional Neural Networks

## MobileNet

deeplearning.ai

# Motivation for MobileNets

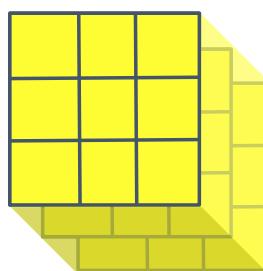
- Low computational cost at deployment
- Useful for mobile and embedded vision applications
- Key idea: Normal vs. depthwise-separable convolutions



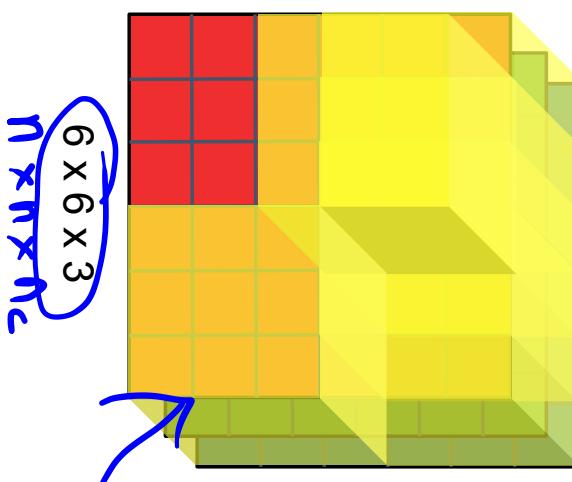
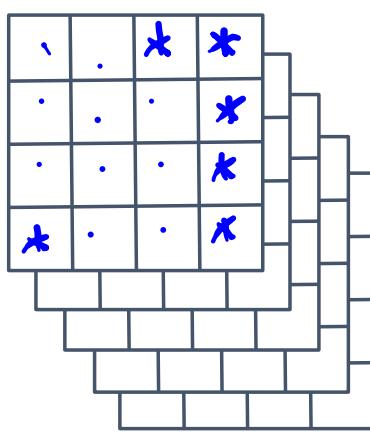
# Normal Convolution

$\nwarrow n_c' \text{ filters}$

\*



=



$\rightarrow 3 \times 3 \times 3$   
 $f \times f \times n_c$

$\nwarrow n_{out} \times n_{out} \times n_c'$

Computational cost

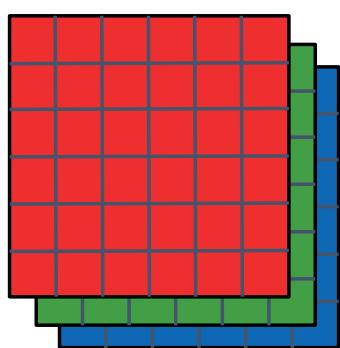
=  $\frac{\# \text{filter params}}{\# \text{filter positions}} \times \# \text{of filters}$

$\rightarrow \underline{\underline{2160}}$

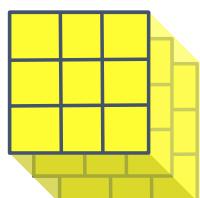
$$= \frac{3 \times 3 \times 3}{3 \times 3 \times 3} \times \frac{4 \times 4}{4 \times 4} \times 5$$

# Depthwise Separable Convolution

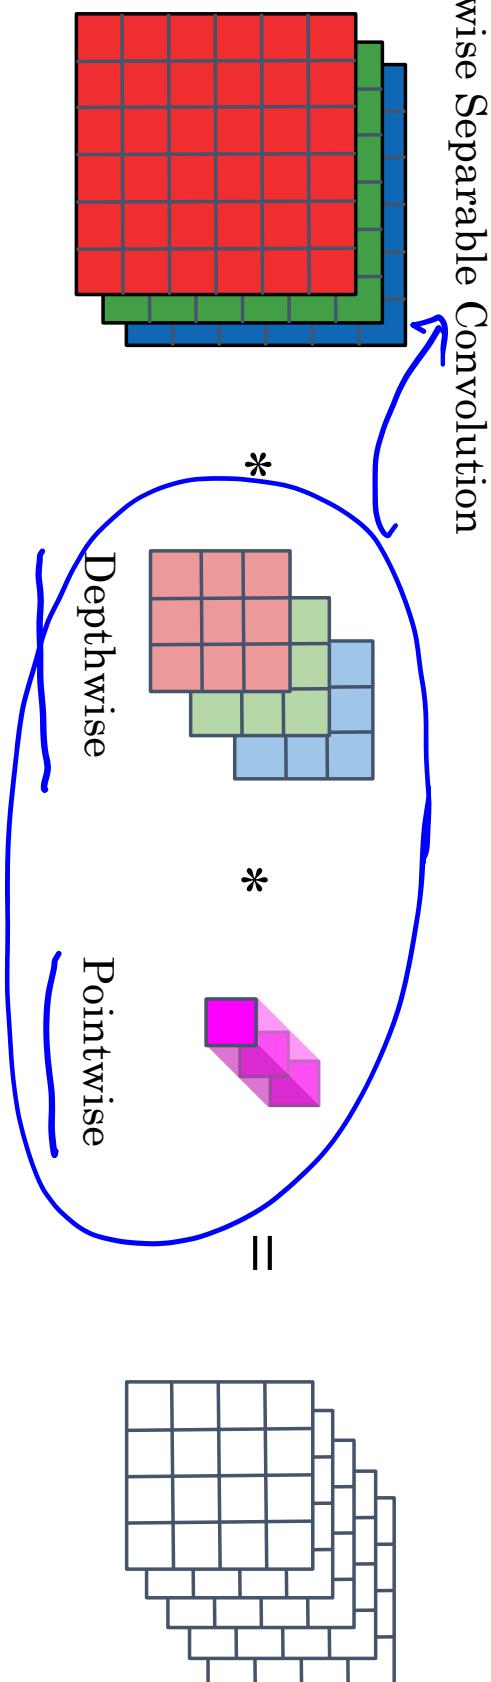
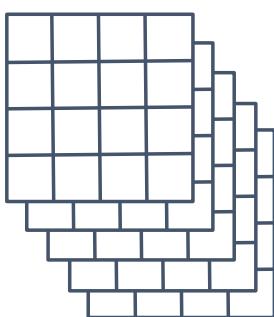
Normal Convolution



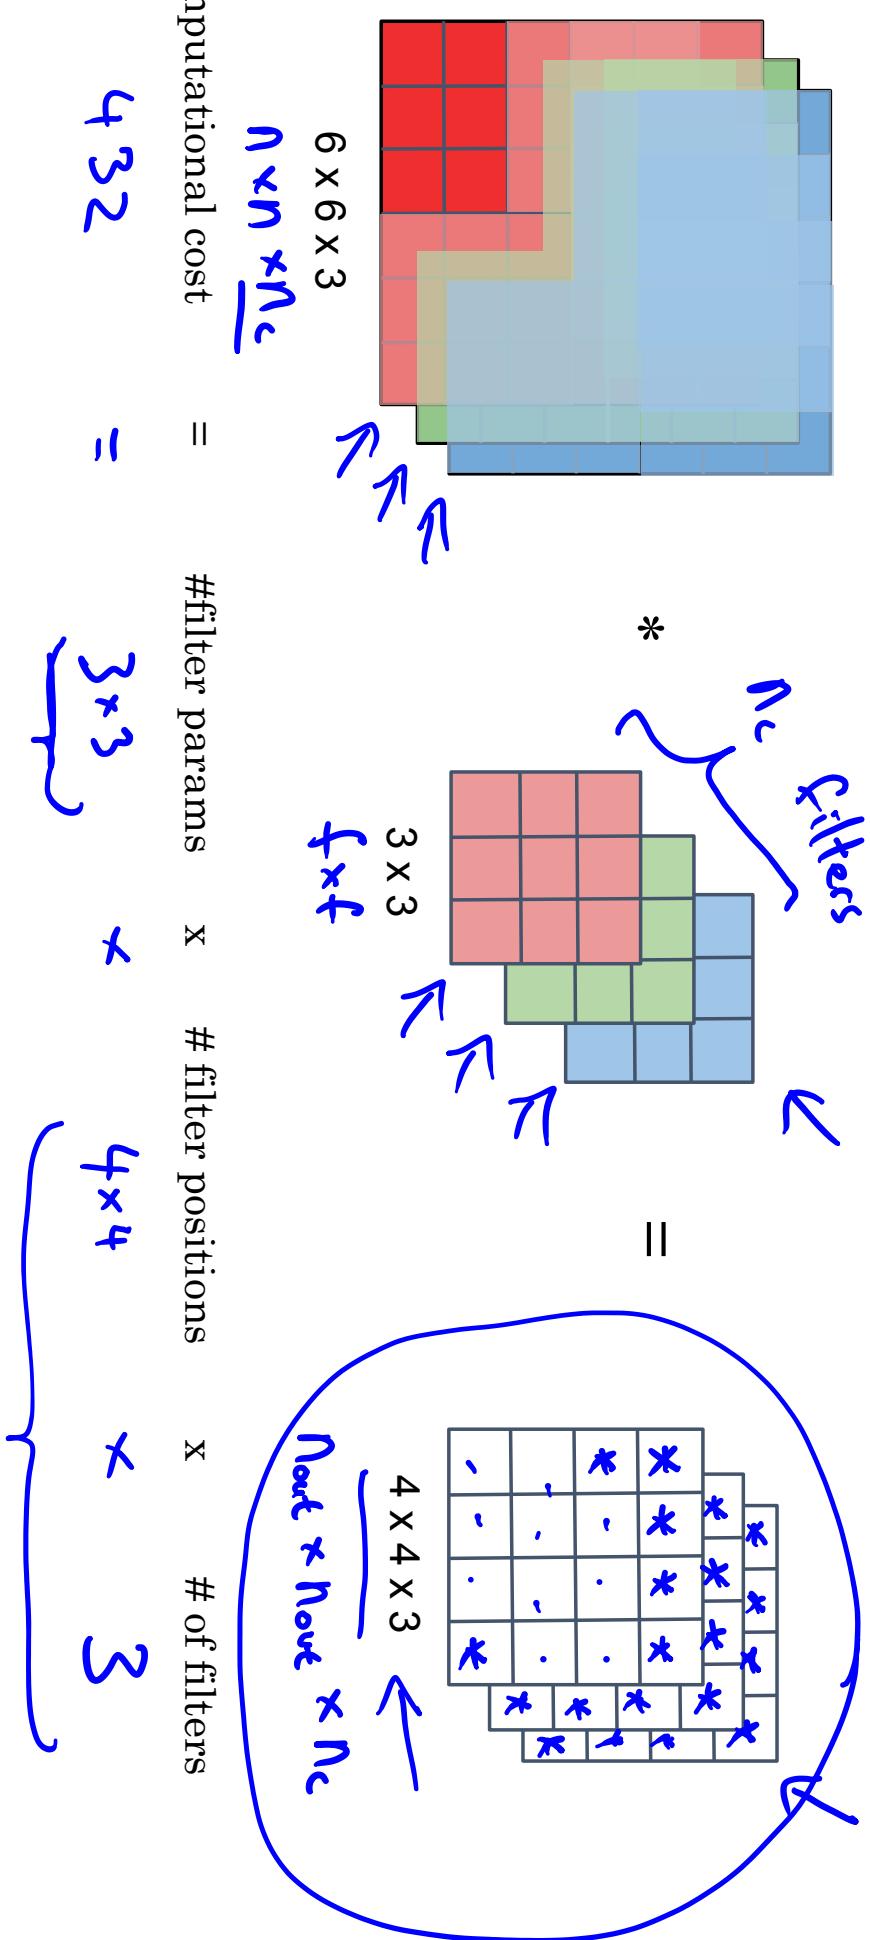
\*



=

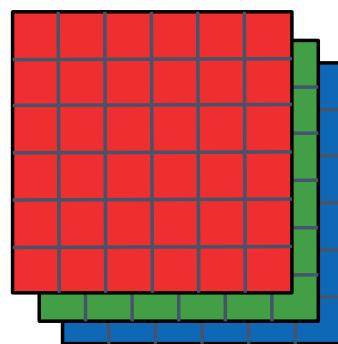


# Depthwise Convolution



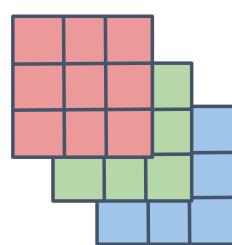
# Depthwise Separable Convolution

Depthwise Convolution



Pointwise Convolution

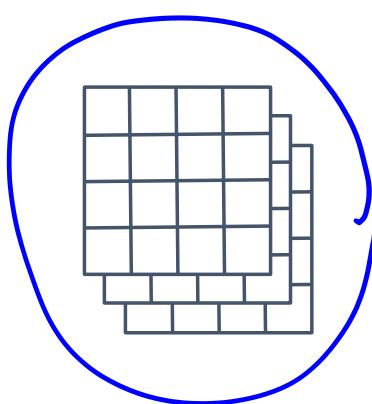
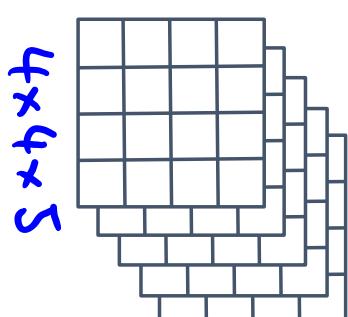
\*



\*



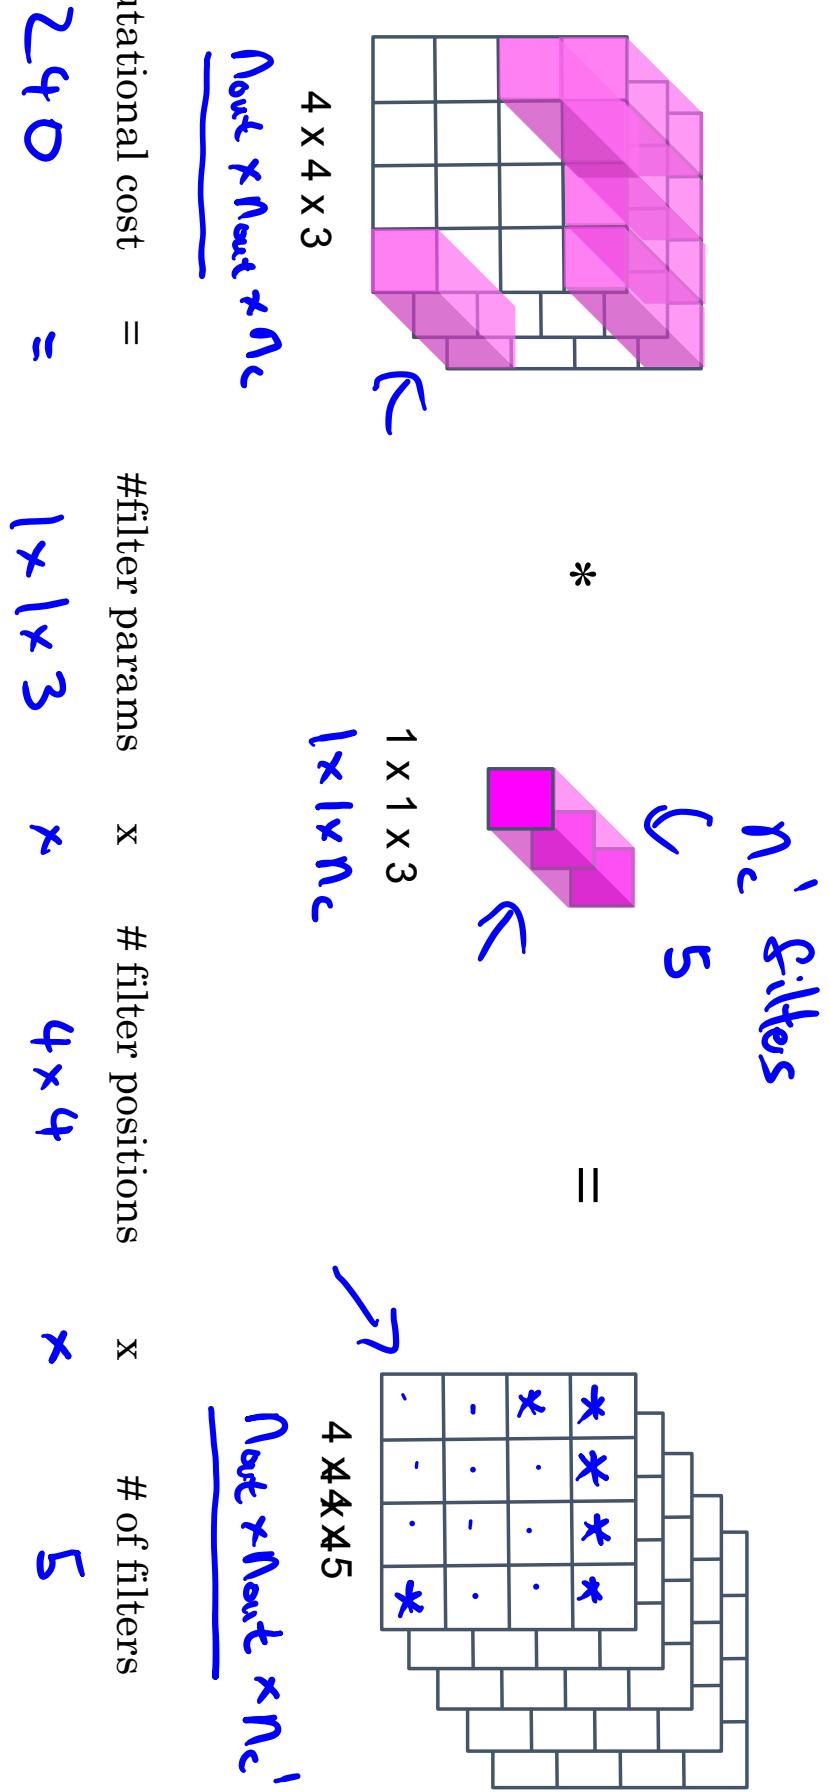
=



=

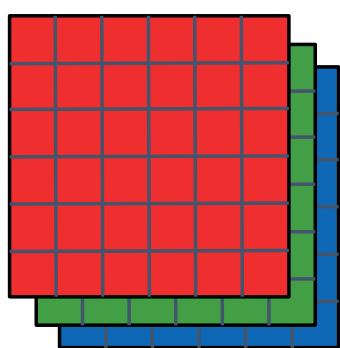
432

# Pointwise Convolution

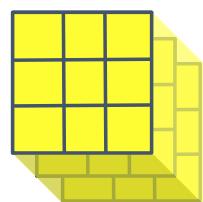


# Depthwise Separable Convolution

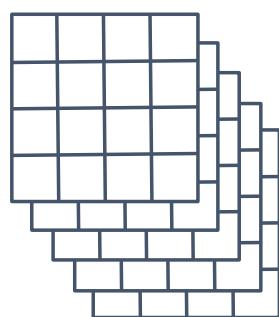
Normal Convolution



\*

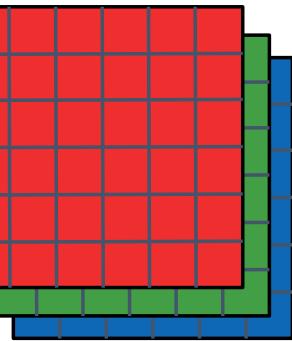


=

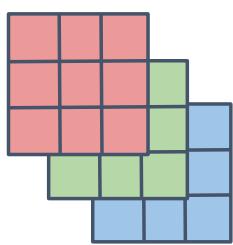


$4 \times 4 \times 3$

Depthwise Separable Convolution



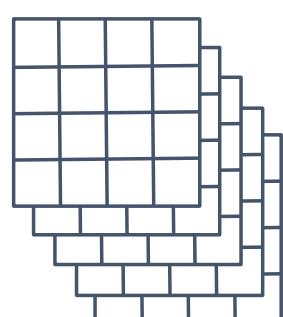
\*



\*



=



Depthwise

Pointwise

# Cost Summary

Cost of normal convolution  $\leftarrow 2160$

Cost of depthwise separable convolution



$$\begin{array}{l} \text{depthwise} \\ 432 \end{array} + \begin{array}{l} \text{pointwise} \\ 240 \end{array} = 672$$

$$= \frac{1}{n_c} + \frac{1}{f^2}$$

$$\frac{1}{s} + \frac{1}{q}$$

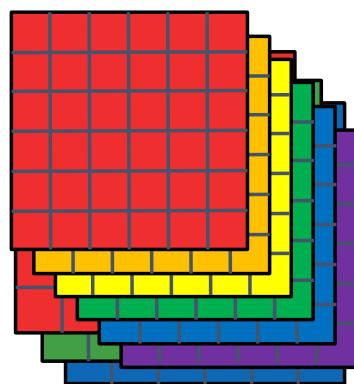
$$= \frac{1}{512} + \frac{1}{3^2}$$

$n/10$  times cheaper

# Depthwise Separable Convolution

Depthwise Convolution

$$6 \times b \times n_c$$



Pointwise Convolution

$$\begin{matrix} * \\ \text{---} \\ 1 \times 1 \times n_c \end{matrix} = \begin{matrix} * \\ \text{---} \\ 3 \times 3 \times n_c \end{matrix} = \begin{matrix} * \\ \text{---} \\ 4 \times 4 \times n_c \end{matrix}$$

The diagram illustrates the computation of a depthwise separable convolution. It shows the input feature map being processed by two stages: a depthwise convolution (1x1 kernel) followed by a pointwise convolution (3x3 kernel). The final output is a 4x4x1 feature map. The input is a 4x4x4 feature map, and the output is a 4x4x1 feature map.



---

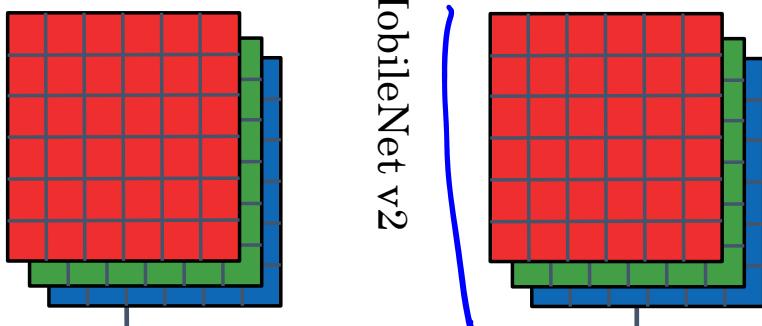
# Convolutional Neural Networks

## MobileNet Architecture

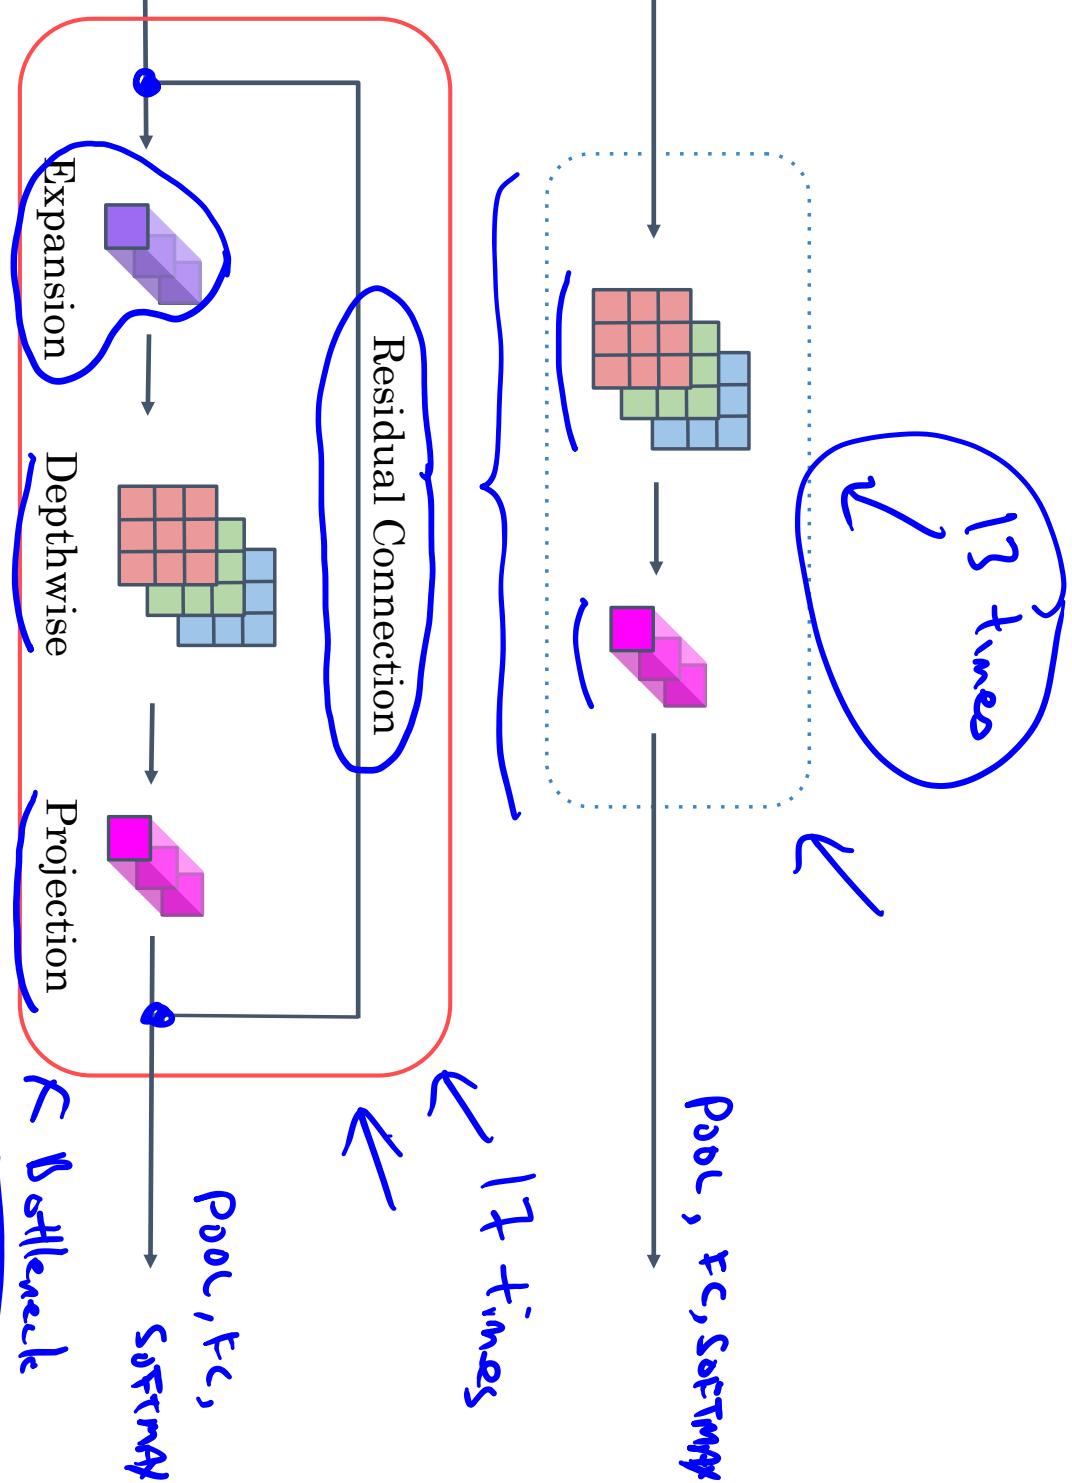
deeplearning.ai

# MobileNet

MobileNet v1



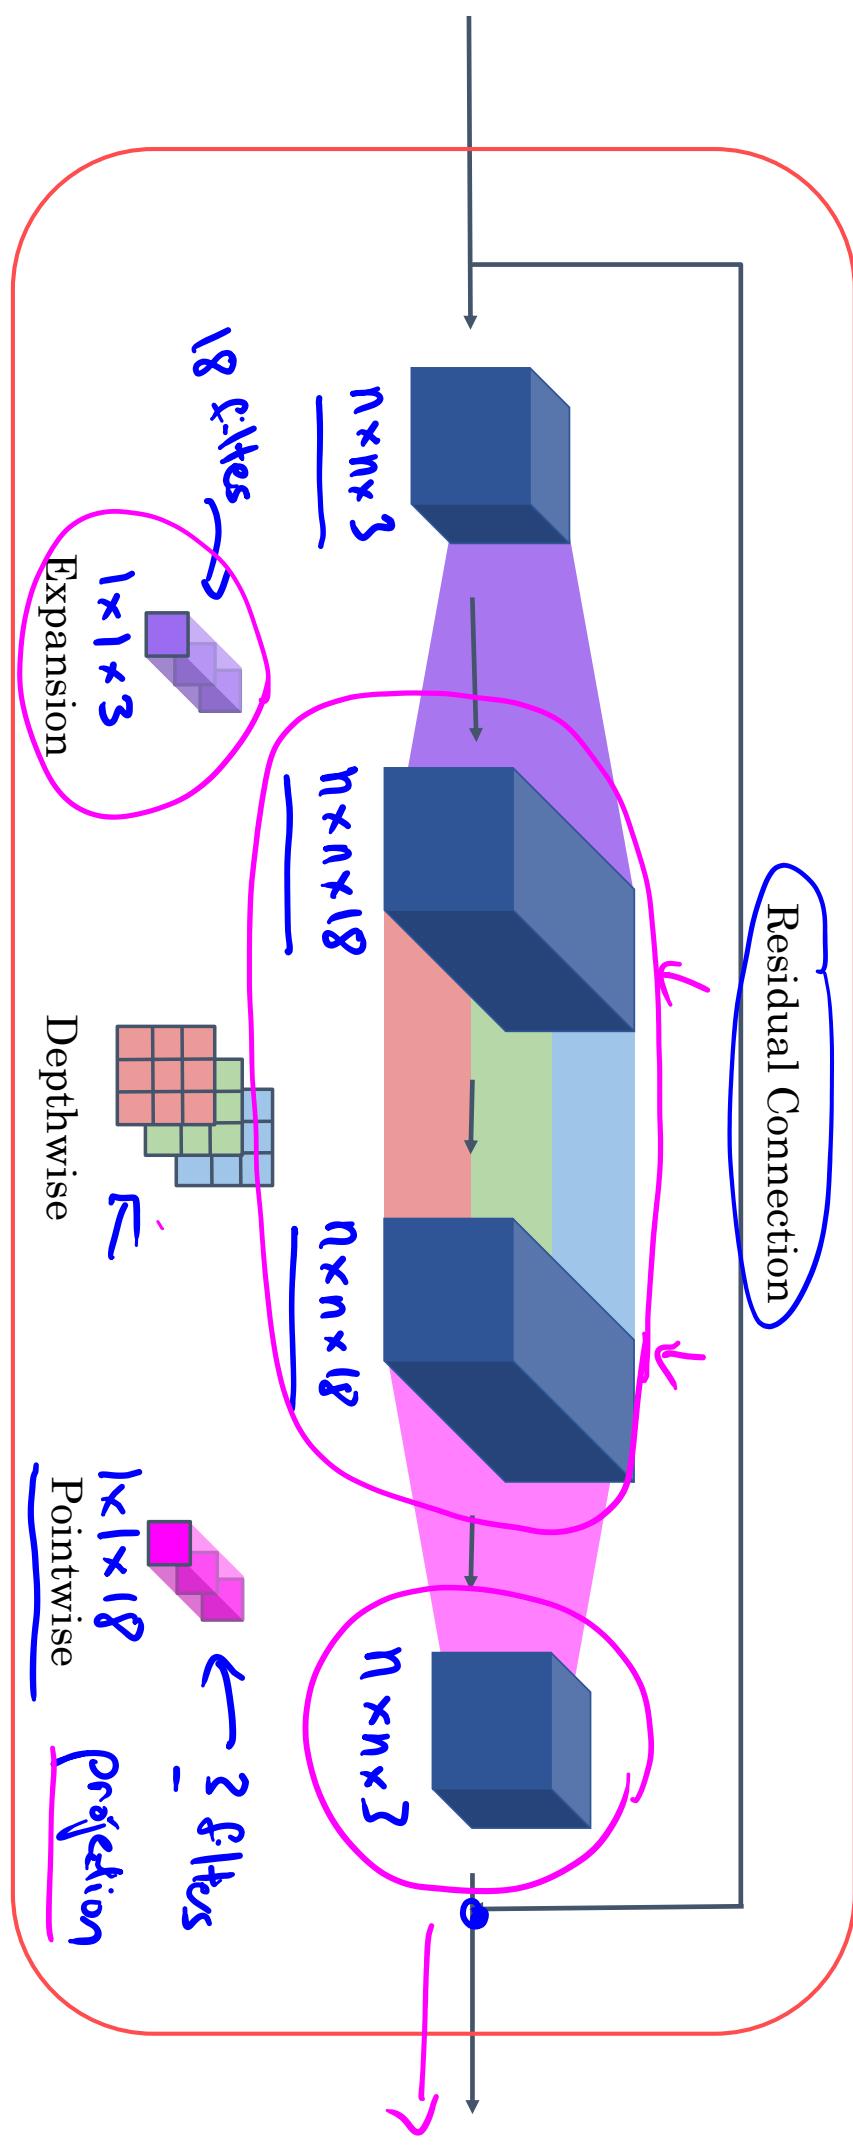
MobileNet v2



[Sandler et al. 2019, MobileNetV2: Inverted Residuals and Linear Bottlenecks]

Andrew Ng

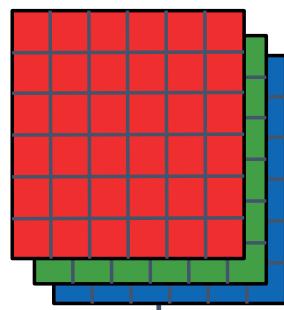
# MobileNet v2 Bottleneck



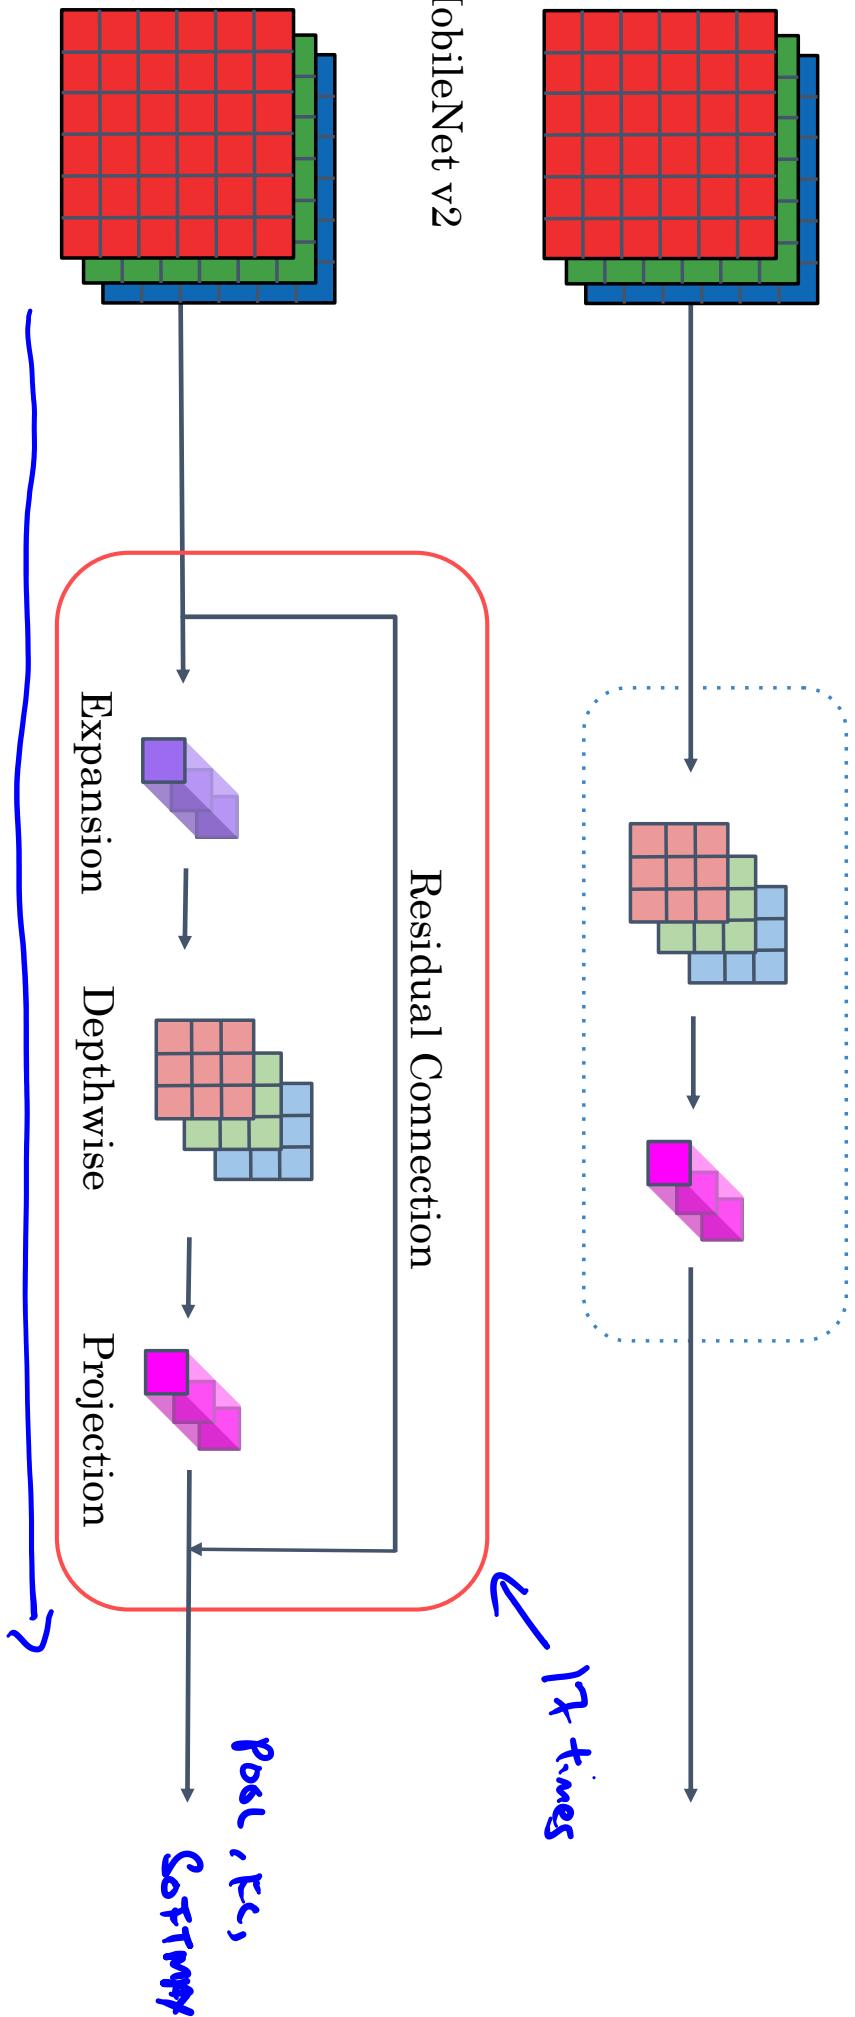
[Sandler et al. 2019, MobileNetV2: Inverted Residuals and Linear Bottlenecks]

# MobileNet

MobileNet v1



MobileNet v2



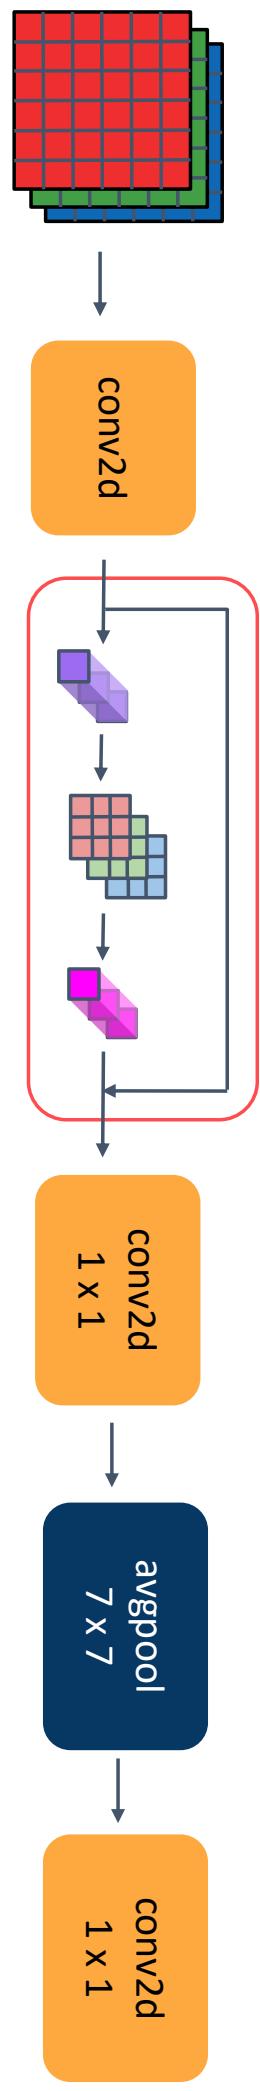
[Sandler et al. 2019, MobileNetV2: Inverted Residuals and Linear Bottlenecks]

Andrew Ng

# MobileNet v2 Full Architecture

[Sandler et al. 2019, MobileNetV2: Inverted Residuals and Linear Bottlenecks]

Andrew Ng





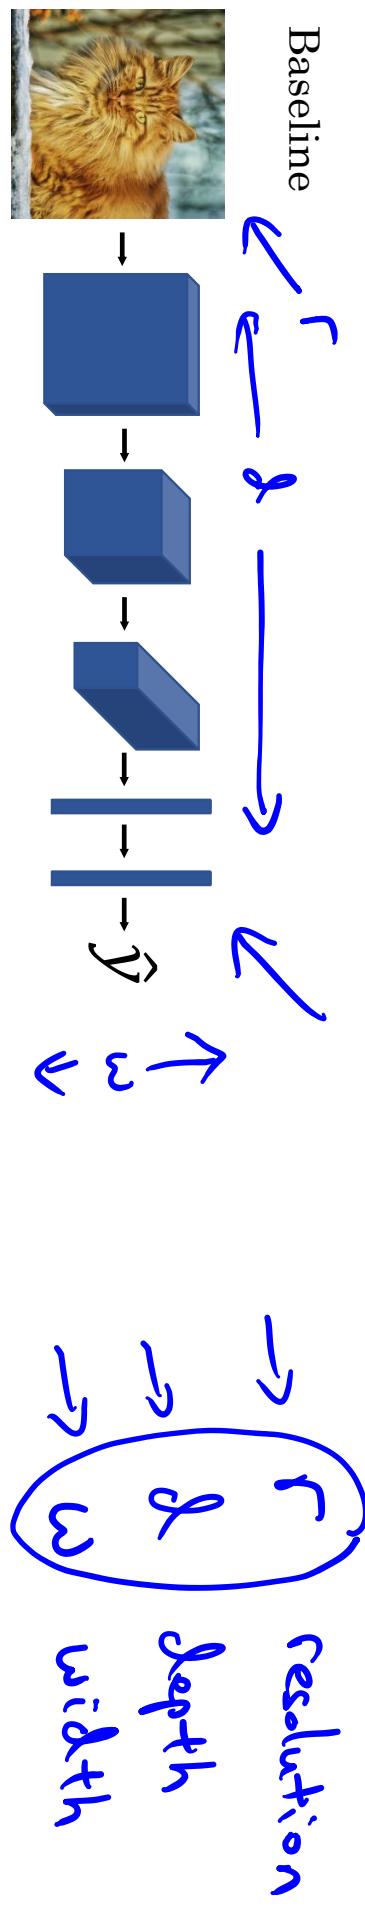
---

# Convolutional Neural Networks

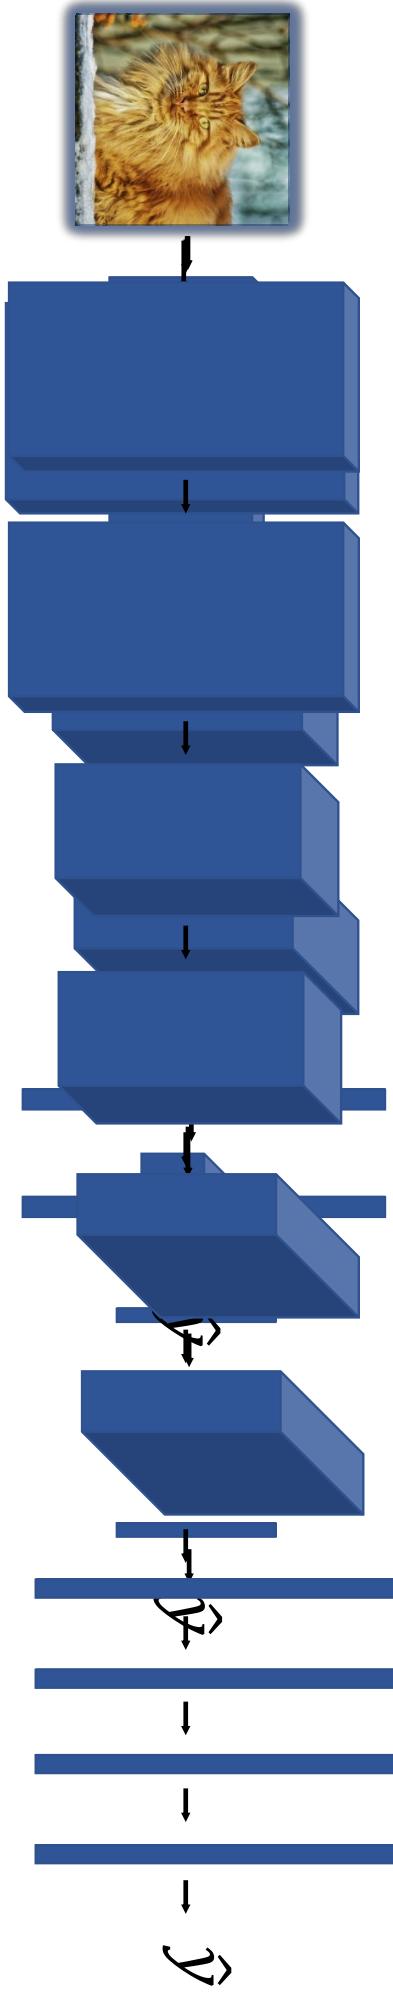
## EfficientNet

deeplearning.ai

# EfficientNet



HighperResetting



[Tan and Le, 2019, EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks]



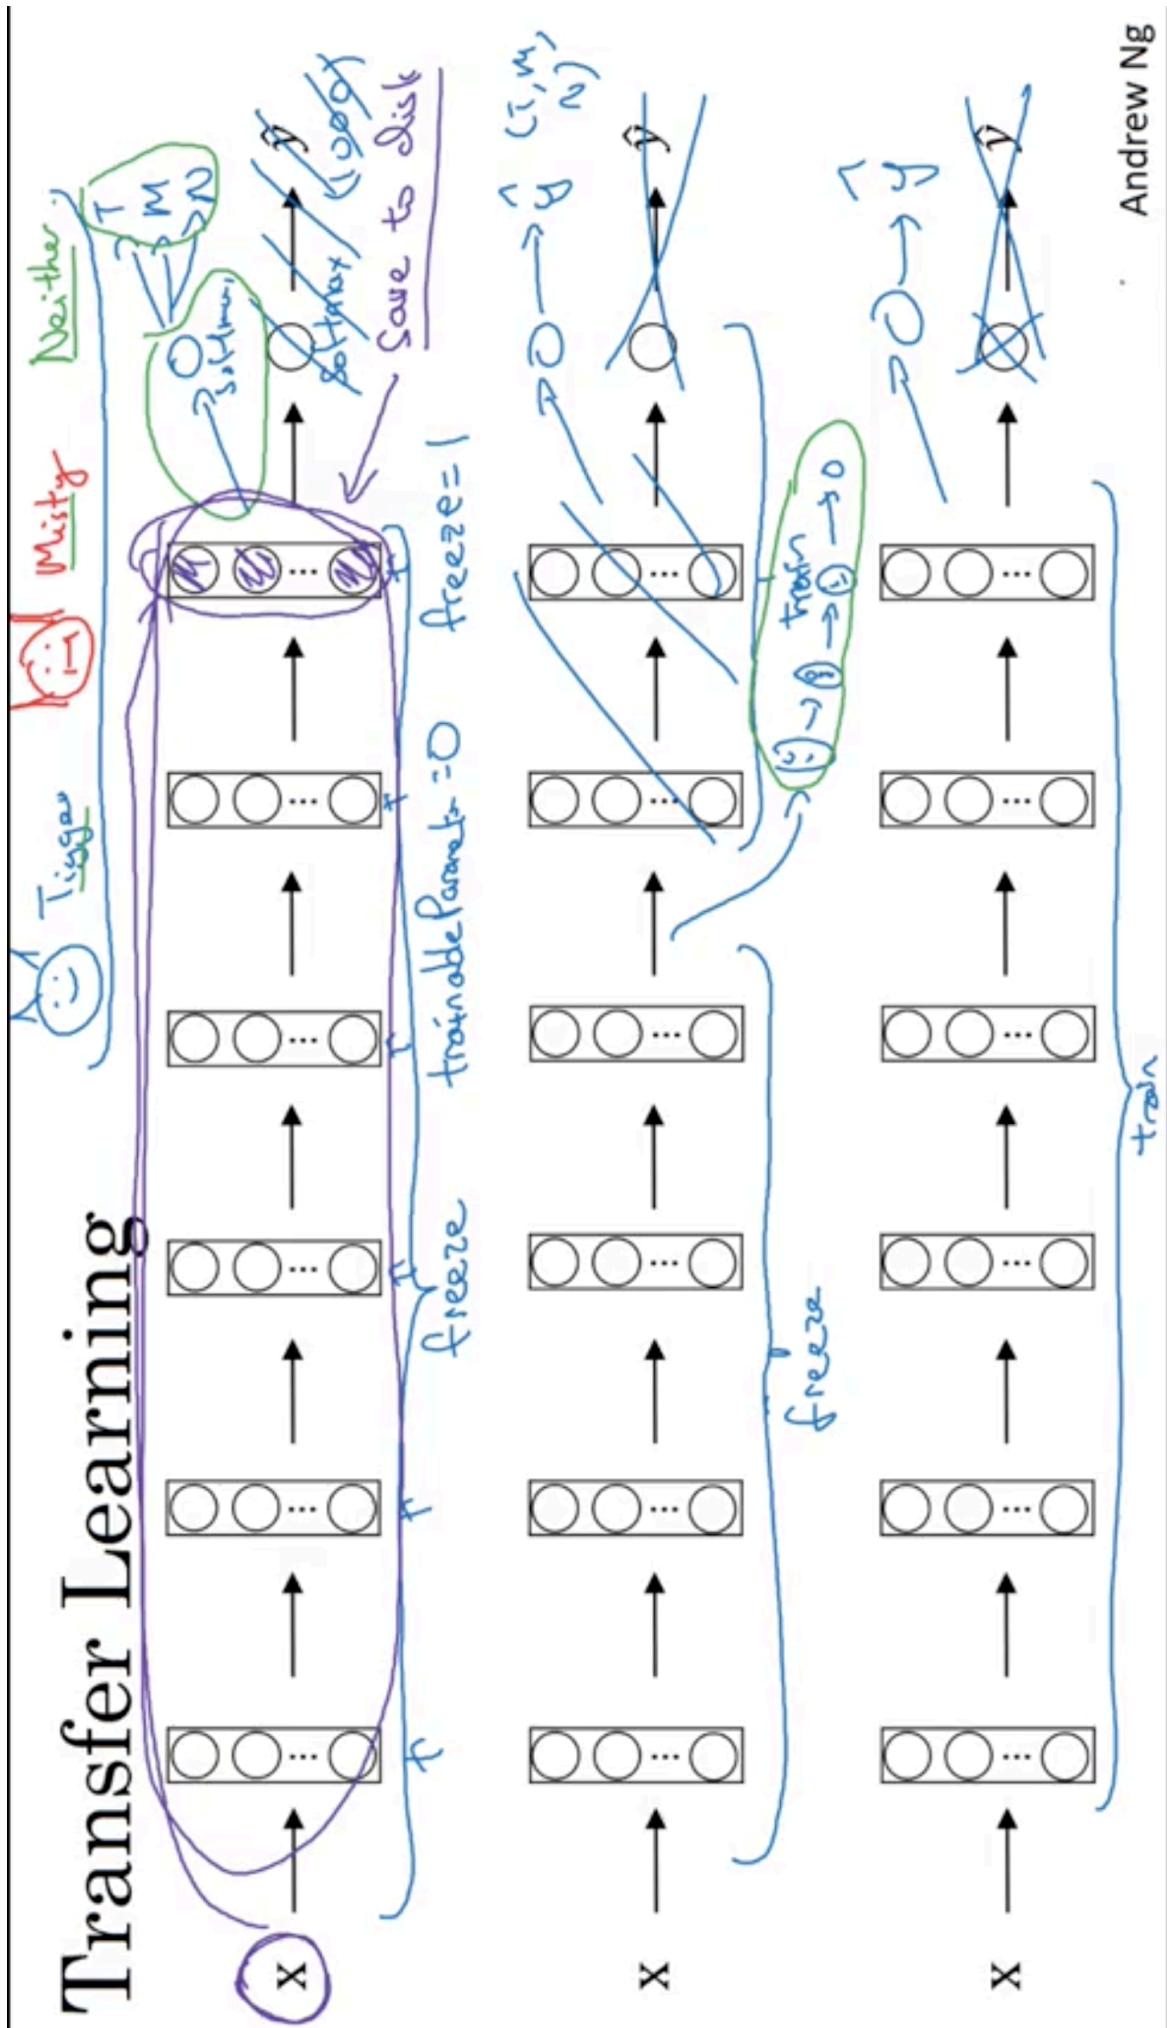
Practical advice for  
using ConvNets

# Transfer Learning

---

deeplearning.ai

Transfer Learning



Andrew Ng

Practical advice for

using ConvNets

---

Data augmentation



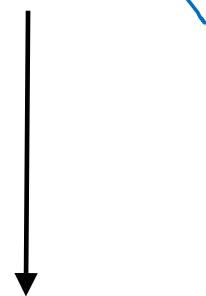
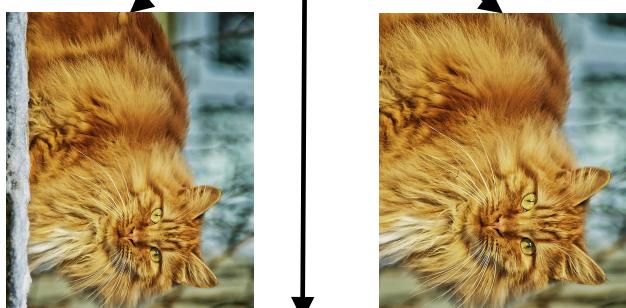
deeplearning.ai

# Common augmentation method

Mirroring



Random Cropping

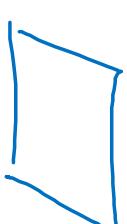


Rotation

Shearing

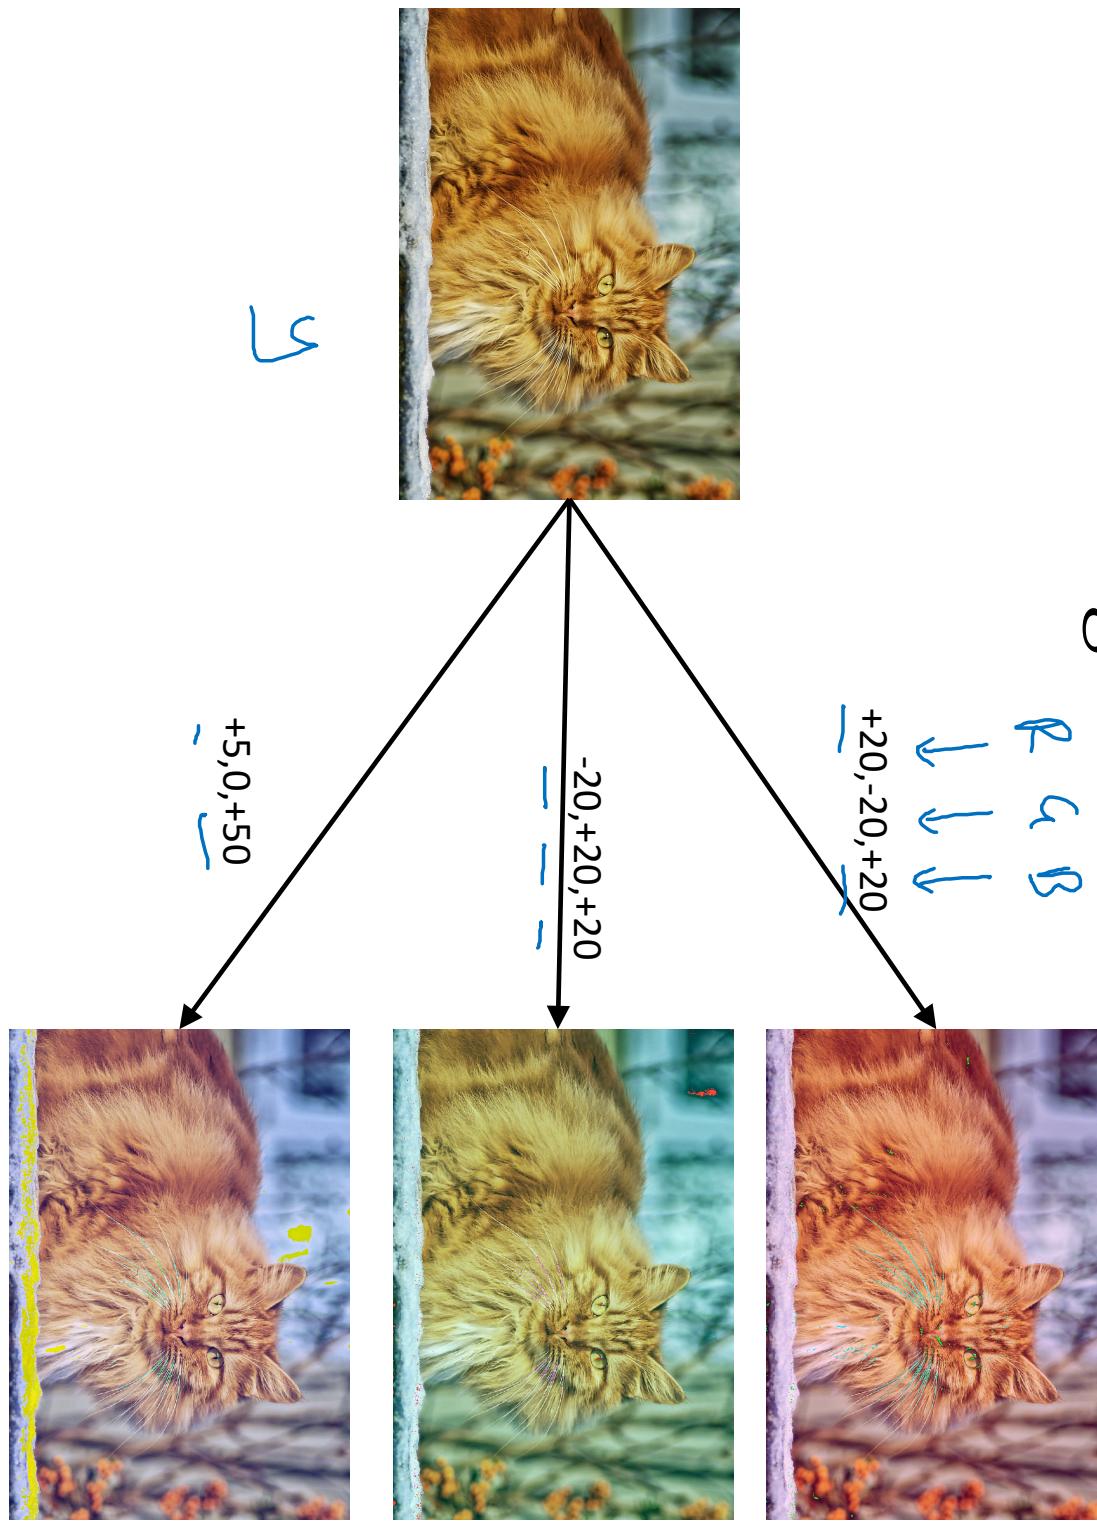
Local warping

...



y

# Color shifting



Advanced:

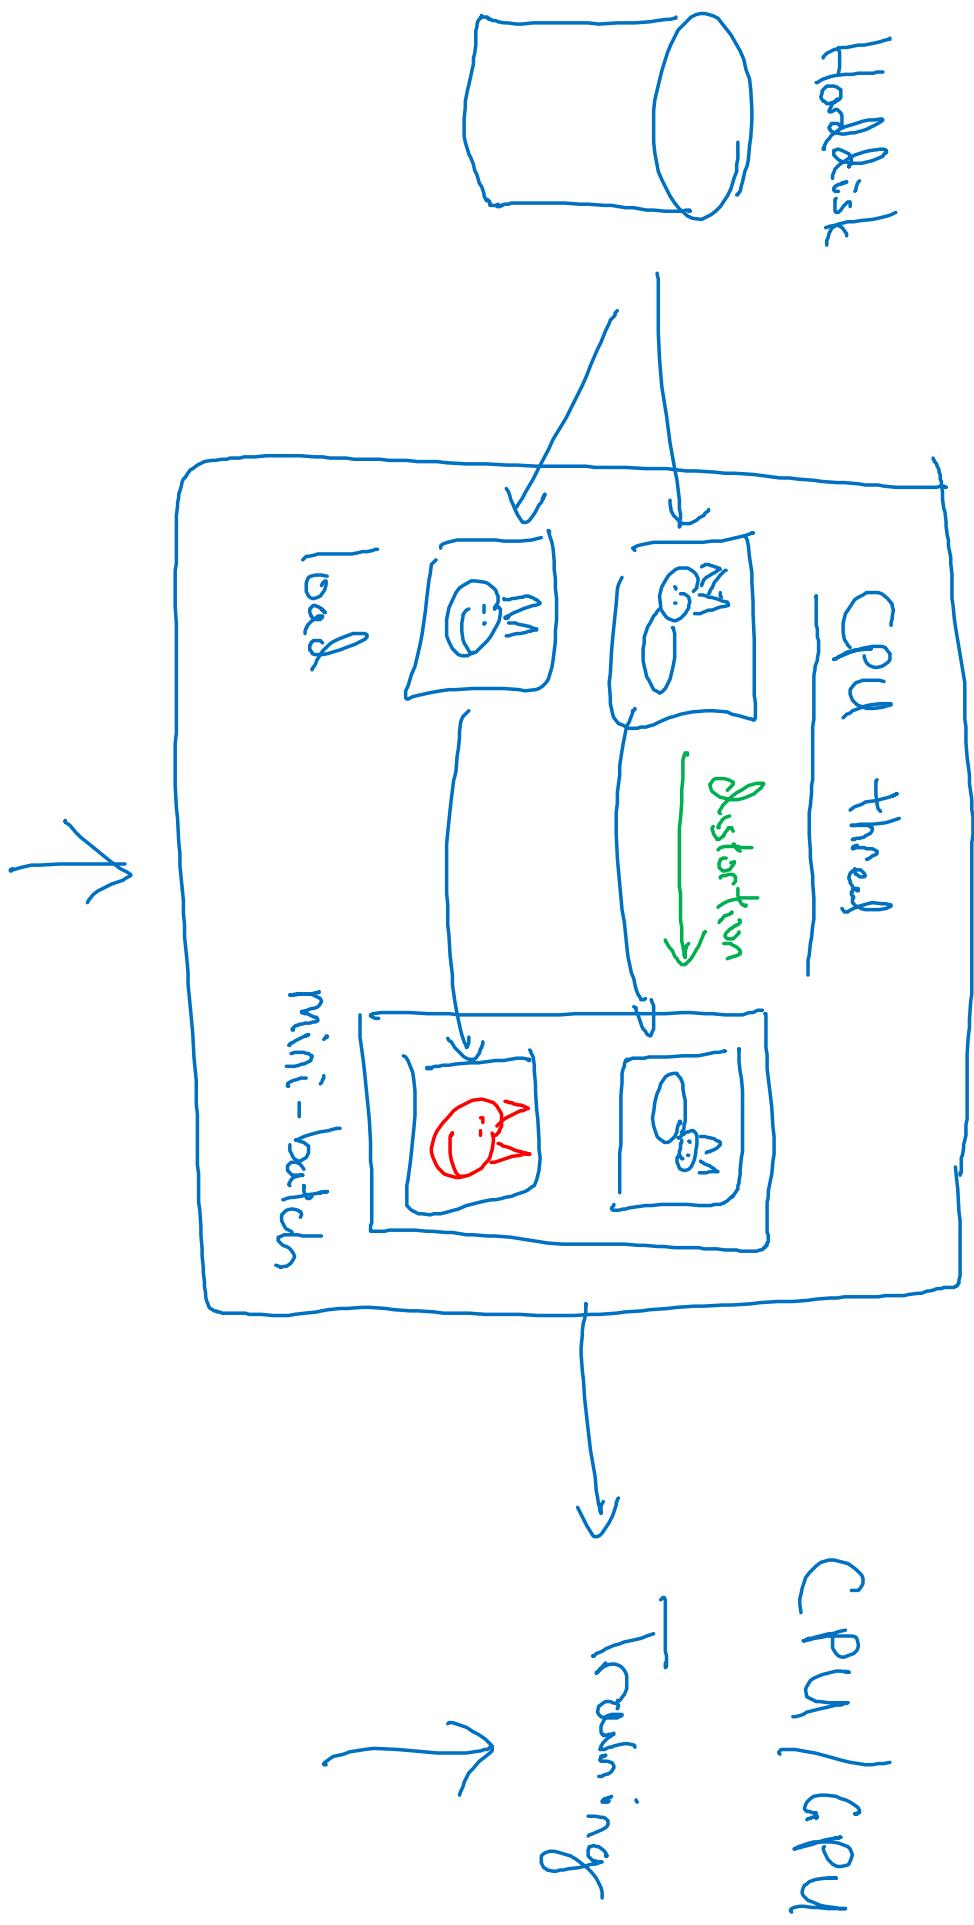
PCA  
ml-class.org

AlexNet Paper

["PCA color  
augmentation"]

R B G

# Implementing distortions during training



Practical advice for

using ConvNets

---

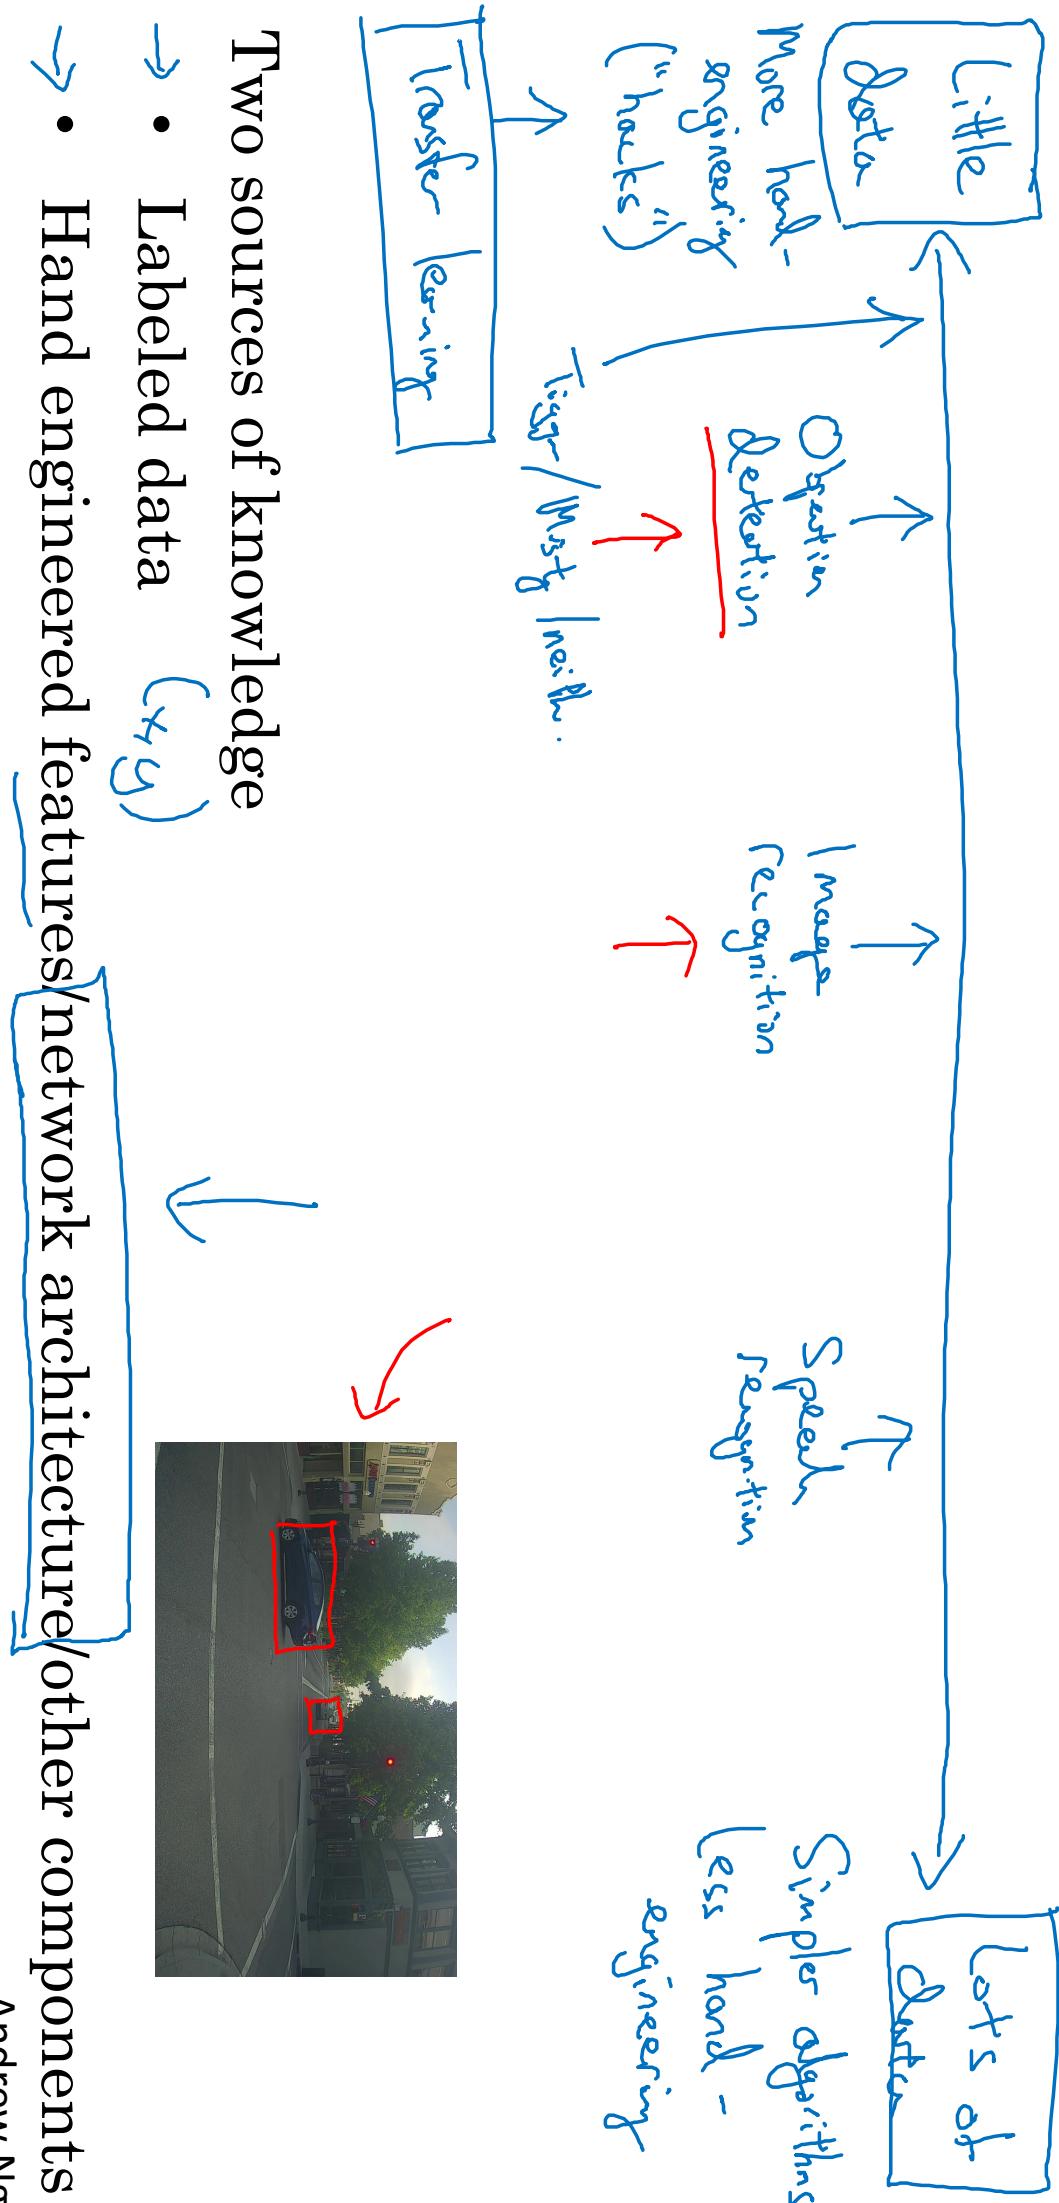
The state of

computer vision

deeplearning.ai



# Data vs. hand-engineering



# Tips for doing well on benchmarks/winning competitions

## Ensembling

- Train several networks independently and average their outputs

3 - 15 networks

$\rightarrow \hat{y}$

## Multi-crop at test time

- Run classifier on multiple versions of test images and average results

10-crop



1

+

4

+

1

+

4

# Use open source code

- Use architectures of networks published in the literature
- Use open source implementations if possible
- Use pretrained models and fine-tune on your dataset